



EE/CSCI 451: Parallel and Distributed Computation

Lecture #2

8/20/2020

Viktor Prasanna

prasanna@usc.edu

ceng.usc.edu/~prasanna

University of Southern California



Outline

- High level view of Processor Core Organization
 - Chapter 2.1, 2.2 in the text
 - (Implicit parallelism)
 - Superscalar
 - Very Long Instruction Word (VLIW) processor
- Multi-core processor organization
- Implicit and Explicit Parallelism



Policies and Procedures

- All classes are recorded and accessible on Blackboard
- Recordings should be used appropriately (<https://policy.usc.edu/scampus-part-c/>)
- Specify your current time zone (and your background info.)
 - Please fill the Google form: <https://forms.gle/yqxoUa7NC7bMq5yv6>



Course Info. (1)

- When & Where
- Lecture: Tuesday, Thursday 3:30 – 4:50 PM, Online
at: <https://usc.zoom.us/j/92090885346?pwd=Ty9PZmFxL2dPVDR3ekNZNmduVFJ0QT09>
 - Meeting ID: 920 9088 5346
- Lab: Friday, 3:30 – 4:50 PM, Online at: <https://usc.zoom.us/j/93866326251>
 - Meeting ID: 938 6632 6251
- Office Hours
 - Tuesday 11:00-12:00 Noon
 - Thursday 10:00-11:00 AM
 - Meeting ID: usc.zoom.us/my/prasanna.zoom



Course Info. (2)

- TA & TA's office hours
 - Meng Yuan
 - Email: ymeng643@usc.edu
 - Office hours:
 - Friday 11:00 AM-1:00 PM or By appointment
 - Meet Yuan at: <https://usc.zoom.us/j/8629150353>
 - Meeting ID: 862 915 0353



Course Info. (3)

- Communicating with me
 - Visit during office hours via zoom
 - Via email prasanna@usc.edu
 - Subject field: EE451

Course Info. (4)



- Grading
 - Homework 10%
 - Homeworks must be done independently
 - **10% late penalty per day** will be assessed with no credit received after the third day
 - Programming Assignments 10%
 - Course Project 15%
 - Midterm I (Sep 25 in lab session, 2 hours) 20%
 - Midterm II (Oct 23 in lab session, 2 hours) 20%
 - Final Exam 25%



Course Info. (5)

- Lab
 - Time: 3:30 – 4:50 PM, Friday
 - Discussion of material covered in the lecture
 - Homework
 - Programming languages
 - Programming assistance
 - Course project discussion
 - Midterms (Sep 25 and Oct 23 in lab session, 330-530pm)
 - Course project presentation



Announcements

- Piazza
 - EE/CSCI 451 Parallel and Distributed Computation Fall 2020
 - Enroll via piazza.com/usc/fall2020/eecsci451
- Lecture notes (ppt) will be uploaded to Piazza one hour before the lecture
- Homework assignments will be uploaded to Piazza
- Please upload your completed homework on Blackboard



Midterms and Final

- Midterm I: 25th Sep in lab session (2 Hours)
- Midterm II: 23rd Oct in lab session (2 Hours)
- Final, as scheduled by the University
 - No makeup exams



Outline



- **High level** view of Processor Core Organization
 - Chapter 2.1, 2.2 in the text
 - (Implicit parallelism)
 - Superscalar
 - Very Long Instruction Word (VLIW) processor
- Multi-core processor organization
- Implicit and Explicit Parallelism

Uniprocessor (Serial Processor)

RAM (Random Access Machine)



A simple view



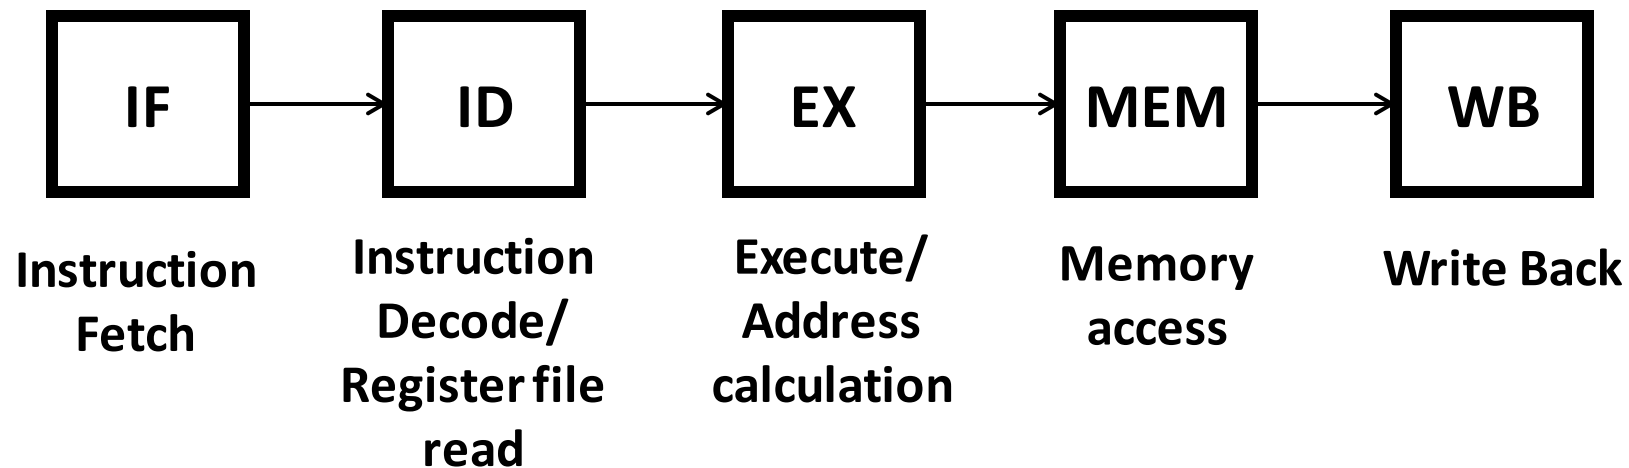
Memory access = 1 cycle for any location
Execution = 1 cycle



Pipelined implementation of processor (1)

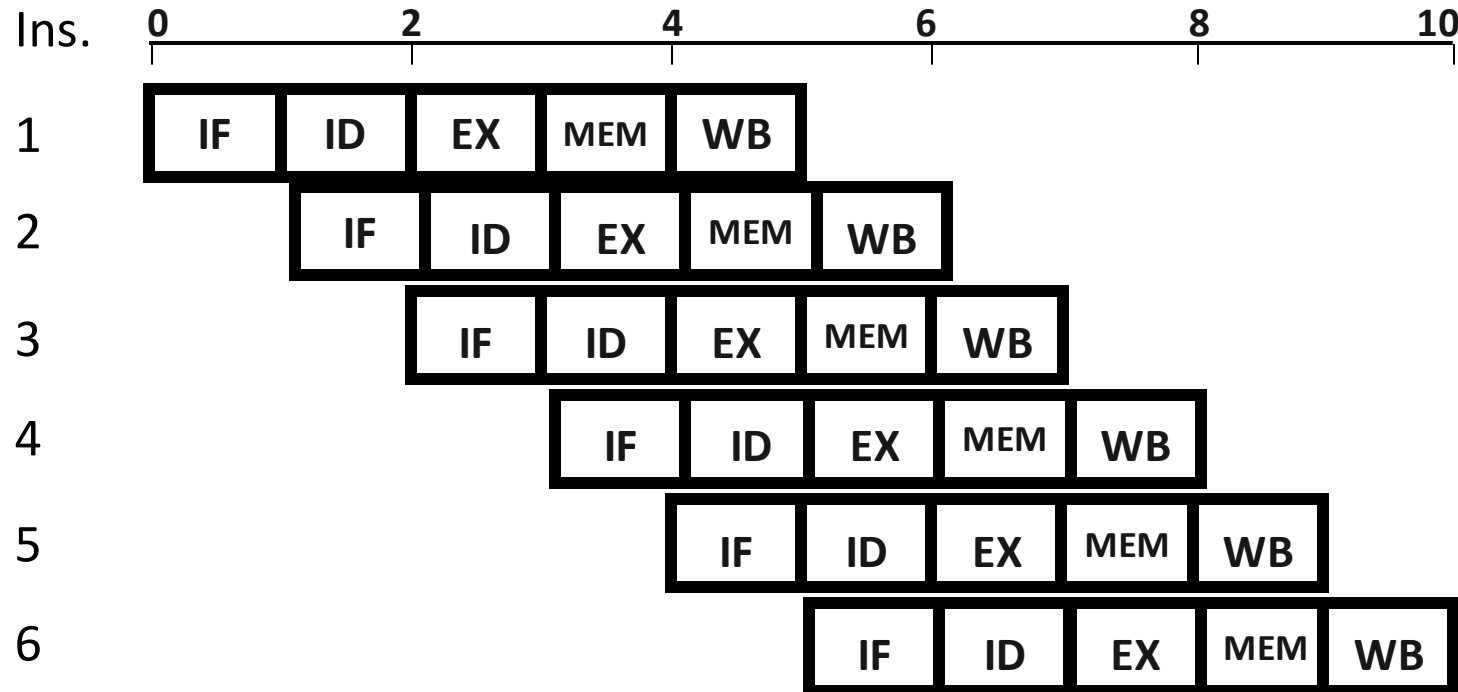
All (modern) processors are pipelined

Five stage pipeline (For example, Hennessey Patterson, EE 457)





Pipelined implementation of processor (2)



1. Load R1, @1000
2. Load R2, @1008
3. Add R1, 100
4. Add R2, 104
5. Add R1, R2
6. Store R1, @2000

Total # of cycles to complete execution = 10

For executing n instructions, $5n \geq \text{total \# of cycles} \geq n$

Ideally, total # of cycles = $5 + (n - 1)$



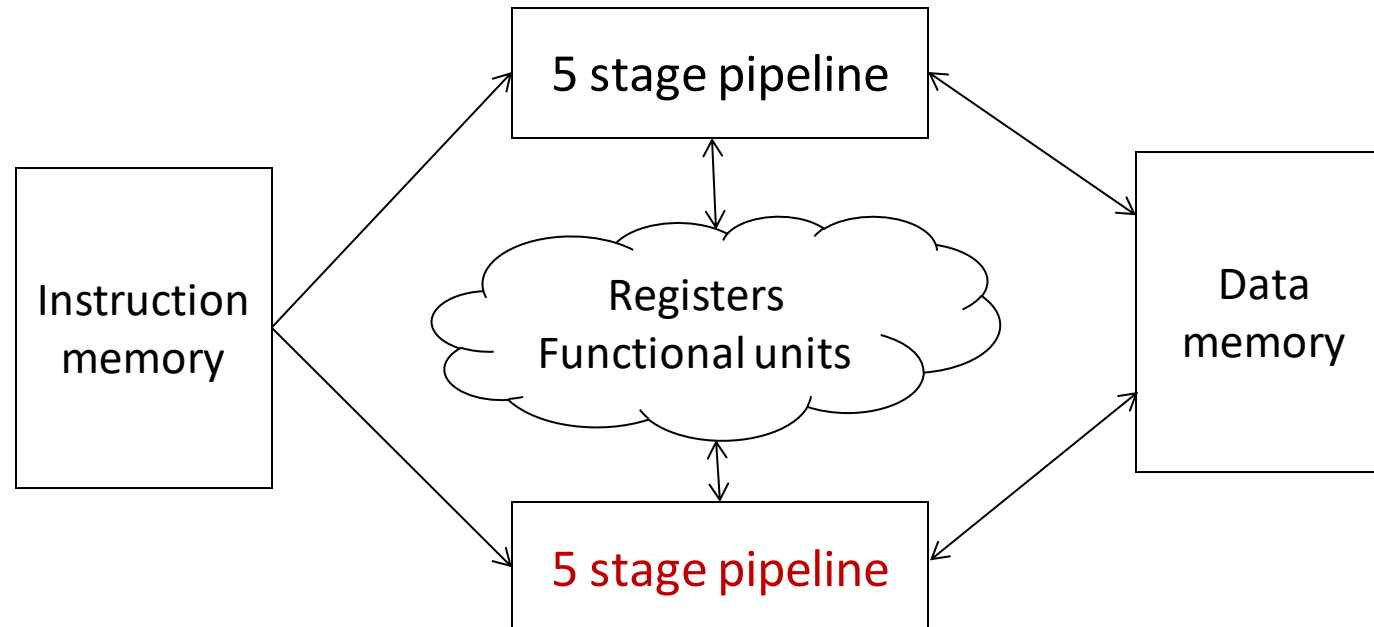
Simple Performance Model

- Instruction exec. rate = clock rate = peak performance ?
- Time to execute n ins. $\geq n / \text{clock rate}$
- Example: 2 GHz processor core (single core)
- **Peak** performance = 2×10^9 Ins. per second (MIPS rating)



Superscalar execution

Example: 2 pipelines (Sharing some hardware resources)
2× improvement?



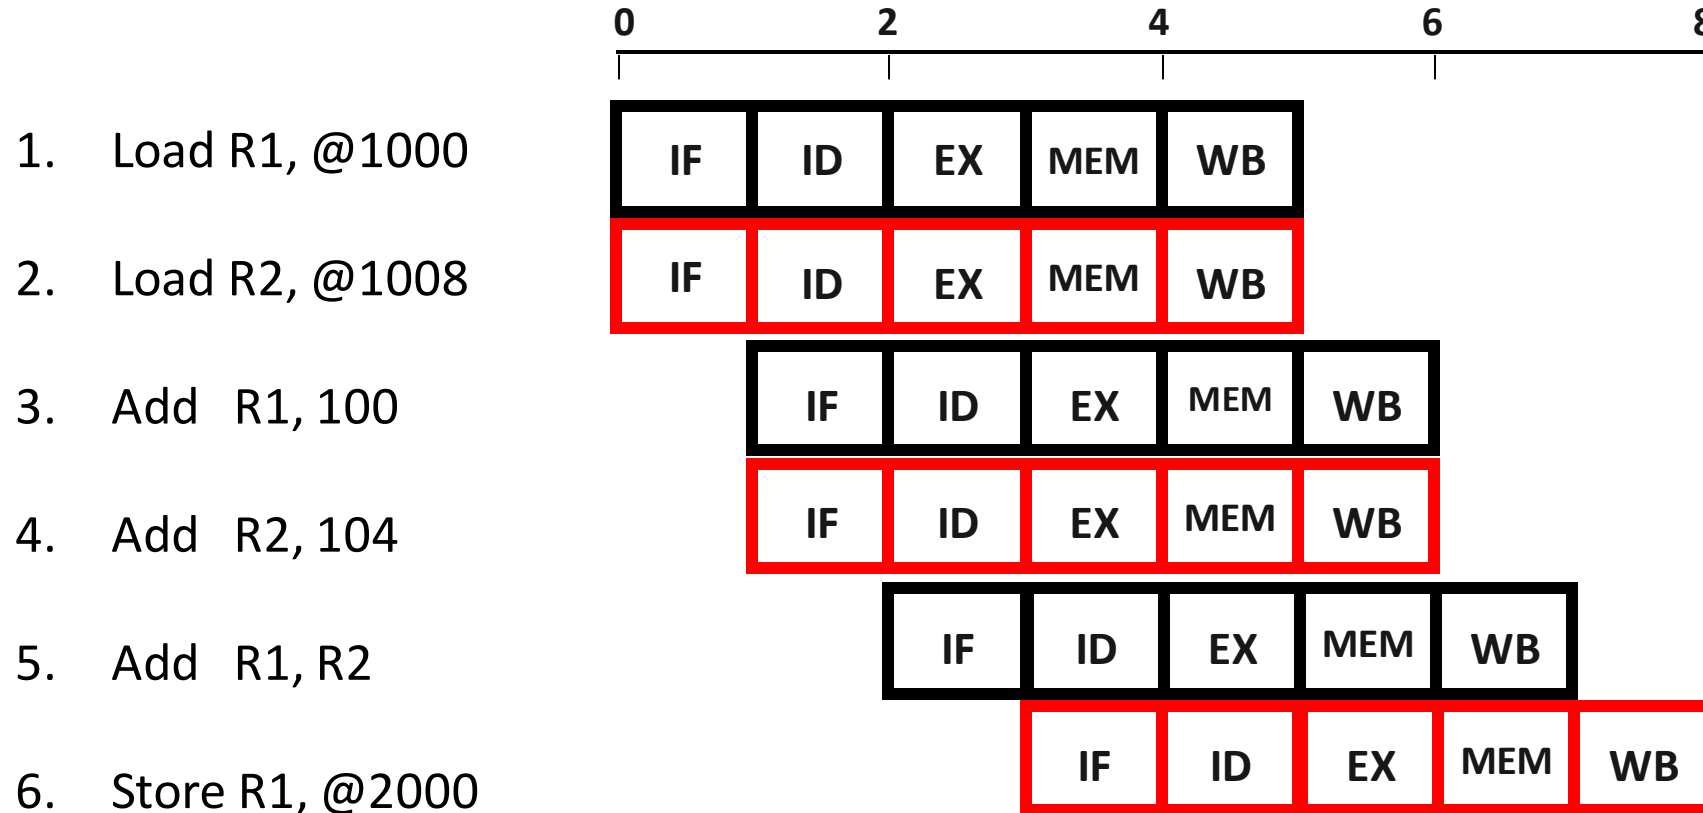
Implicit Parallelism



Implicit Parallelism

- 5 stage pipeline
 - $5n \rightarrow 5 + n - 1$ cycles
- k-way superscalar
 - $5 + n/k - 1$?

2-way super scalar execution (almost 2 Ins. Issued per cycle)



Total # of cycles = 8

Note: Ins. 5 and 6 cannot be issued in the same cycle



Data dependencies (1)

Output of Ins i is an input to Ins k ($k > i$)

Potential problem: Ins k is close to Ins i **during** execution (during execution Ins k and Ins i are in the hardware pipeline concurrently)

Example: Assume Single Pipeline

1. Load R1, @1000
2. Load R2, @1008
3. Add R1, 100
4. Add R2, 104
5. Add R1, R2



Simple solution = Stall the pipeline -> Performance loss

Hardware solution = Internal Forwarding


- Ex.: Dependency between Ins. 4 and 5 can be handled without stalling the pipeline.



Data dependencies (2)

- Data dependencies can have significant impact on performance
- We assumed a simple model of execution: **one cycle to access memory**
- In reality, main memory (DRAM) access latency is (very) high (50-100 cycles)


Example: (Execution of Ins 3 may have to be delayed by 10's of cycles!)

1. Load R1, @1000
 2. Load R2, @1008
 3. Add R1, 100
 4. Add R2, 104
 5. Add R1, R2
- 



Out of Order Issue

Example: 2-way superscalar

1. Load R1, @1000
 2. Add R1, 100
 3. Load R2, @1008
 4. Add R2, 104
- 

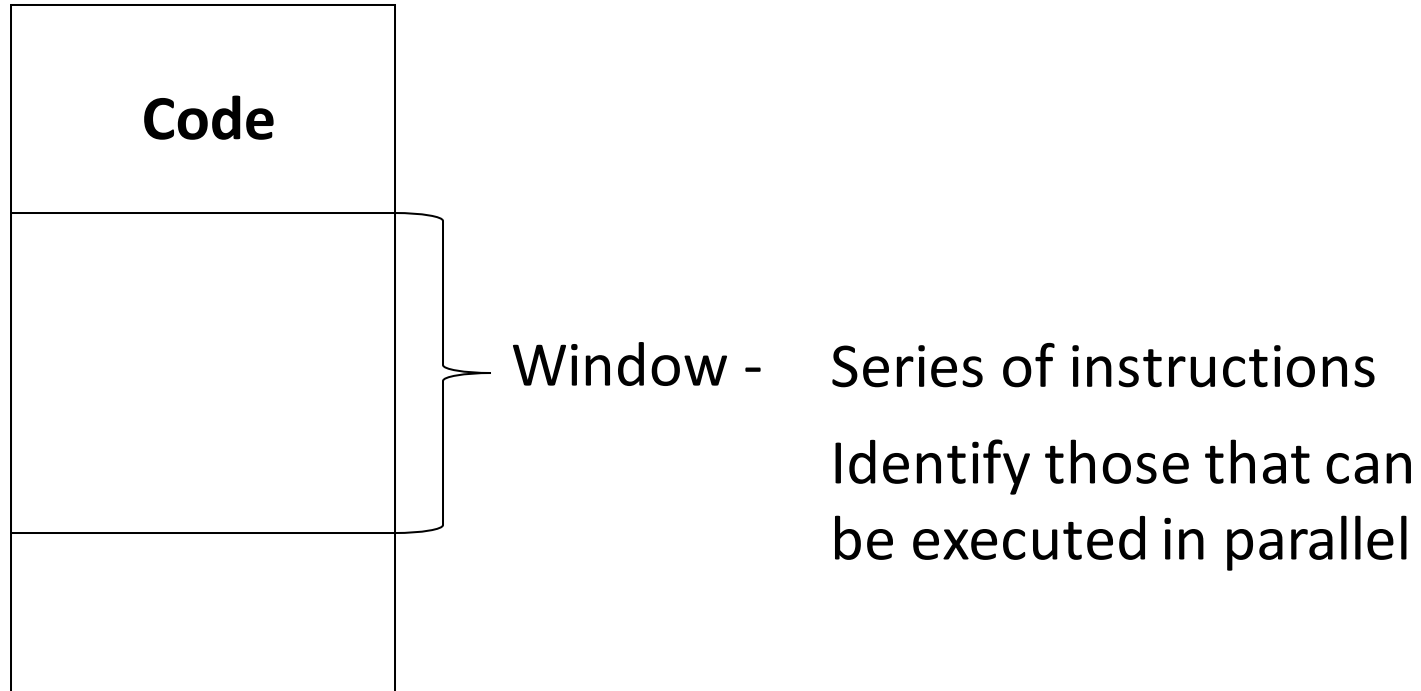
- Cannot issue Ins. 1, 2 simultaneously
- Cannot issue Ins. 3, 4 simultaneously

However, can do: Initiate Ins. 1, Ins. 3 in clock cycle 0

Initiate Ins. 2, Ins. 4 in clock cycle 1



Dynamic Instruction Issue




Note: Hardware does scheduling at run time

Way of exploiting instruction level parallelism



Resource dependency

Example: Suppose 2-way super scalar execution
one hardware Floating Point Unit (FPU)



[Load R1, @1000
Load R2, @1004
FP Mult R1, 100
FP Mult R2, 104

Execution hardware detects such dependencies
and schedules the execution accordingly

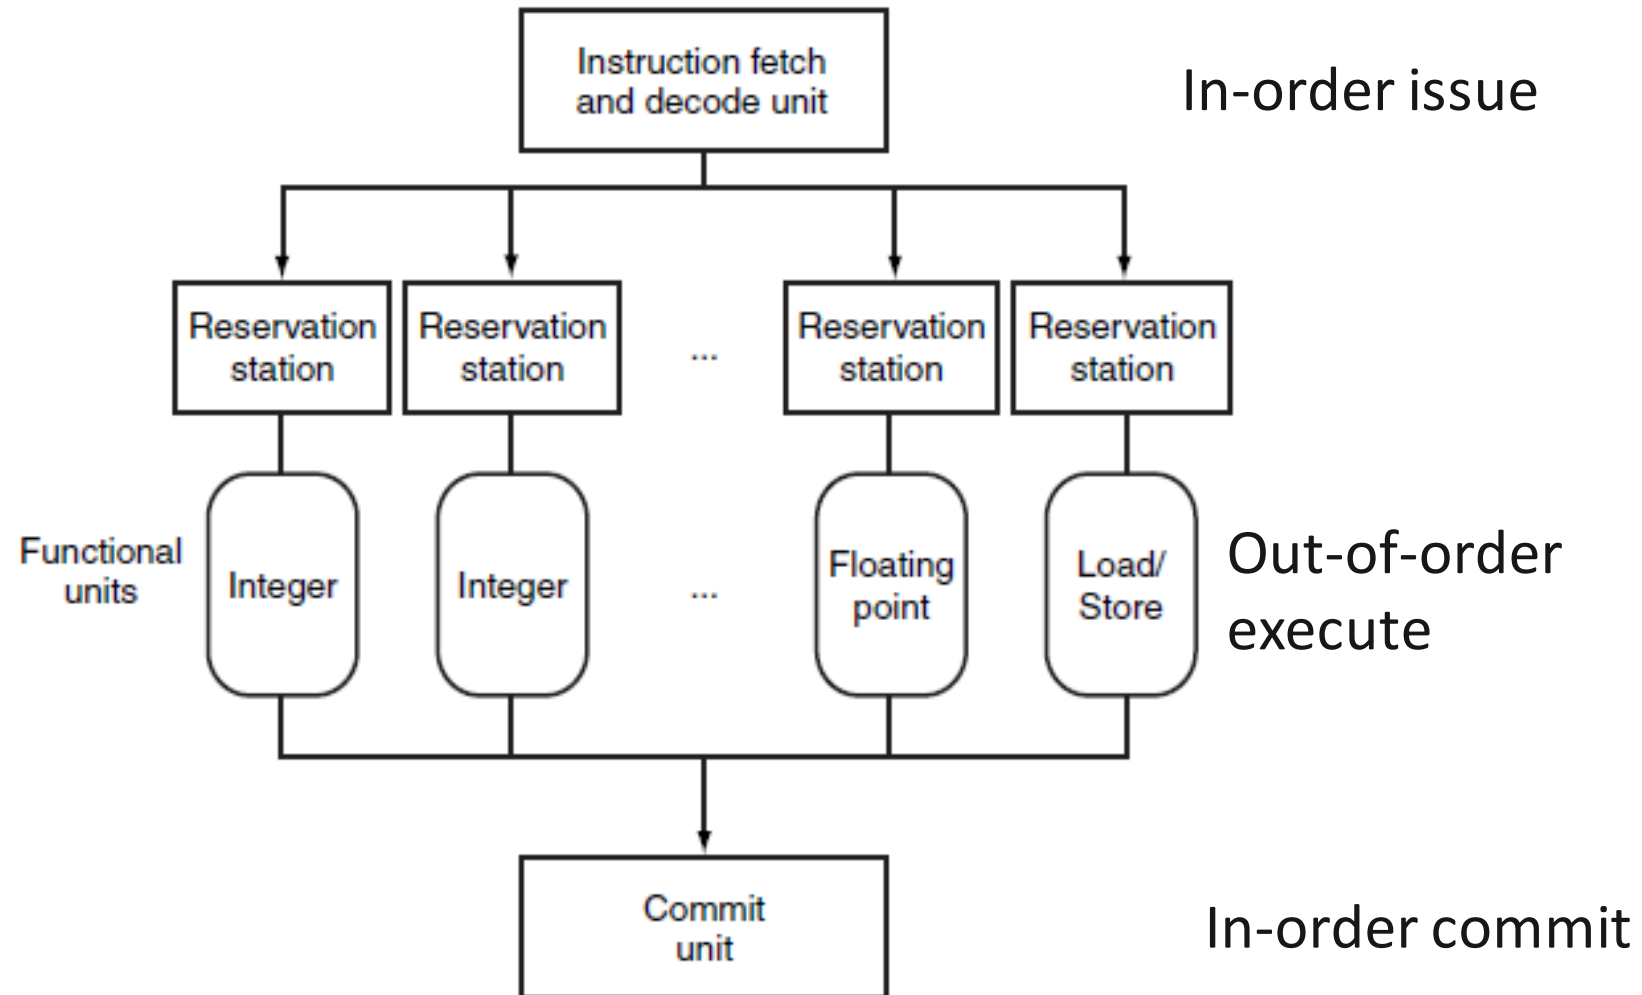


Superscalar (1)

- Instruction-level parallelism
 - Multiple function units (Ex. arithmetic logic unit, bit shifter)
 - One or more instructions issued per clock
 - In-order issue
 - Out-of-order execution
 - In-order commit



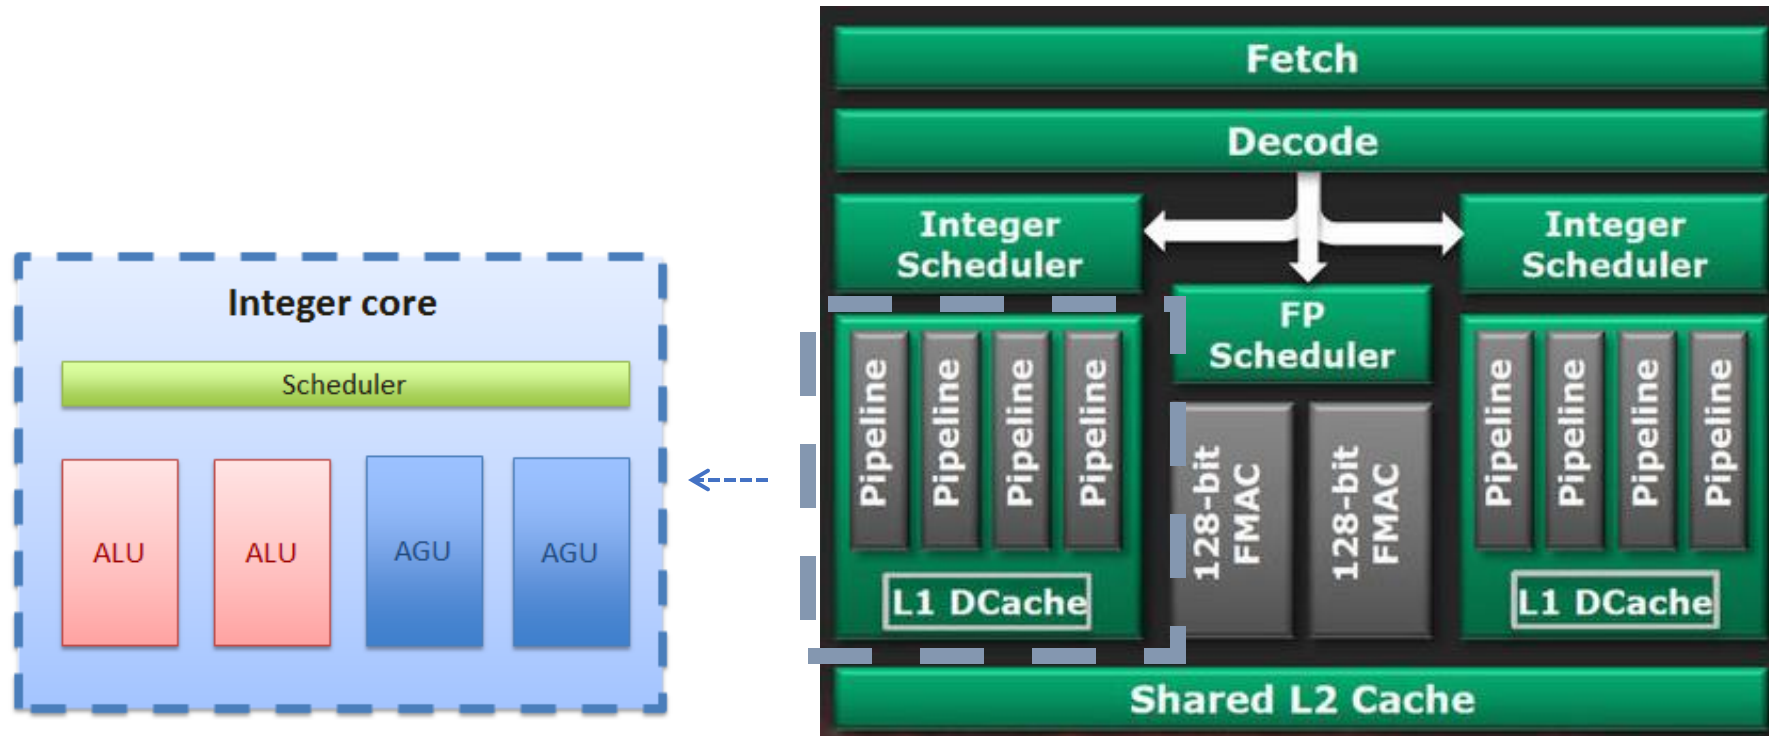
Superscalar (2)





Example Superscalar: Zen2 Architecture

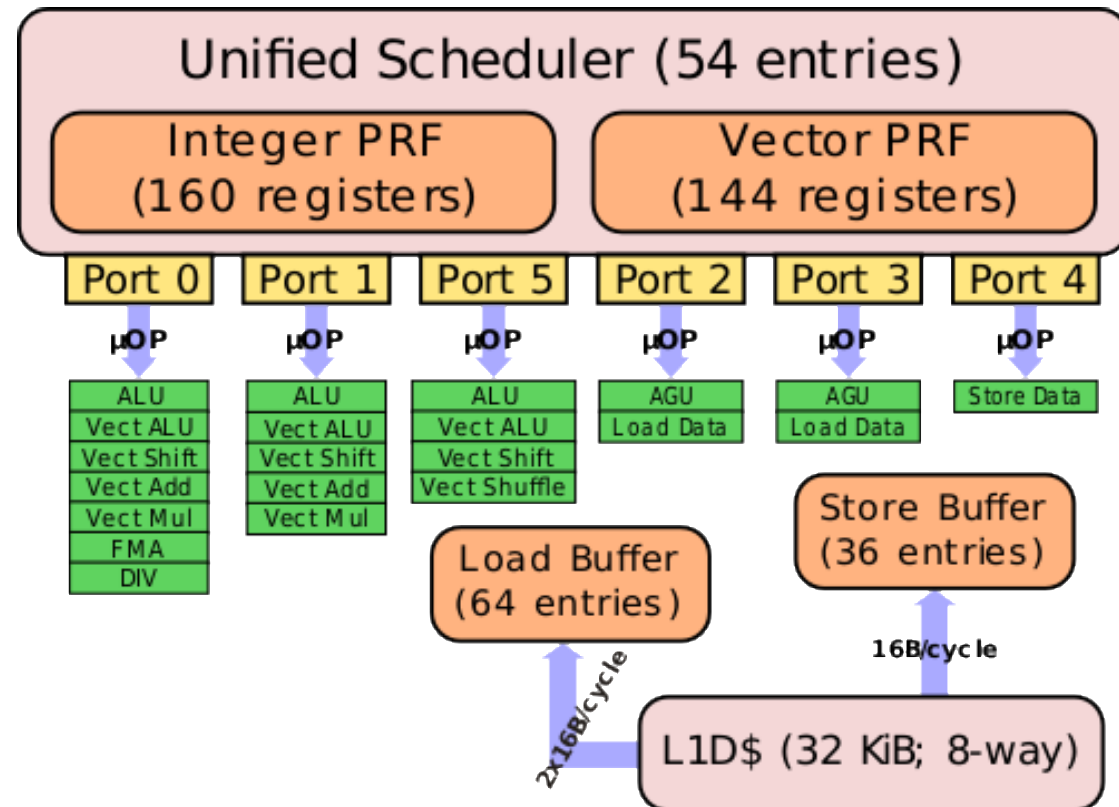
- Can handle 4 independent arithmetic and memory operations per clock for each integer core
- Issue bandwidth = 4





Example Superscalar: Sandy Bridge

- Microarchitecture by Intel
- Can handle up to 5 independent arithmetic and memory operations per clock for each core






Very Long Instruction Word (VLIW) Processor

Superscalar processor

- Dynamic scheduling of Ins.
- Scheduling **hardware** is expensive
 - Dependence window size \times # of pipelines

VLIW: Do scheduling at compile time (in **software**)



Specify parallelism as parallel activities (using long word instruction format)





Processor Execution Model

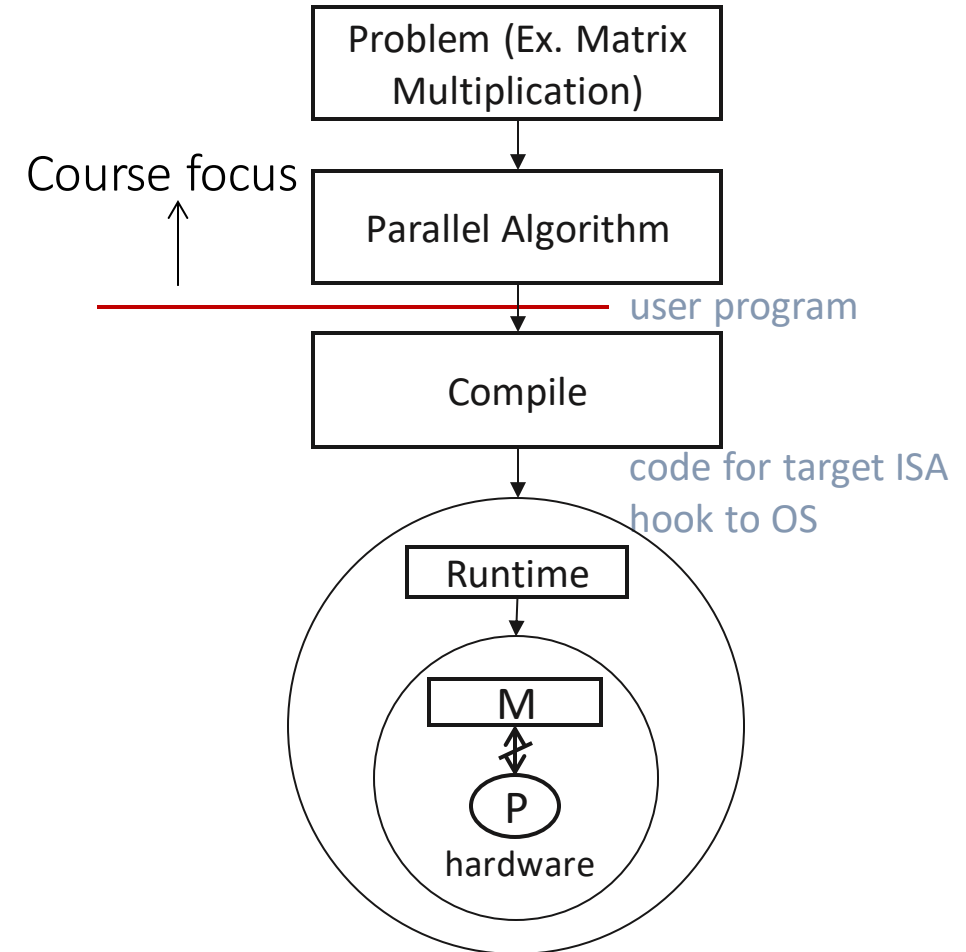
- SISD – Single Instruction Single Data
 - (Random Access Machine -- RAM)
- Instruction Set Architecture (ISA)
 - Abstract model of a computer, exposed to the users, compiler





Designing and Executing a Parallel Program


- Performance metric
 - **Execution time**
- Execution time depends on
 - Compile time optimizations (Compiler level)
 - Execution time optimizations (Processor level—Implicit)
- Many optimizations are possible at architecture and hardware level (EE 457, EE 557)





Example Multi-core Processor (1)

- AMD Ryzen 9 3900X processor
 - 12 cores
 - 24 threads
 - 3.8 GHz
 - L3 cache: 64 MB
 - Memory: 2 Channels
 - PCIe: 24 Gen4 lanes






Multi-core Processor (2)

- Parallelism
 - Pipelined execution of Instructions (in each core)
 - Instruction-level parallelism (in each core)
 - Thread-level parallelism (across cores)
- Data parallelism
- Loop parallelism



Multi-core Processor (3)

- Data parallelism
 - Same operation on large data set
 - 512-bit Advanced Vector Extensions instructions (AVX-512) -- eight 64-bit or sixteen 32-bit integer operations per clock cycle





Multi-core Processor (4)

- Loop parallelism
 - Most execution time (of a serial program) is spent in loops
 - (90/10 rule)—90% of the time in 10% of the code
 - Extract parallel tasks from loops (for example, unrolling the loop)
 - Dynamically (hardware)
 - Statically by compiler (software)
 - Need to determine instruction dependence (data dependencies)
 - In C++: (compiler directives)
 - `#pragma unroll`
 - `#pragma unroll(n)`
 - `#pragma nounroll`




Explicit parallelism for Matrix Multiplication (1)

- Matrix Multiply

$C = A \times B$ (A, B, C are $n \times n$ matrices with $C[:,:] = 0$)

```
for i = 1 to n
  for j = 1 to n
    for k = 1 to n
      D(i,j,k) = A(i,k) × B(k,j)
      C(i,j) = C(i,j) + D(i,j,k)
```



$2n^3$ operations



Explicit parallelism for Matrix Multiplication (2)

- **Data parallelism**

Compute $A(i,k) \times B(k,j)$ for various k in parallel within one instruction

- AVX-512 with 64-bit integer (for $k=1,9,17,\dots$):

- Can do 8 ops/ins using 64 bit arithmetic
- No data dependencies

A		$A(i,k)$	$A(i,k+1)$	$A(i,k+2)$	$A(i,k+3)$	$A(i,k+4)$	$A(i,k+5)$	$A(i,k+6)$	$A(i,k+7)$
B	\times	$B(k,j)$	$B(k+1,j)$	$B(k+2,j)$	$B(k+3,j)$	$B(k+4,j)$	$B(k+5,j)$	$B(k+6,j)$	$B(k+7,j)$
D	$=$	$D(i,j,k)$	$D(i,j,k+1)$	$D(i,j,k+2)$	$D(i,j,k+3)$	$D(i,j,k+4)$	$D(i,j,k+5)$	$D(i,j,k+6)$	$D(i,j,k+7)$

- Compute $C(i,j)$ from $D(i,j,k)$
- Explicit (parallel) program specification to exploit Implicit Parallelism (AVX-512) of the architecture



Explicit parallelism for Matrix Multiplication (3)

- **Loop parallelism**

Compute $A(i,k) \times B(k,j)$ for various k in parallel using multiple threads

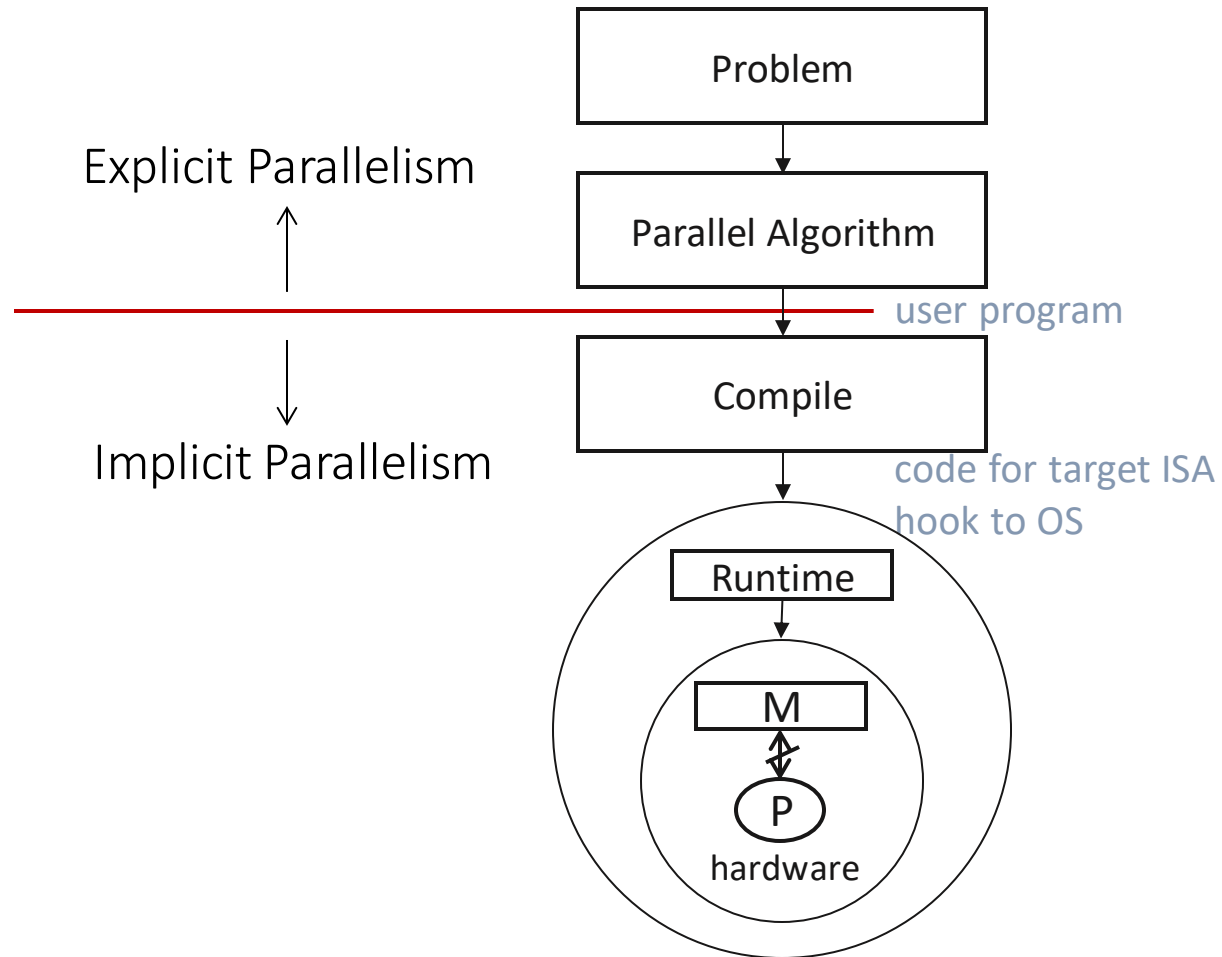
- Assume a two-thread multicore:

Restructured loop program (by a compiler)

```
for i = 1 to n
  for j = 1 to n
    for k = 1,3,...,n-1
      D(i,j,k) = A(i,k) × B(k,j) //thread 0
      D(i,j,k+1) = A(i,k+1) × B(k+1,j) //thread 1 } compute in parallel, no dependencies
      C(i,j) = C(i,j) + D(i,j,k)
      C(i,j) = C(i,j) + D(i,j,k+1) } thread coordination, data dependencies--can be handled
```



Summary (1)





Summary (2)

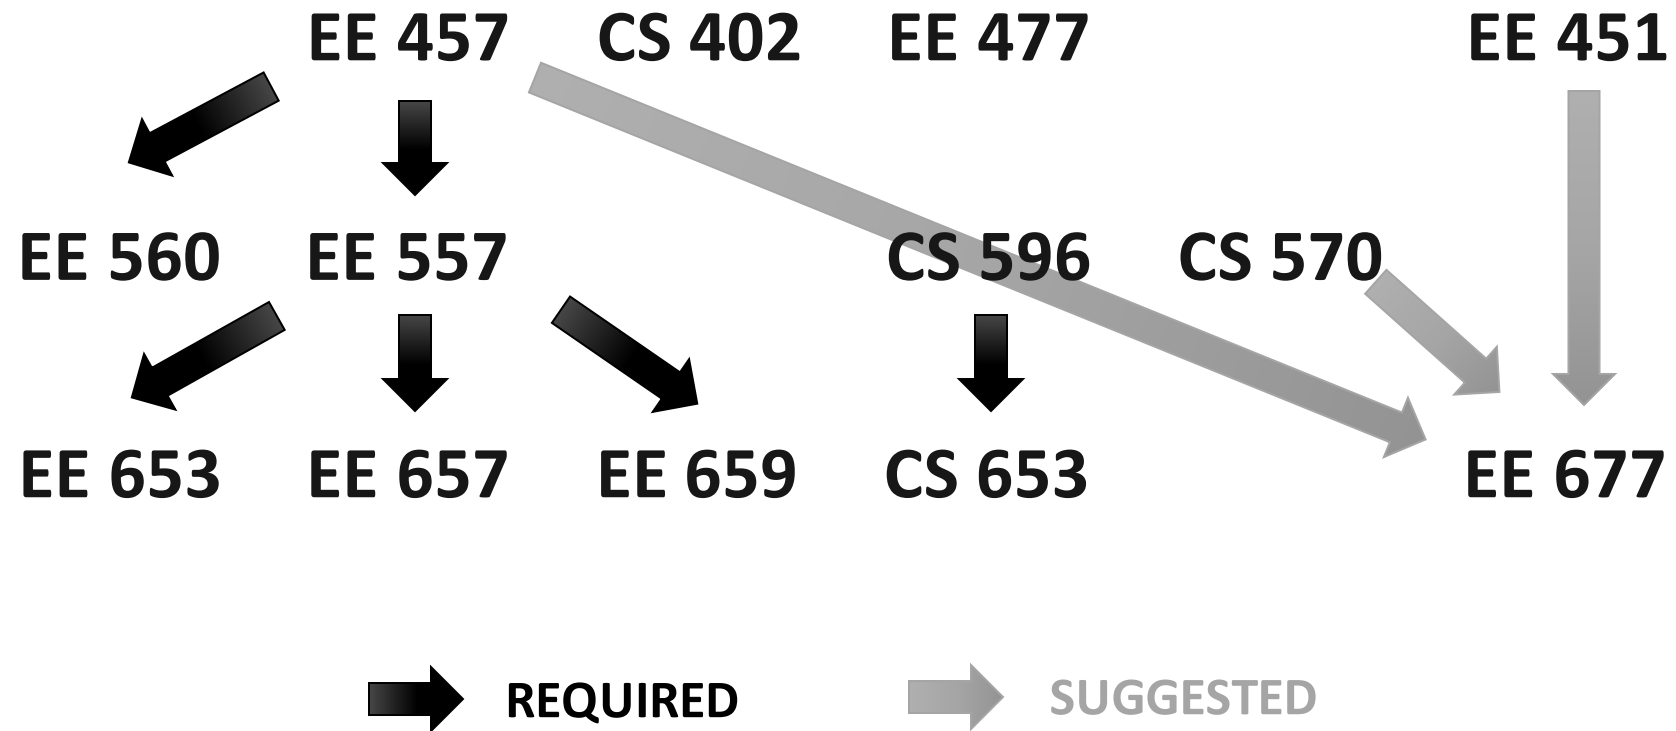
- SISD execution model
 - Processor core organization
 - Pipelined execution
 - Data dependencies
 - Multi-core architecture
 - Explicit and Implicit Parallelism
-
- Course Focus: Explicit Parallelism



Backup Slides




Course Dependencies





Challenges (1)


- Processor - Memory performance gap





Challenges (2)

- Limit to achievable speedup
 - Amdahl's Law



$$\text{Speedup} = \frac{\text{Serial time}}{\text{Parallel time}} = \frac{1}{\alpha}.$$

Example: 50 % of the code is parallizable,

Speedup ≤ 2 (no matter how many processors are used)




Challenges (3)

- Co-ordination

Do in parallel
 $P_1,$
 $P_2,$
 $\dots,$
 P_k
end

Synchronization


overhead = ?





Challenges (4)


- Communication
 - Compute
 - Communicate
 - Wait?





Challenges (5)

- Load balance
 - Example: Matrix multiplication
 - P_{ij} – compute output matrix ij , $0 \leq i, j < \sqrt{p}$
 - $p = \text{number of processors}$





Challenges (6)

- Key performance metrics
 - $Speedup = \frac{Serial\ time}{Parallel\ time}$
 - Given $p > 1$ processors, can we get p fold improvement performance?
 - Scalability



Course Outline (1)

- Overview / Course Introduction
- Architectural Principles for Application Developers
 - Pipelined processor organization: data and control hazards, ILP, out of order execution, multithreading.
 - Memory systems: DRAM organization, cache organization.
 - Impact on software performance, locality, multithreading, prefetching.



Course Outline (2)

- Analytical Models for Parallel Systems
 - Architecture performance metrics: CPI, MIPS, SpecMark. Software performance benchmarks: Peak performance, sustained performance, LinPack, Bandwidth benchmarks.
 - Limits on achievable performance, Amdahl's Law, scalability definitions, work optimality, Iso-efficiency function, Order notation.



Course Outline (3)

- Analytical Models for Parallel Systems(cont.)
 - Communication costs in parallel machines: start-up cost, throughput, latency. Routing mechanisms: packet routing, cut through, virtual channels. Modeling message passing and shared address space machines. Data layouts and graph embeddings.
 - Multi-core, many-core and GPU architectures, data parallel programming abstraction of GPUs.



Course Outline (4)

- PRAM and Data Parallel Algorithms
 - PRAM model of computation, Brent's theorem, other models of computation, illustrative examples
 - Max, Scan operations
 - Recursive doubling, graph algorithms
 - Performance analysis, scalability, relation-ship to practical platforms
 - Sorting, FFT



Course Outline (5)

- Basic Communication Primitives
 - Broadcast and all-to-all, communication costs on various topologies
 - Personalized communication, Reduce, prefix sum, and scatter and gather
 - Graph embeddings



Course Outline (6)

- Message Passing Programming Model
 - Message passing abstraction, send receive primitives, blocking and non-blocking commands, collective operations
 - Illustrative examples: Canon's algorithm, overlapping computation and communication, Odd-even merge sort, Review for Midterm 2
- Shared Address Space Programming Models
 - Pthreads , OpenMP, Illustrative examples



Course Outline (7)

- Parallel Dense Algebra
 - Matrix vector, matrix matrix computations
 - Parallel Dense Algebra
- Parallel Search and Sorting
 - Parallel search, illustrative example applications, throughput optimization, Multi-dimensional search, decision tree and decomposition
 - Sorting techniques, Bitonic sort, row-column sort, Mapping onto parallel architectures



Course Outline (8)

- Cloud, Big Data and Map Reduce
 - Cloud as a computing platform, Large data sets and organization
 - Discovery of knowledge, annotation of data
 - Computational and programming models for Big Data
 - Map Reduce as a parallel programming model, Hadoop, Spark
 - Illustrative examples from data science domain



Acknowledgement

- Lecture notes compiled over the years
- My research team
- Course TA and mentors during Spring '14, '15, '16, and '17
- EE451 Students in Spring '14, '15, '16, and '17
- Introduction to Parallel Computing (2nd Ed.), Vipin Kumar, Ananth Grama
-



Research

- Research Projects
 - High Performance Networking
 - Big Data Analysis
 - Social Network Analysis
 - Energy Efficient Computing
- Webpages
 - ceng.usc.edu/~prasanna/
 - <http://fpga.usc.edu/>
 - <http://pdc.usc.edu/>
 - <http://dslab.usc.edu/>



Related Courses

- EE/CSCI 451 Introduction to Parallel and Distributed Computation (Prasanna)
- EE 454 Introduction to Systems-on-Chip (Bogdan)
- EE 457 Computer Systems Organization (Puvvada)
- CSCI 402 Operating Systems (Cheng)
- EE 557 Computer Systems Architecture (Annavaram/Dubois)
- EE 560L Digital System Design — Tools and Techniques (Puvvada)
- EE 532 Wireless Internet and Pervasive Computing (Hwang)
- EE 533 Network Processor Programming and Design (Cho)
- CSCI 558L Internetworking and Distributed System Laboratory (Cho)
- CSCI 570 Analysis of Algorithms (Shamsian/Adamchik)
- CSCI 596 Scientific Computing and Visualization (Nakano)
- EE 653 Advanced Topics in Microarchitecture (Dubois)
- CSCI 653 High Performance Computing and Simulations (Nakano)
- EE 659 Interconnection Networks (Pinkston)
- EE 657 Parallel and Distributed Computing (Dubios)
- EE 677 VLSI Architectures and Algorithms (Prasanna)

Sample course work with focus on Computer Architecture



- EE/CSCI 451 Introduction to Parallel and Distributed Computation (Prasanna)
- EE 454 Introduction to Systems-on-Chip (Bogdan)
- EE 457 Computer Systems Organization (Puvvada)
- CSCI 402 Operating Systems (Cheng)
- EE 557 Computer Systems Architecture (Annavaram/Dubois)
- EE 533 Network Processor Programming and Design (Cho)
- EE 599 Cyber-Physical Systems (Bogdan)
- CSCI 570 Analysis of Algorithms (Shamsian/Adamchik)
- EE 653 Advanced Topics in Microarchitecture (Dubois)
- EE 659 Interconnection Networks (Pinkston)
- EE 677 VLSI Architectures and Algorithms (Prasanna)

Sample course work with focus on Hardware, embedded systems



- EE/CSCI 451 Introduction to Parallel and Distributed Computation (Prasanna)
- EE 454 Introduction to Systems-on-Chip (Bogdan)
- EE 457 Computer Systems Organization (Puvvada)
- EE 477 MOS VLSI Circuit Design (Nazarian)
- EE 577a VLSI System Design (Pedram/Nazarian)
- EE 577b VLSI System Design (Pedram)
- EE 560L Digital System Design — Tools and Techniques (Puvvada)
- EE 532 Wireless Internet and Pervasive Computing (Hwang)
- EE 533 Network Processor Programming and Design (Cho)
- EE 599 Cyber-Physical Systems (Bogdan)
- CSCI 570 Analysis of Algorithms (Shamsian/Adamchik)
- EE 677 VLSI Architectures and Algorithms (Prasanna)

Sample course work with focus on parallel and scientific computing



- EE/CSCI 451 Introduction to Parallel and Distributed Computation (Prasanna)
- EE 454 Introduction to Systems-on-Chip (Bogdan)
- EE 457 Computer Systems Organization (Puvvada)
- CSCI 402 Operating Systems (Cheng)
- EE 557 Computer Systems Architecture (Annavaram/Dubois)
- EE 542 Internet and Cloud Computing (Hwang)
- CSCI 570 Analysis of Algorithms (Shamsian/Adamchik)
- CSCI 596 Scientific Computing and Visualization (Nakano)
- CSCI 653 High Performance Computing and Simulations (Nakano)

Sample course work with focus on distributed systems and big data



- EE/CSCI 451 Introduction to Parallel and Distributed Computation (Prasanna)
- EE 450 Introduction to Computer Networks (Zahid/Touch)
- EE 542 Internet and Cloud Computing (Hwang)
- CSCI 402 Operating Systems (Cheng)
- CSCI 555 Advanced Operating Systems (Govindan)
- CSCI 567 Machine Learning (Sha/Liu)
- CSCI 570 Analysis of Algorithms (Shamsian/Adamchik)
- CSCI 585 Database Systems (Raghavachary/Ghyam)
- CSCI 657 Advanced Distributed Systems (Lloyd)