



EE/CSCI 451: Parallel and Distributed Computation

Lecture #3

8/25/2020

Viktor Prasanna

prasanna@usc.edu

ceng.usc.edu/~prasanna

University of Southern California



Outline

- From last class
 - Processor organization
 - Implicit parallelism
- Today (Chapter 2.2)
 - Memory systems
 - Latency, Bandwidth: performance implications
 - Cache: impact on performance
 - Data layout
 - Latency Hiding: Multithreading, Prefetching




Announcements

- Homework 1
 - Due: September 1 2020, midnight AOE (Any time On Earth)
 - Submit typewritten/scanned copy via blackboard
- Programming Homework 1
 - Due: August 28 2020, midnight AOE
 - Submit zip file via Blackboard



Memory system

- System performance (execution time) depends on the “rate” at which data can be accessed by the program

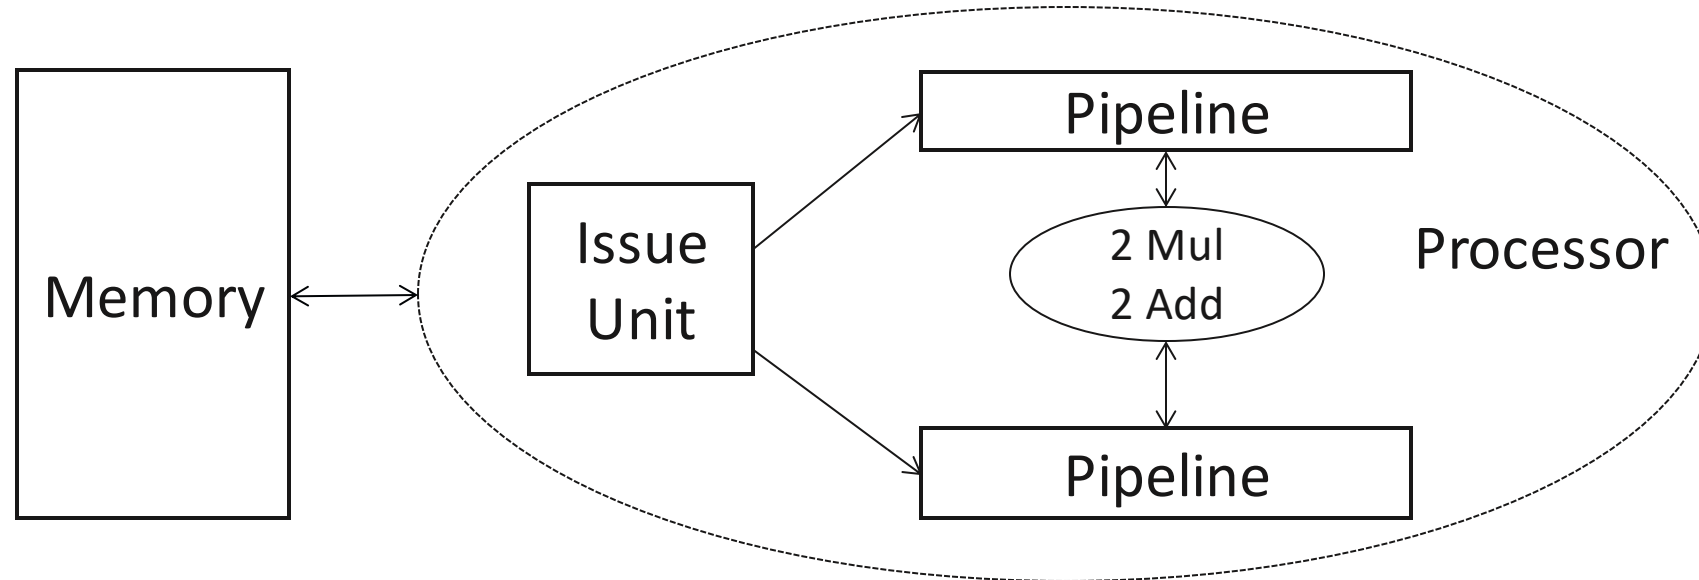


- **(DRAM) Hardware Latency** ~ 10 ns
- **(Peak) Processor-Memory Bandwidth** ~ 64 bits at 1 GHz
(64 Gbits/sec or 8 GB/sec)
- Single cycle processor

- Effective (DRAM) bandwidth depends on access stride
e.g. access consecutive addresses $i, i + 1, \dots$ [0,1,2,3,...]
or arbitrary sequence $i_1, i_2, \dots, i_k, \dots$ [0,10,50,2,...]



System Organization



Performance of an application depends on

- Memory b/w and latency
- # of pipelines
- # of Functional Units (FPUs)



Example memory system performance (1)

Memory

Latency = 10 ns

(Peak) Bandwidth = 64

bits at 1 GHz

(64 Gbits/sec or 8 GB/sec)

(bus frequency)

Processor


2 GHz, 1 word = 64 bits (8 bytes)

Unit of data access = 1 word

Cycle time = 0.5 ns (processor cycle)

2-issue superscalar, single cycle processor

2 Double Precision multiply-add (2 multipliers, 2 adders) FPU's



Peak performance of the processor = $2 \times 4 = 8$ GFlops/s
(Raw compute power)

Clock rate Total # of FP ops / cycle

Peak performance can also be computed as $2 \times 2 = 4$ GFlops/s

Clock rate 2 pipelines



Example memory system performance (2)

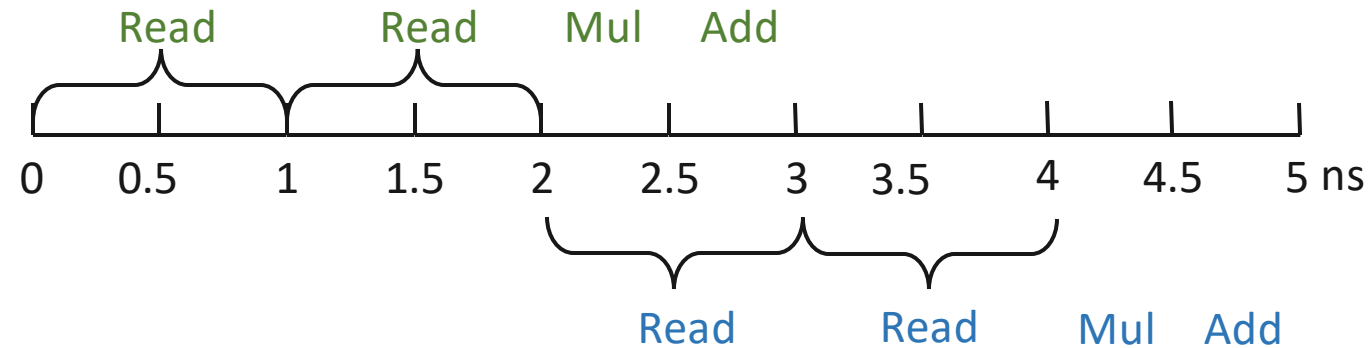
Best Case ?

Example: Inner product: $a \cdot b = \sum_{i=1}^n a_i b_i$ (data in external memory)

2 data fetches for each multiply, add

Processor can do 8 mult or add/ns = 4 FP ops /cycle

Note: issue bandwidth limits performance to 2 FP ops/cycle (if no Fused mult add)



If we can stream the data (ie no DRAM latency (10ns) per access),

In the **best case**, i.e. , data is streamed:

Sustained performance (possible best case) = 2 ops over 4 processor cycles = 1 GFlop/s


Note: Initial access latency can be ignored for large n

Processor memory bandwidth determines possible best case sustained performance

Note: Processor idles, waiting for data.

Example memory system performance (3)

Worst Case ?



Memory Latency = 10 ns
(Peak) Bandwidth = 64
bits at 1 GHz
(64 Gbits/sec or 8 GB/sec)

Repeat

- 1. read $a_i \leftarrow 10 \text{ ns}$
- 2. read $b_i \leftarrow 10 \text{ ns}$
- 3. multiply (use register data)
- 4. add (use register data)

n times

Overlap Instr. 3, 4 with Instr. 1, 2 (Pipelining)

Sustained performance (possible worst case)


$$\frac{\text{Total \# of FP ops}}{\text{Total time}} \approx \frac{2n}{20n \times 10^{-9}} = \mathbf{0.1 \text{ GFlops/s}}$$

← Memory access time



Example memory system performance (4)

Inner product



Processor Peak = 8 GFlops/s (Raw compute power)

Processor Peak = 4 GFlops/s (Processor organization)

(Best case)

Sustained = 1 GFlop/s (Memory b/w bottleneck)

(Worst case)

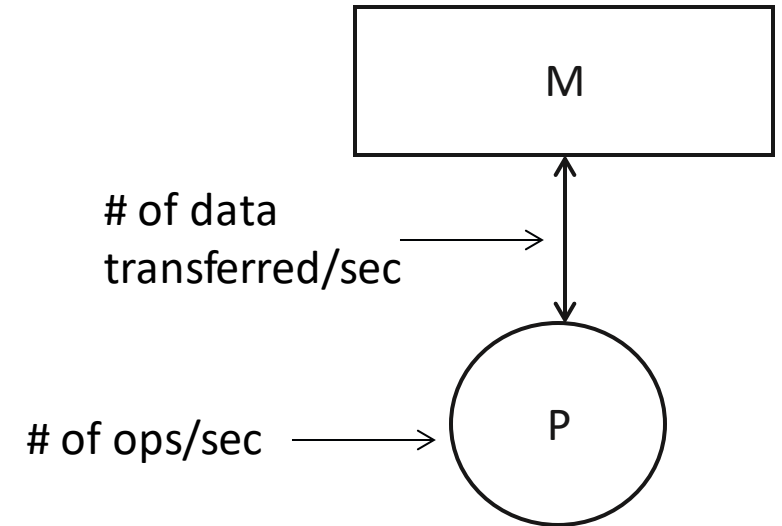
Sustained = 0.1 GFlops/s (Memory latency bottleneck)

Memory Performance dictates overall application performance

Compute Bound/Memory (bandwidth) Bound (1)




- Memory gap/Memory wall
- **Processor Performance \gg Memory Performance**
- Processor may idle waiting for data



- Compute intensity = For every fetched data (from memory) how many ops make use of them

Compute Bound/Memory (bandwidth) Bound (2)

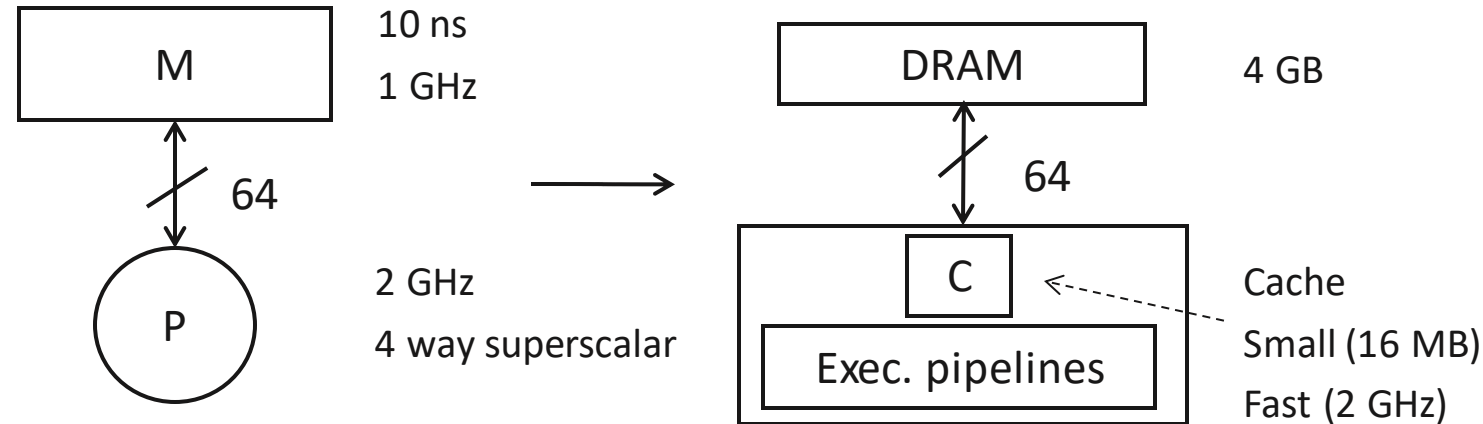


- **Compute Bound:** Performance is limited by available compute resources i.e. if we add more compute resources, the performance may improve (for a given bandwidth)



Cache (1)

Improving effective memory access latency



- Why it works?

- Data reuse


- Accessing from DRAM is expensive (DRAM latency)
 - Repeatedly use the data in cache, if possible (fast access)

- Hit ratio: fraction of the memory references served by the cache



Cache (2)

- Cache characteristics
 - Cache size (M), Cache line size (L)
 - Unit of DRAM access—line size
 - Cache organization (virtual to physical mapping)
 - Direct Mapped
 - Set Associative
 - Fully Associative
 - Cache replacement policy (what is stored in cache?)
 - First in first out
 - Most recently used
 - Random replacement
 - ...
 - Write policy
 - Write through (to cache and to DRAM)
 - Write back






Cache (3)

Locality of references (program behavior)


- **Spatial locality**

- If location i is referenced at time t , then locations near i are referenced in a small window of time following t
- Data in a cache **line** is effectively used



- **Temporal locality**

- In a small window of time, repeated references to a small set of data items
- Reduces traffic to main memory





Cache (4)

Example: bubble sort

```
for  $i$  from 1 to  $N$ 
  for  $j$  from 0 to  $N-2$ 
     $R_1 \leftarrow a[j]$ 
     $R_2 \leftarrow a[j+1]$ 
    if  $R_1 > R_2$ 
      swap ( $R_1$  and  $R_2$ )
    end if
     $a[j] \leftarrow R_1$ 
     $a[j+1] \leftarrow R_2$ 
  end for
end for
```

Assumptions:

- Direct mapped
- First in first out policy
- Cache line size = 1 data element
- Cache size 2 data elements

$i = 1$ Read $a[0], a[1] \rightarrow$ miss, miss
 Read $a[1], a[2] \rightarrow$ hit, miss
 Read $a[2], a[3] \rightarrow$ hit, miss
 ...
 Read $a[6], a[7] \rightarrow$ hit, miss

$i = 2$ Read $a[0], a[1] \rightarrow$ miss, miss
 Read $a[1], a[2] \rightarrow$ hit, miss
 ...

Cache hit ratio = $\frac{\text{\# of times data found in cache}}{\text{Total \# of accesses}}$

$i = 1$, Read hit ratio = $\frac{N-2}{(N-1)*2} = \frac{6}{7*2} = 0.43$




Impact of cache on performance (1)

- Matrix multiplication with on-chip cache

$$\mathbf{C} = (\mathbf{A} \times \mathbf{B})_{32 \times 32}$$


- 4 GHz Processor + DRAM
- 2 issue/cycle
- 5 stage pipeline
- Cache access latency = processor cycle time (0.5 ns)
- Multiple access to cache





Impact of cache on performance (2)

- Steps:
 - Read **A** into cache
 - Read **B** into cache
 - Perform n^3 multiply and n^3 add to compute **C**
 - Store the results in cache
 - Store the results from cache to DRAM
 - Cache size $\geq 3K$ words for 32×32 matrix multiplication





Impact of cache on performance (3)

- Assume
 - No cache misses (no conflicts) – doable by intelligent data placement
 - Ideal case (no data dependencies) – doable by program restructuring
 - There is some overhead (load, compute) – 2 instructions / compute step
- Case 1: 10 ns latency / memory access
 - Case 1a: with cache
 - Case 1b: without cache
- Case 2: 1 ns latency / memory access (streaming)
 - Case 2a
 - Case 2b



Impact of cache on performance (4)

- Case 1a: **worst case access latency to memory with cache**

- Read **A** into cache: $(32 \times 32) \times 10 \text{ ns} = 10 = 10 \mu\text{s}$
- Read **B** into cache: $10 \mu\text{s}$
- Perform $2n^3$ mult/add to compute **C** and store in cache:
 $2 \times 0.25 \text{ ns} \times 2 \times 32^3 / 2 = 16.384 \mu\text{s}$
- Store the results from cache to DRAM: $10 \mu\text{s}$

Load
Compute

- Total time: $46.384 \mu\text{s}$
- Total number of ops: $2n^3 = 2(32)^3 = 65,536$


Sustained performance: $\frac{65536}{46.384 \times 10^{-6}} = 1.41 \text{ GFlops/s}$

Note: we used worst case access latency to external memory



Impact of cache on performance (5)

- Case 1b: **worst case access latency to memory** **without** cache
- Sustained performance
 - Latency of 10 ns per access to each data
 - 2 words lead to one Multiply and one Add
 - Overlap read with arithmetic operations
 - = $2/20$ GFlops/s
 - = 0.1 GFlops/s





Impact of cache on performance (6)

- Case 2a: **best case access latency to memory with cache**

- Total time = time to compute + data access
= $16.384 + 1 + 1 = 18.384 \mu s$ (time to load matrix A and B = $1 \mu s$)

- Sustained Performance = $\frac{65536}{18.384 \times 10^{-6}} = 3.56 \text{ GFLOPS/s}$

- Case 2b: **best case access latency to memory without cache**

- Two operations for two data access, streaming access at 1 ns per data

- Sustained Performance = $\frac{2}{2 \times 1 \times 10^{-9}} = 1 \text{ GFLOPS/s}$




Impact of cache on performance (7)

- Performance improvement due to cache
 - Sustained performance (with cache)/Sustained performance (without cache)
 - Case 1: **$1.41/0.1 \approx 14x$**
 - Case 2: **$3.56/1 \approx 3.5x$**
- Note
 - Parallel algorithm should be carefully designed to achieve such speed-ups by exploiting cache
 - Idea can be extended for large scale matrix multiplication (matrix size > cache size)



Cache pollution and sustained performance (1)

- Cache pollution
 - **All** fetched data in a cache line is not needed/used by the program
 - Reason: Mismatch between the way data is stored in DRAM and access pattern of the program (Low Spatial Locality)






Data layout and data access pattern (1)


Storing a 2-dimensional array in memory

Row major order



$$A(i, j) \rightarrow \text{Memory}(i \cdot n + j)$$
$$0 \leq i, j < n$$

Column major order



$$A(i, j) \rightarrow \text{Memory}(i + n \cdot j)$$
$$0 \leq i, j < n$$



Data layout and data access pattern (2)

- Example


- Cache line 4 words
- Sum each row of a matrix

$$B(i) = \sum_{j=0}^{n-1} A(i, j) \quad 0 \leq i < n$$

- Assume A is stored in row major order
- Implementation 1:

1. $B(i) \leftarrow 0, 0 \leq i < n$
2. for $j = 0$ to $n - 1$ do
3. for $i = 0$ to $n - 1$ do
4. $B(i) \leftarrow B(i) + A(i, j)$

Column major access pattern



$$A(i, j) \rightarrow \text{Memory}(i \cdot n + j) \\ 0 \leq i, j < n$$




Data layout and data access pattern (3)

- Example (cont.)
 - Poor spatial locality
 - cache pollution
 - memory cache line access: only one word is used



Access pattern of
the program



$$A(i, j) \rightarrow \text{Memory}(i \cdot n + j)$$
$$0 \leq i, j < n$$



Data layout and data access pattern (4)



- Example –alternate implementation

- Assume A is stored in row major order, cache line = 4 words

- Implementation 2:

1. $B(i) \leftarrow 0, 0 \leq i < n$
2. for $i = 0$ to $n - 1$ do
3. for $j = 0$ to $n - 1$ do
4. $B(i) \leftarrow B(i) + A(i, j)$

**Row major
access pattern**



- High spatial locality

- All 4 words of data accessed in each cycle are needed
 - Up to 4 times the performance as Implementation 1

$$A(i, j) \rightarrow \text{Memory}(i \cdot n + j) \\ 0 \leq i, j < n$$

Data layout and data access pattern (5)



Note: can also do column major layout and use Implementation 1
→ high performance



Memory Systems Performance Summary

- Sustained performance
 - Main memory: latency and bandwidth affect performance
 - Caches can be very effective: Improvement depends on how cache is effectively accessed/exploited (Cache friendly/Cache oblivious algorithms)
 - Hit ratio
 - Data Locality
 - spatial locality
 - temporal locality
 - Cache pollution
 - Data layout and program access pattern affect
 - Effective processor-memory bandwidth
 - Sustained performance (GFlops/sec)

STREAM benchmark for memory performance (1)



- A simple synthetic benchmark program that measures sustainable memory bandwidth (in MB/s) and the corresponding computation rate for simple vector kernels
 - designed to be more indicative of the performance of very large, vector style applications
- It counts how many bytes the user asked to be read plus how many bytes the user asked to be written

STREAM benchmark for memory performance (2)



- The table below shows how many Bytes and FLOPs are counted in each iteration of the STREAM loops.

name	kernel	bytes/iter	FLOPS/iter
COPY:	$a(i) = b(i)$	16	0
SCALE:	$a(i) = q * b(i)$	16	1
SUM:	$a(i) = b(i) + c(i)$	24	1
TRIAD:	$a(i) = b(i) + q * c(i)$	24	2




Modern CPU STREAM benchmark results

Processor	# of cores	COPY	SCALE	ADD	TRIAD
AMD EPYC 7601	64	283 GB/s	286 GB/s	292 GB/s	290 GB/s
Intel Xeon Platinum 8160	40	175 GB/s	150 GB/s	170 GB/s	175 GB/s
IBM POWER9	40	240 GB/s	250 GB/s	255 GB/s	260 GB/s

<https://www.amd.com/system/files/2017-06/AMD-EPYC-SoC-Delivers-Exceptional-Results.pdf>
<https://www.mdpi.com/2079-9292/9/6/1035/pdf>



Architectural support for latency hiding (1)




no interactions

- **Program:** Collection of *independent sequence of instructions* ➔ *thread*
- Architecture → switch between threads when there is access to main memory by a thread
 - Typically hardware supports single-cycle switch between threads



Architectural support for latency hiding (2)



- Example

- Assume 20 cycles for DRAM access [0.5 ns processor, 10 ns DRAM]
- Assume 1 cycle overhead for context switch
- Single thread: can be blocked for 20 cycles
- Multiple threads (e.g., $p \geq 10$): we can make sure the processor is kept busy executing some thread (i.e. doing useful work)

Multi-threading



Prefetching to hide memory latency

Example

1. for $i = 0$ to $n - 1$ do
2. $C(i) \leftarrow A(i) + B(i)$

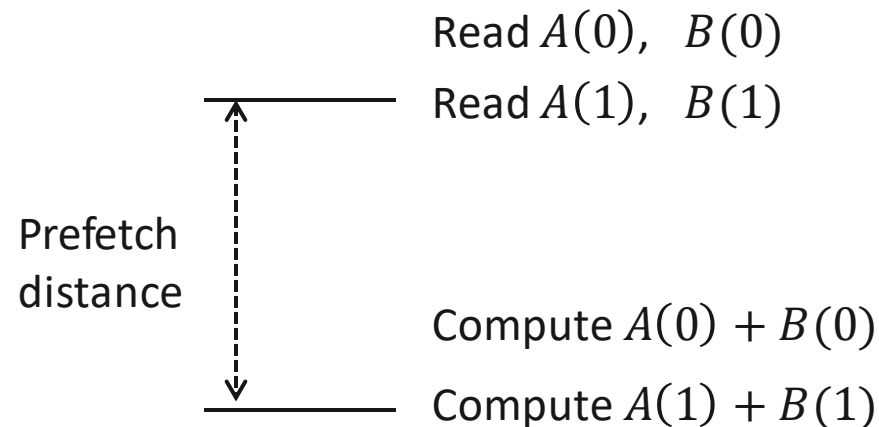
$A(0), B(0) \rightarrow$ **cache miss**

Issue request to read $A(0), B(0)$ in advance

Issue request to read $A(1), B(1)$ in advance

to mask 20-cycle memory latency

... } inserted by compiler





Summary

- Memory systems
 - Main memory latency, throughput
 - Cache
 - Data reuse
 - To achieve high performance (reduce execution time)
 - Latency hiding
 - Reduce processor – memory traffic (increase data reuse)



Backup Slides



Compute intensity example

$A \leftarrow B * C$ $n \times n$ matrix multiplication

Do $i = 1$ to n

 Read i^{th} row of B

 Do $j = 1$ to n

 Read j^{th} col of C

 Inner product $B(i,*)$, $C(*,j)$

Compute intensity of row of $B = n$

Compute intensity of col of $C = 1$