



# EE/CSCI 451: Parallel and Distributed Computation

Lecture #4

8/27/2020

Viktor Prasanna

prasanna@usc.edu

[ceng.usc.edu/~prasanna](http://ceng.usc.edu/~prasanna)

University of Southern California



# Policies and Procedures

- Specify your current time zone and your background info.
  - Please fill the Google form: <https://forms.gle/yqxoUa7NC7bMq5yv6>



# Grades (anxiety?)

## Grade Distribution: Fall 2019


- Midterm 1:
  - A & A- (70-100 pts) : 2 (4%)
  - B & B- (40-70 pts): 33 (73%)
  - C (30-40 pts): 9 (20%)
  - D (0-30pts): 1 (3%)
- Midterm 2:
  - A & A- (80-100 pts) : 5 (11%)
  - B & B- (50-80 pts): 34 (76%)
  - C (0-50 pts): 6 (13%)
- Final grade:
  - A & A- : 33 (73%)
  - B & B- : 11 (24%)
  - C : 1 (3%)



# Outline

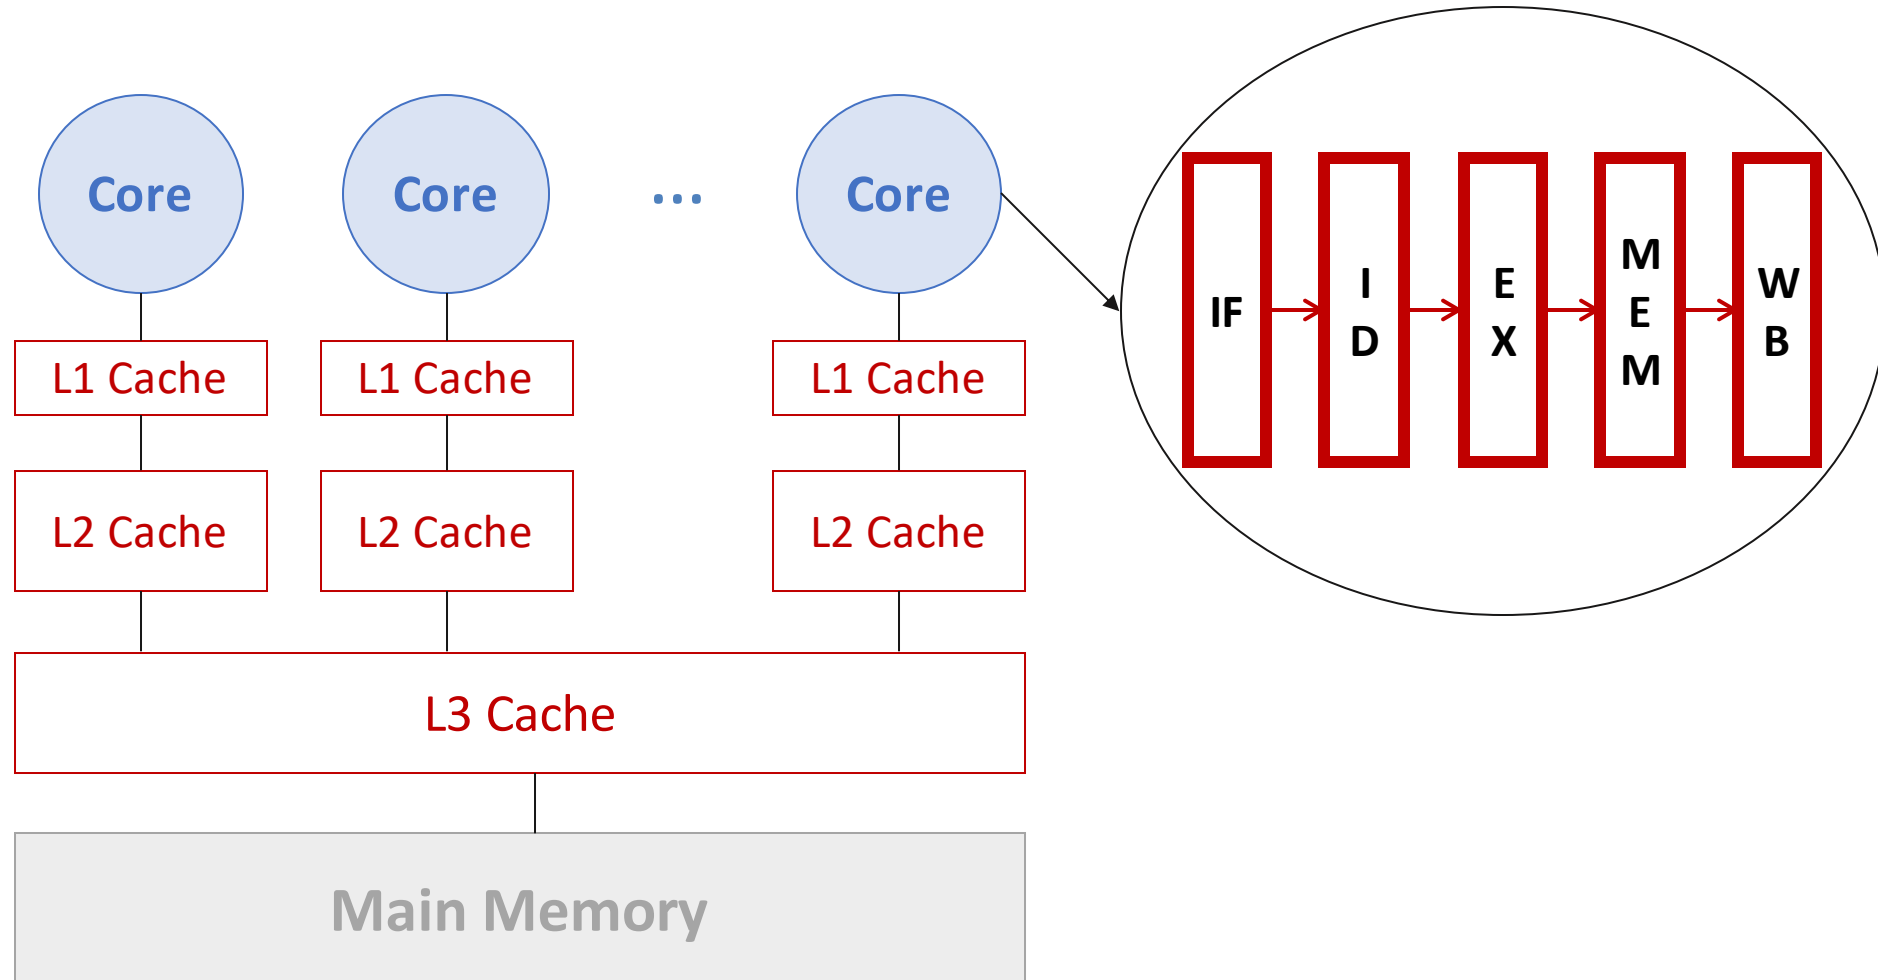
- From last class
  - Memory system
    - Latency, Bandwidth: performance implications
    - Cache: impact on performance
    - Data layout
    - Impact of memory system on performance
      - Random/streaming access
      - Cache
    - Latency hiding
- Today
  - Shared memory programming model
  - Scalability of a parallel solution
  - A simple model of shared memory parallel computation
  - Example shared memory programs

# Generic Parallel Architecture





# Node Architecture





# Parallelism

- In each node (multiple cores in parallel)
- Across the  $p$  nodes ( $p$  nodes in parallel)
- Implicit parallelism



# Distributed Shared Memory


- The address space is distributed across the local external memory of the nodes
- Access to shared data provided by the architecture
  - Load R, X
    - ↑  
May be local or in remote node
- Access to (local and remote) data is implicit





# Shared Address Space Programming (1)

- All threads have access to the same global, shared memory
- Threads can also have their own private data
- Programmer is responsible for synchronizing access (protecting) globally shared data (to ensure correctness of the program)






# Shared Address Space Programming (2)

- Thread

- An independent stream of instructions
- Can be scheduled to run concurrently and/or independently by the operating system
- Exists within a process and uses the process resources
- Lightweight: most of the overhead already been accounted for through the creation of the process
  - Example: POSIX Threads (Pthreads)





# Shared Address Space Programming (3)

- Example 1

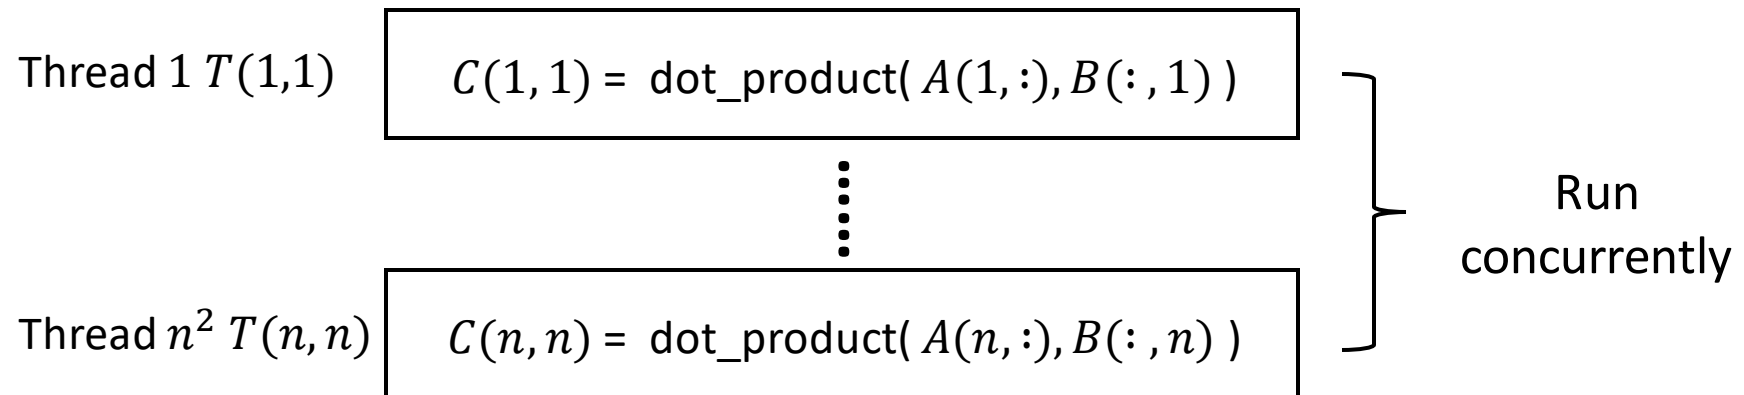
- $n \times n$  Matrix multiplication

$i$  from 1 to  $n$

$j$  from 1 to  $n$


$$C(i, j) = \text{dot\_product}(A(i, :), B(:, j))$$

- $n^2$  independent iterations  $\rightarrow n^2$  threads (no dependency among threads)






# Shared Address Space Programming (4)





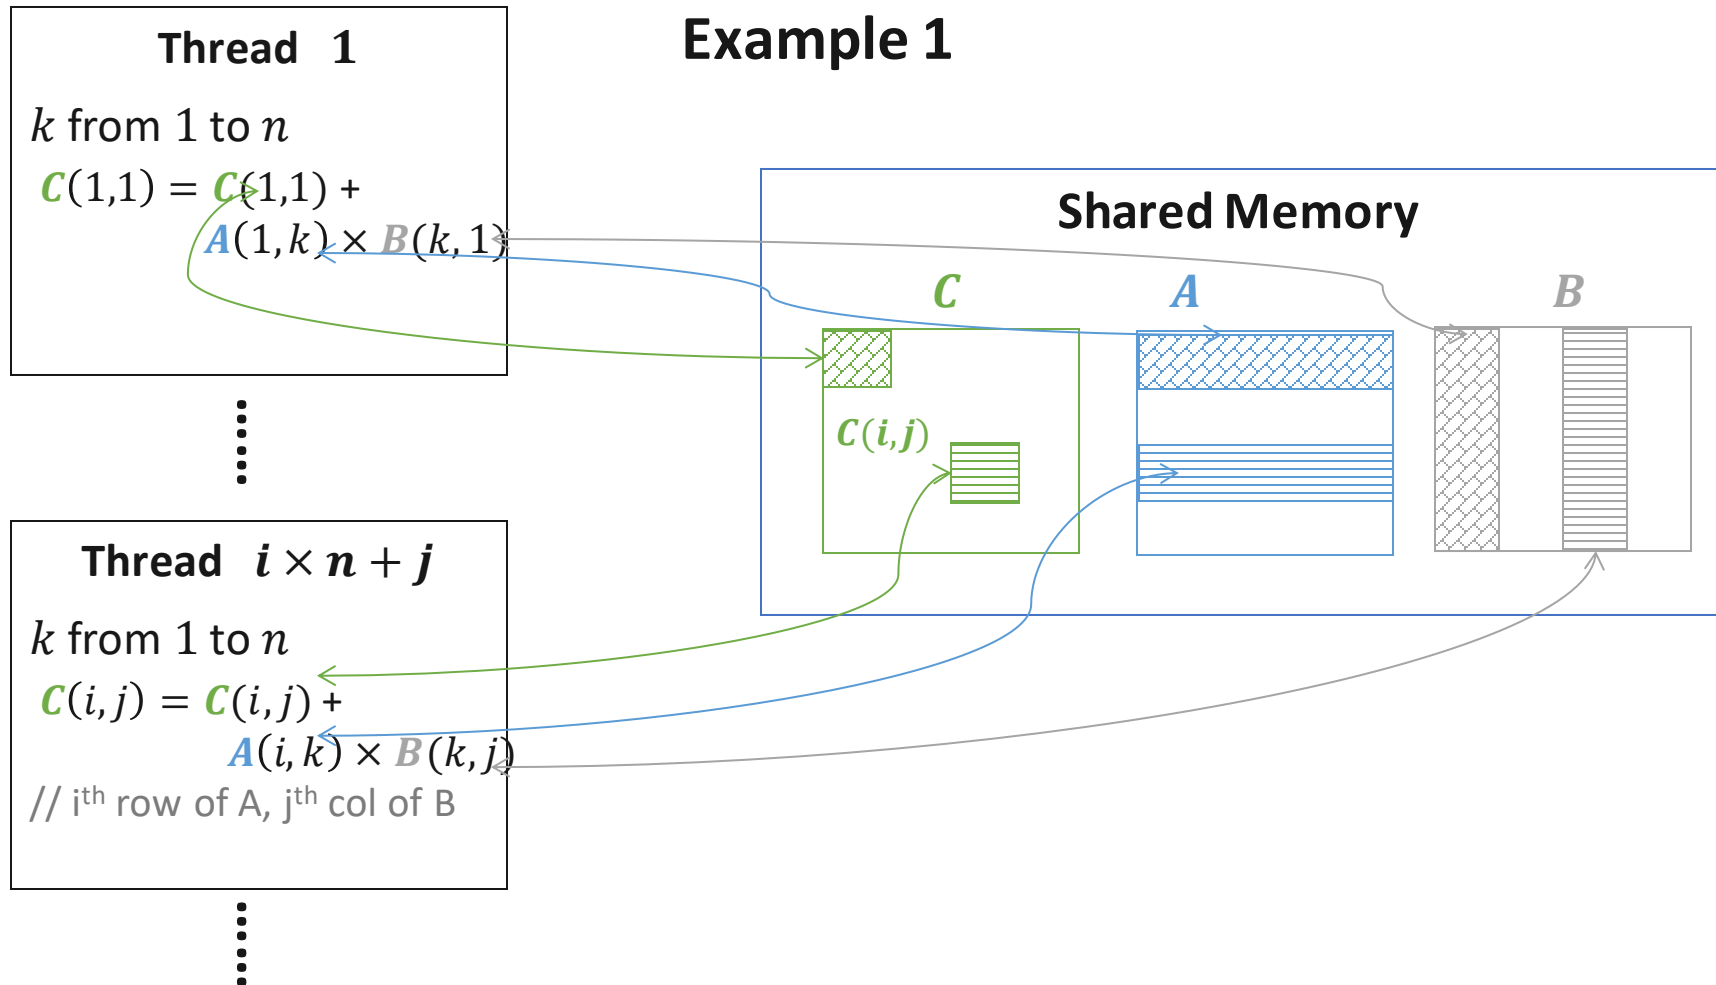
# Shared Address Space Programming (5)

- Assume  $k = n^2$
- At time  $t$ , all cores execute the same instruction ??
  - No Guarantee
- Execution of the threads is **asynchronous** even if all the cores have the same clock
  - OS scheduling
  - Cache state





# Shared Address Space Programming (6)





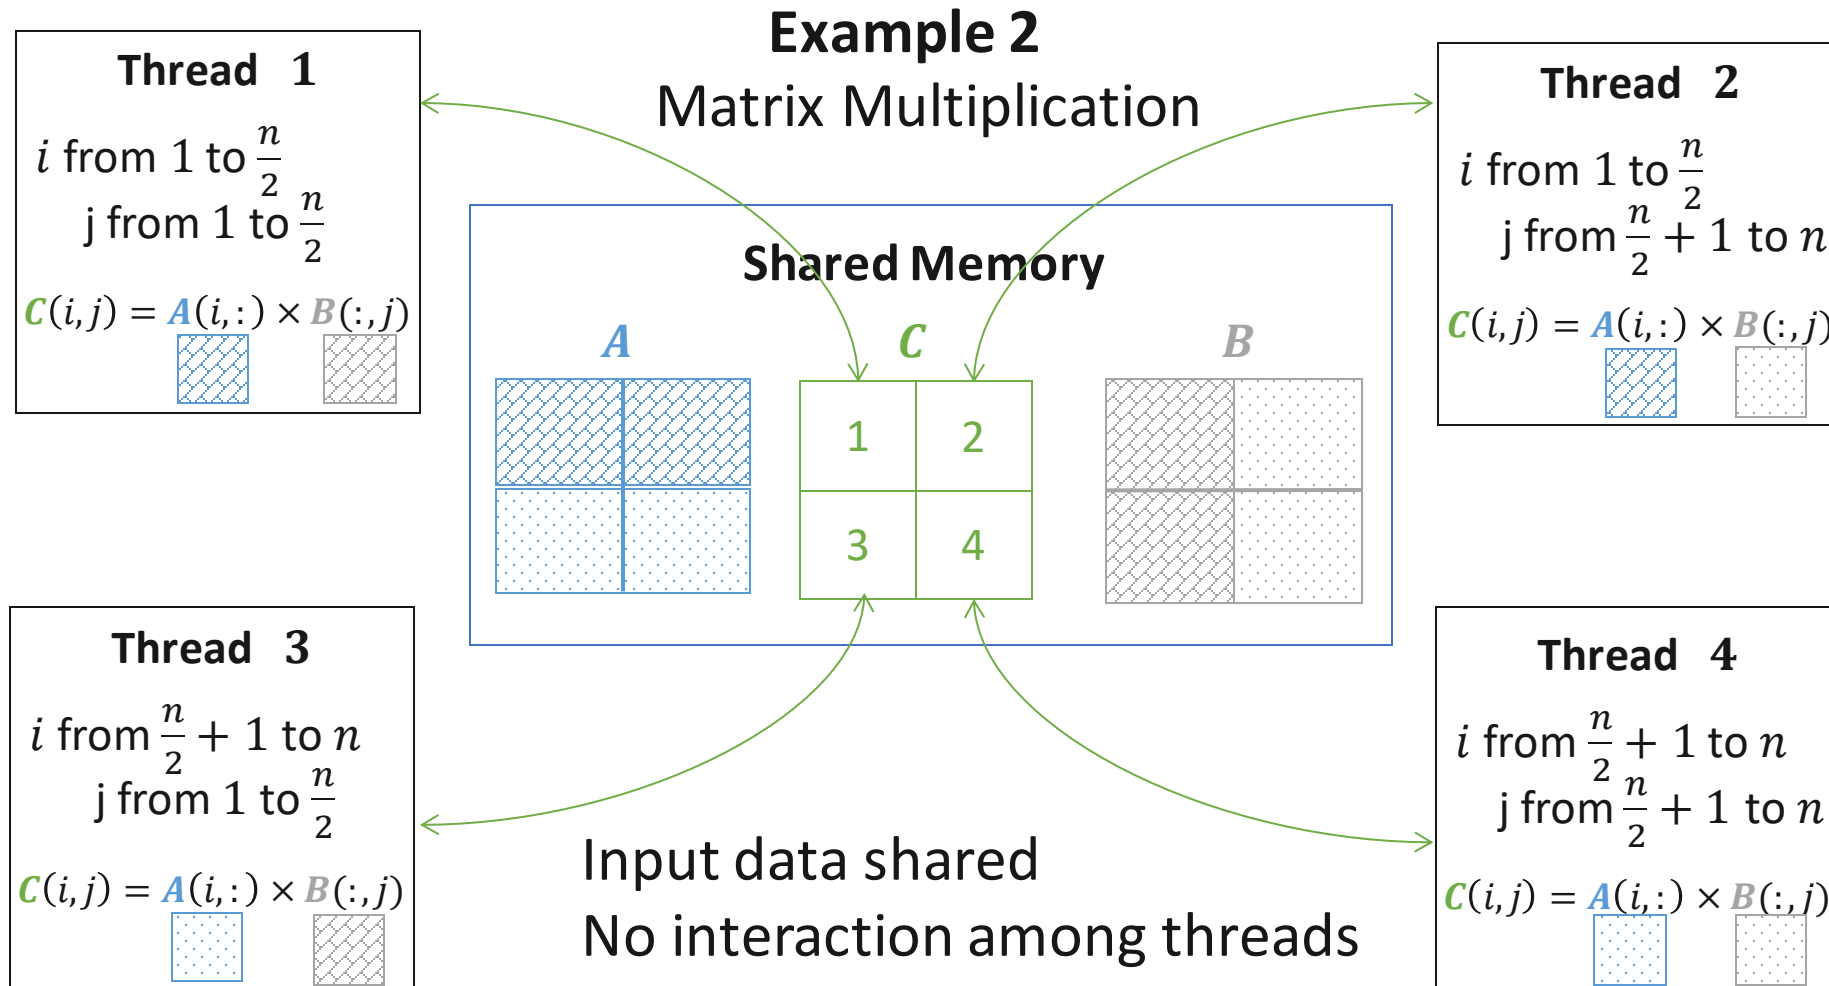
# Shared Address Space Programming (7)

## Note:

- Input data ( $A, B$  matrix) shared by the threads
- Output data ( $C$  matrix) **not** shared by threads
  - Each element of output matrix is computed (updated) by only one thread
    - => **no** dependencies between threads
    - => **no** coordination is needed between the threads
- Very simple shared address space program (Embarrassingly parallel)
- Memory access cost?
  - If  $k$  threads execute ( $k \leq n^2$ ),  $2k$  concurrent memory accesses by  $k$  threads



# Shared Address Space Programming (8)









# Shared Address Space Programming (9)

## Example 3



$T(i, j)$  computes  $(i, j)^{th}$  output block  
Each block of size  $\frac{n}{p} \times \frac{n}{p}$

```
Thread →  
  Do  $i = 1$  to  $p$   
    Do  $j = 1$  to  $p$   
       $T(i, j)$  computes  $(i, j)^{th}$  output block  
    end  
  end  
end
```



# Shared Address Space Programming (10)

Thread  $(i, j)$

Output block  $C(i, j) \leftarrow 0$

Do  $k = 1$  to  $p$

Read Block  $(i, k)$  of  $A$

Read Block  $(k, j)$  of  $B$

Perform matrix multiplication using the blocks

Update output block  $C(i, j)$

end

$$\text{Block size} = \frac{n}{p} \times \frac{n}{p}$$

$$\text{Total \# of threads} = p^2$$



# Shared Address Space Programming (11)

Thread  $(i, j)$  performs  $p \left(\frac{n}{p}\right)^3$  Add & Multiply operations  $= 2 \frac{n^3}{p^2}$  operations

Total # of ops performed by all the threads  $= p^2 \left(\frac{2n^3}{p^2}\right)$   
 $= 2n^3$  operations

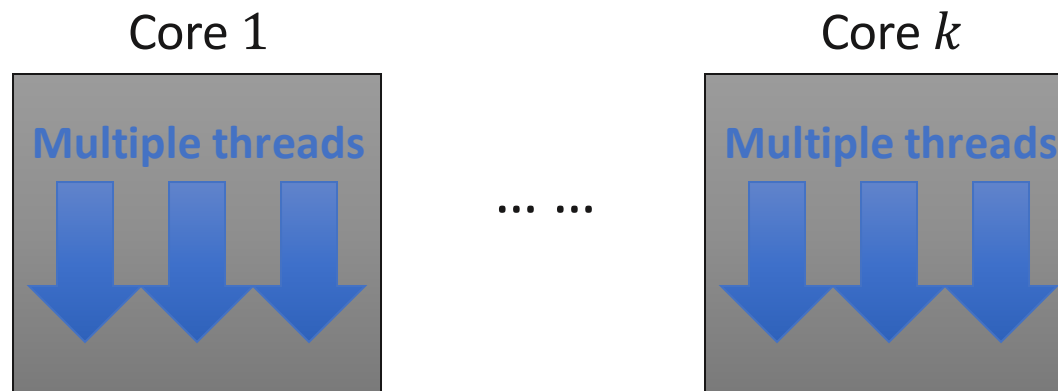
Total data fetched from shared memory  
 $= p^2 \left( p \cdot \left[ \frac{n}{p} \times \frac{n}{p} \right] \cdot 2 \right)$   
**Total # of threads**  $= 2 \overbrace{pn^2}^{\text{Block size}}$



# Threads on Multi-core platform

- Operating system scheduler maps threads/ processes to cores.
- We can bind a thread to a specific core.
- Multiple threads can run on the same core.
- More threads  $\rightarrow$  resource contention  $\uparrow$ ; thread switch overhead  $\uparrow$ ; cache hit ratio  $\downarrow$

$\rightarrow$  Can hide access latency





# Scalability (1)


- Poor Scalability?
  - Communication Cost?
  - Coordination among the processors?
- Example:  $n \times n$  matrix multiplication

$C \leftarrow A \times B$   $A, B$  are  $n \times n$  matrices

$p^2$  processors

Each processor computes a  $n/p \times n/p$  sub-matrix of  $C$


(See Example 3 earlier)






# Scalability (2)

## Algorithm:



$T(i, j)$  computes  $(i, j)^{th}$  output block  
Each block of size  $\frac{n}{p} \times \frac{n}{p}$



Thread →

```
Do i = 1 to p
  Do j = 1 to p
    T(i, j) computes (i, j)th output block
  end
end
```



# Scalability (3)

Time to compute  $(i, j)^{\text{th}}$  output block

$$= p \times 2\left(\frac{n}{p}\right)^3$$

$$= \frac{2n^3}{p^2}$$



# Scalability (4)

Total volume of data communicated



$$\begin{aligned} &= p^2 \times (n/p \times n/p) \times p \times 2 \\ &= 2p \times n^2 \end{aligned}$$

Suppose bandwidth between shared memory and the processors =  $p^2$

- $p^2$  threads access  $p^2$  values (1 value/PE) in 1 unit of time

Time for communication  $\propto n^2/p$

Time for computation  $\propto n^3/p^2$







# Scalability (5)

Total time = Computation time + Communication time

$$\begin{aligned} &= \frac{n^3}{p^2} + \frac{n^2}{p} \\ &= \frac{n^2}{p} (n/p) + \frac{n^2}{p} \quad (1) \end{aligned}$$

Suppose Total number of processors  $p^2 \leq n^2$

**Time for computation dominates time for communication**

$$\text{Parallel time} \propto \frac{\text{Serial Time}}{\# \text{ of processors}}$$

**Scalable Solution**




# Shared access and Synchronization

## 1. Shared variable access

Update shared variable  $x$

Speed of execution of threads?


Order of execution?



Produce correct output independent of execution speed of  $T_1$  and  $T_2$

## 2. Synchronization


Threads should agree they all reached an  
(agreed upon) execution state before proceeding further





# Barrier (1)

- Synchronization method
- Barrier objects can be created at certain places in the program
- Any thread which reaches the barrier stops until all the threads have reached the barrier





## Barrier (2)

Example:


Find max in array  $A[0:N-1]$ :

CreateThread(find\_max,  $A[0]$ ,  $A[N/2-1]$ , max1) {T1}

CreateThread(find\_max,  $A[N/2]$ ,  $A[N-1]$ , max2) {T2}

Barrier()

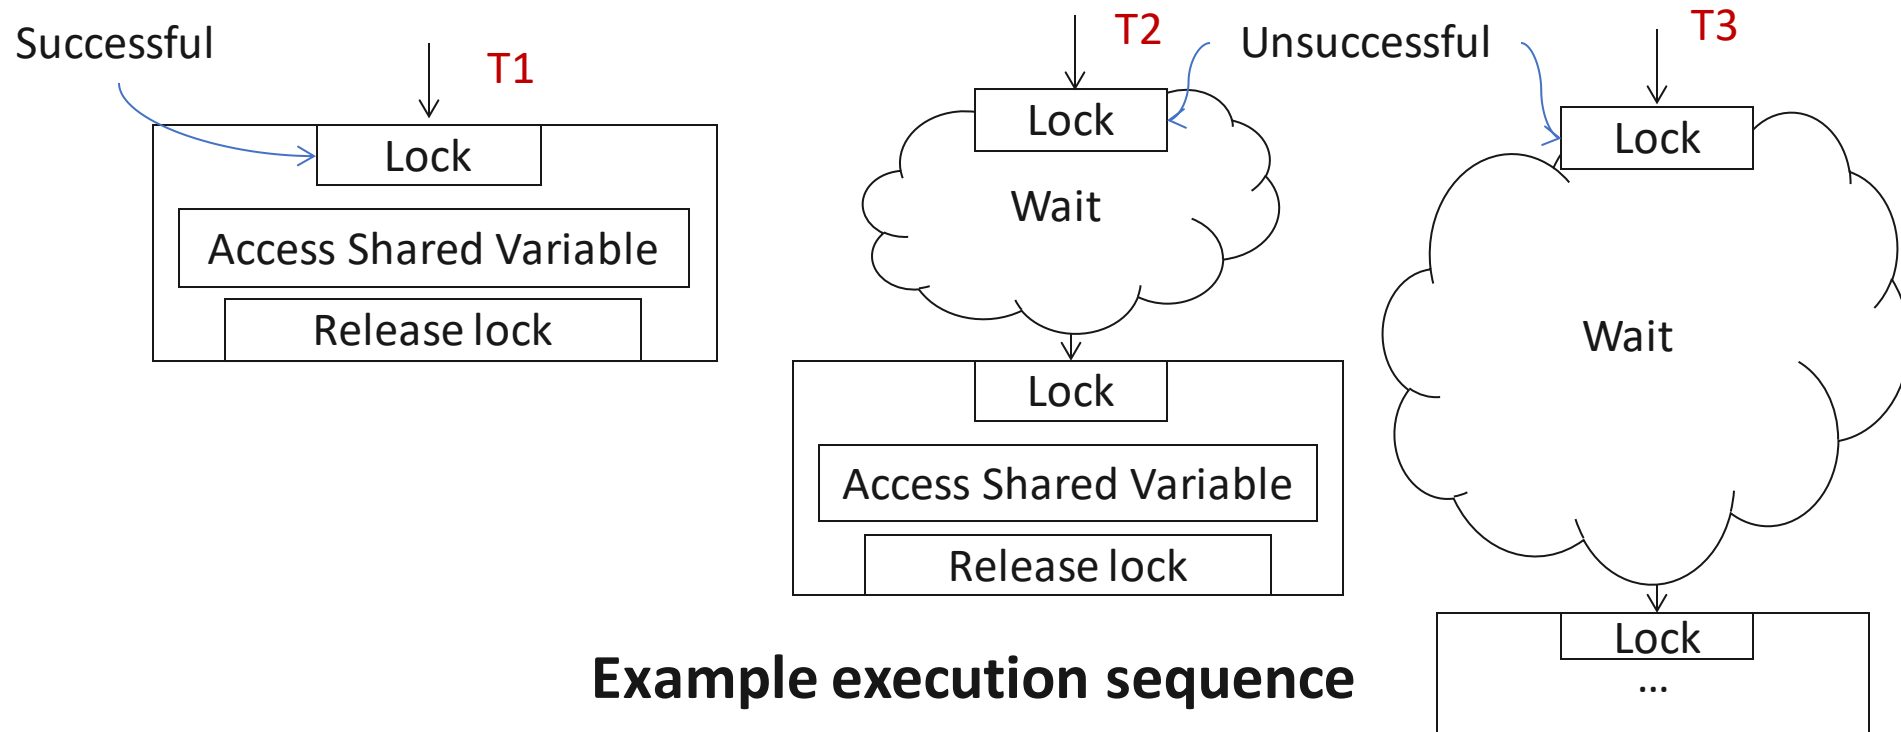
Return max(max1, max2)





# Shared Variable Access (1)

- Threads acquire lock to modify a shared variable
- Release lock when done
- Only 1 thread can acquire a lock at any time





# Shared Variable Access (2)

Example: Find max between two threads

Each thread has a local value  $i, j \geq 0$

Initialize Max (Max is a shared variable) = 0

## Thread 1

1. Acquire\_lock(Max)
2. If ( $i > \text{Max}$ )
3.  $\text{Max} = i$
4. Release\_lock(Max)

## Thread 2

- Acquire\_lock(Max)
- If ( $j > \text{Max}$ )
- $\text{Max} = j$
- Release\_lock(Max)



# Shared Variable Access (3)

Possible execution sequence depends on execution speed of threads

Max  $\leftarrow$  0

Thread 1

Thread 2

Max  $\leftarrow$  0

Thread 2

Thread 1

**In both cases correct output produced**



# Shared Variable Access (4)

**Example: Find max between two threads**

Each thread has a local value  $i, j \geq 0$

Initialize Max (Max is a shared variable) = 0

Suppose lock is **not** used

**Thread 1 ( $T_1$ )**

1. If ( $i > \text{Max}$ )

2.  $\text{Max} = i$

**Thread 2 ( $T_2$ )**

1. If ( $j > \text{Max}$ )

2.  $\text{Max} = j$

Suppose  $i = 10, j = 20$

Possible execution sequence that will produce incorrect result:

	Core 1 ( $T_1$ )	Core 2 (Faster) ( $T_2$ )
		$T_2$ Ins 1      Max = 0
Max = 0	$T_1$ ins 1	
		$T_2$ Ins 2      Max = 20
Max = 10	$T_1$ Ins 2	

Final result: Max = 10





# Shared Variable Access (5)

## Atomic Instruction

- Ex. Acquire\_lock, Release\_lock
- Single Instruction, once execution starts it is completed without interruption
- If  $i > \text{Max}$  then  $\text{Max} = i$  is **not** an atomic instruction

### Thread 1 ( $T_1$ )

1. If ( $i > \text{Max}$ )
2.  $\text{Max} = i$

### Thread 2 ( $T_2$ )

1. If ( $j > \text{Max}$ )
2.  $\text{Max} = j$

- Example code for Thread 1

1	LD R1, Max
2	CMP R1, adrs of $i$
3	IF Flag JMP Update
4	Exit
5	Update:STORE Max, R1



# Possible execution sequence

Suppose  $i = 10, j = 20$

Core 1 ( $T_1$ )

1

2

3

5

Core 2 ( $T_2$ ) Faster

1

2

3

5

Max = 10      ← incorrect



# Implementing Lock (1)

**Test\_and\_Set** Instruction  $\leftarrow$  Atomic Instruction

(Only one thread can execute this instruction at any time)

Test\_and\_Set (*lock*)

Meaning of *lock* = 0  $\rightarrow$  Shared variable associated with *lock* is **not** being used

*lock* = 1  $\rightarrow$  Shared variable associated with *lock* is being used

Execution of Test\_and\_Set (*lock*): Returns value of *lock* into a Reg (R0)  
Sets *lock* to 1

**Atomic**



# Implementing Lock (2)

Initially *lock* = 0

Acquire\_lock(*lock*)


Spin: Test\_and\_Set(*lock*)      ← Spinning (waiting)  
If R0 = 1, go to Spin      ← Atomic instructions

Release\_lock(*lock*)

Set *lock* to 0      ← Atomic instruction



# Correct Parallel Program




For **all** data inputs, for **all** execution sequences,  
correct output is produced



# Order in which instructions are executed

$T_1$	$T_2$
$I_1$	$J_1$
$I_2$	$J_2$
$I_3$	$J_3$
$I_4$	$J_4$

## Example execution sequence

$I_1$	$J_1$	
	$J_2$	
$I_2$		
$I_3$	$J_3$	
$I_4$		
	$J_4$	Time

## Execution sequence

- Can be dependent on input data
- For the same input data, can result in a different execution sequence on different parallel platforms
- For the same input data, can result in different execution sequence on the same parallel platform when executed again



# Deadlock May Occur

Thread 1


Lock( $x$ )  
    Lock( $y$ )  
        Access  $x, y$

Thread 2

Lock( $y$ )  
    Lock( $x$ )  
        Access  $x, y$



# Race Condition




Output depends on the rate at which the two parallel threads execute (for a given input)

⇒ leads to incorrect output





# Matrix Multiplication using Shared Variable (1)



Thread( $i, j$ ) accesses  $A(i, j)$  {owns  $A(i, j)$ } i. e., local data

All other data is shared


Each thread is responsible to update all outputs  $C(i, k)$ ,  $1 \leq k \leq n$ , to which  $A(i, j)$  contributes to

$$\text{Eg. } C(i, 1) = A(i, 1) * B(1, 1) + \dots \boxed{A(i, j)} * B(j, 1) + \dots A(i, n) * B(n, 1)$$

Contributes to  $C(i, k)$ ,  $1 \leq k \leq n$



## Matrix Multiplication using Shared Variable (2)





# Matrix Multiplication using Shared Variable (3)

## **Thread(*i*, *j*)**

//Uses local data  $A(i, j)$  to update **all**  $C(i, k)$

Do  $k$  from 1 to  $n$

Acquire\_lock ( $C(i, k)$ )

$C(i, k) \leftarrow C(i, k) + A(i, j) * B(j, k)$

Release\_lock ( $C(i, k)$ )

End



# Matrix Multiplication using Shared Variable (4)

## **Main Program**

Initialize  $C(i,j)$  to 0

Create  $n^2$  threads

Wait for all threads to complete

End

# Matrix Multiplication using Shared Variable (5)

## Correctness?



### Main Program

Initialize  $C(i, j)$  to 0

Create  $n^2$  threads

End

### Thread( $i, j$ )

Do  $k$  from 1 to  $n$

Acquire\_lock ( $C(i, k)$ )

$C(i, k) \leftarrow C(i, k) + A(i, j) * B(j, k)$

Release\_lock ( $C(i, k)$ )

End

### Main Program

Create  $n^2$  threads

End

### Thread( $i, j$ )

Initialize  $C(i, j)$  to 0

Do  $k$  from 1 to  $n$

Acquire\_lock ( $C(i, k)$ )

$C(i, k) \leftarrow C(i, k) + A(i, j) * B(j, k)$

Release\_lock ( $C(i, k)$ )

End

# Matrix Multiplication using Shared Variable (6)

## Correctness?



### Main Program

Create  $n^2$  threads

Wait for all threads to complete

End

### Thread( $i, j$ )

Initialize  $C(i, j)$  to 0

Barrier

Do  $k$  from 1 to  $n$

Acquire\_lock ( $C(i, k)$ )

$$C(i, k) \leftarrow C(i, k) + A(i, j) * B(j, k)$$

Release\_lock ( $C(i, k)$ )

End



# Matrix Multiplication using Shared Variable (7)

## Notes:

- Each output  $C(p, q)$  is shared by  $n$  threads
- Order in which  $C(p, q)$  is updated can vary (depends on how the threads are scheduled and execution speed of the processors)
- In principle, any  $C(p, q)$  can be updated by  $n$  threads in any of the  $n!$  ways



# Programmer (Application Developer)'s Responsibility


- Ensure shared memory access is coordinated
- Ensure there is no race condition
- Insert appropriate barrier(s)
- Ensure there is no deadlock





# Pthreads Program Structure

Initialize  
Create threads





# Asynchronous Execution

- No global clock coordinating execution
- Order of execution of instructions depends on
  - input data
  - scheduling algorithm (OS)
  - speed of the processors
  - speed of communication network



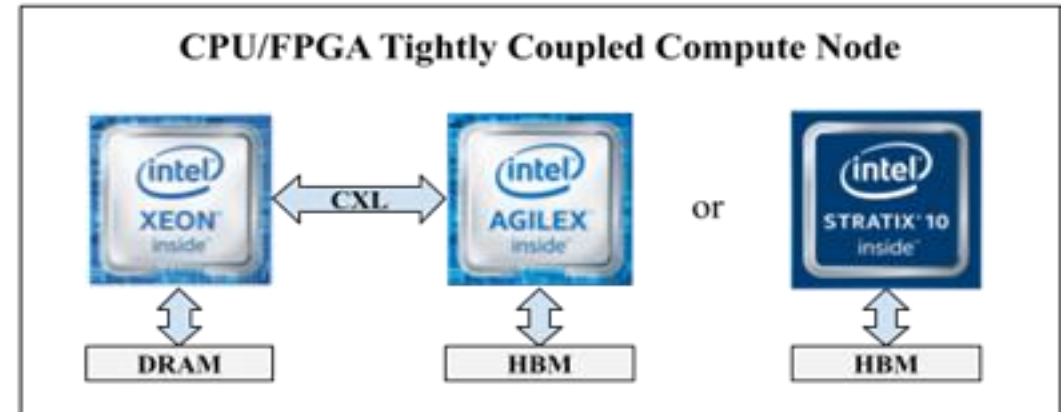
# Summary

- Shared memory programming model (Pthreads)
  - Thread
  - Thread vs. core
  - Shared variables, local variables
  - Asynchronous execution
  - Programming abstractions
    - Lock
    - Barrier
  - Coordination
  - Overheads
  - Correctness
- Examples of shared memory programs, syntax,... (Discussion session)



# Tightly coupled CPU/FPGA accelerators

- Programming model
  - One (or more) CPU thread(s) to coordinate tasks running on FPGAs.
  - One (or more) CPU thread(s) for CPU compute
  - Explicit API for synchronization and communication.
- Programming language support
  - Host:
    - OpenCL
    - OneAPI
  - FPGA:
    - HLS (Vivado HLS; OpenCL)
    - Verilog/VHDL





# CPU / FPGA threads interaction

