



EE/CSCI 451: Parallel and Distributed Computation

Lecture #5

9/1/2020

Viktor Prasanna

prasanna@usc.edu

ceng.usc.edu/~prasanna

University of Southern California



Outline

- Announcement
 - PHW2 released today, due 9/14 AOE
 - HW2 to be released on Thursday 9/3, due 9/10 AOE
- From last class
 - Shared memory programming model
 - Scalability of a parallel solution
- Today
 - A simple model of shared memory parallel computation
 - Example shared memory programs
 - OpenMP
 - OpenMP programming model
 - OpenMP directives
 - Examples

A Simple Model of Shared Address Space Parallel Machine (PRAM) (1)




1 unit of time = Local access
shared memory access

Synchronous model

Parallel time = total number of cycles

Pthreads programming model?

Asynchronous shared memory ??







PRAM (2)

Random Access Machine (RAM)


- Random Access
 - Access to **any** memory location (Direct access)
- Time
 - Access to memory \longrightarrow 1 unit of time
 - Arithmetic/Logic operation \longrightarrow 1 unit of time
- Serial time complexity $T_s(n)$
 - Example:
Merge Sort $T_s(n) = O(n \cdot \log n)$





PRAM (3)

Parallel Random Access Machine (PRAM)



Random Access Machines, executing in parallel using a common clock



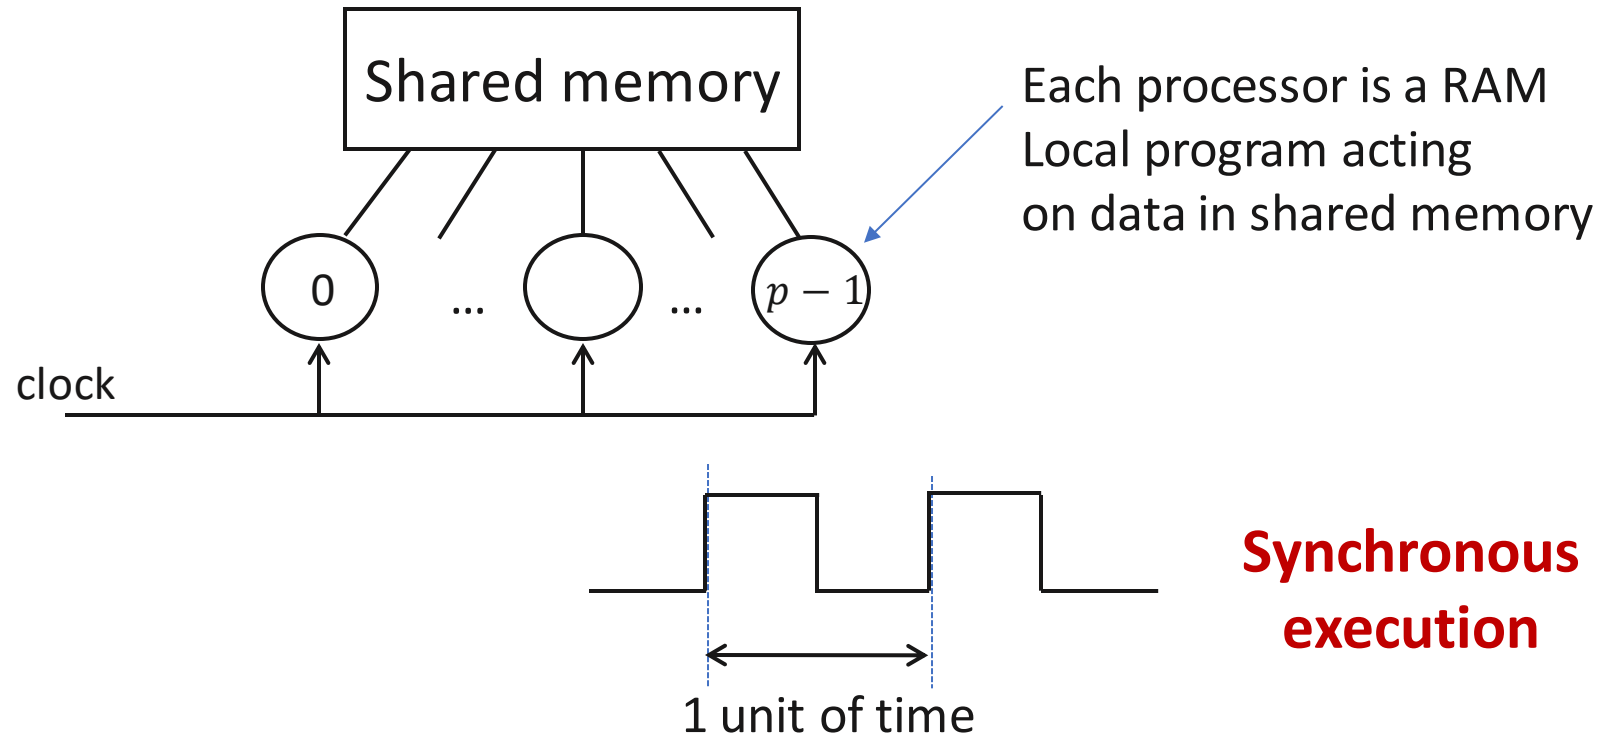
PRAM (4)

- Parallel Random Access Machine (PRAM)
 - Use shared memory for communication between processors
 - Synchronous Model (Global clock)
- Time
 - Access to **any** memory location → 1 unit of time
 - Execute one instruction in each processor (arithmetic/logic operation) → 1 unit of time



PRAM (5)

PRAM is a **synchronous** model



For all i , i -th instruction in the execution sequence is executed in the i -th cycle by **all** the processors




PRAM (6)

PRAM features

p = # of Processors

- In each cycle, p locations can be accessed (in parallel)
- All communication overheads ignored
- Ideal shared memory
- In one cycle
 - in each processor
 - Read (from shared memory)
 - Compute (using data in the processor)
 - Write (into shared memory)
- The code executed by different processors can be different
- Lock step execution
- p operations in parallel






Adding on PRAM (1)

Simple shared memory algorithm for adding n numbers

$$\text{Output} = \sum_{i=0}^{n-1} A(i) \quad \text{in } A(0)$$






Adding on PRAM (2)

Key Idea

For $n = 8$

i = iteration #





Adding on PRAM (3)

Example: $n = 8$

Active processors		iteration # (i)			(time)
Processor # (j)		0	1	2	
	0	✓	✓	✓	
	2	✓			
	4	✓	✓		
	6	✓			

$$\begin{aligned}\text{Total \# of additions performed} &= \frac{n}{2} + \frac{n}{4} + \dots + 1 \\ &= n - 1\end{aligned}$$

= Total # of additions performed by
a serial program



Adding on PRAM (4)

- Algorithm

Program in processor $j, 0 \leq j \leq n - 1$

1. Do $i = 0$ to $\log_2 n - 1$
2. If $j = k \cdot 2^{i+1}$, for some $k \in N$
 then $A(j) \leftarrow A(j) + A(j + 2^i)$
3. end


Note:

A is shared among all the processors

Synchronous operation [For ex. all the processors execute instruction 2 during the same cycle, $\log_2 n$ time]

N = set of natural numbers = $\{0, 1, \dots\}$


Parallel time = $O(\log n)$





Adding on PRAM (5)

- i^{th} iteration:
 - data which are 2^{i+1} distance apart added
 - $\frac{n}{2^{i+1}}$ partial results produced






Addition on PRAM, $p < n$ processors

1. Add within block of size $\frac{n}{p}$
2. Apply Algorithm 1

$$T_p = \text{Parallel Time} = O\left(\frac{n}{p} + \log p\right)$$

$$T_s = \text{Serial Time} = O(n)$$





Synchronous and Asynchronous

Synchronous (e.g. PRAM)

Clock

0 $y \leftarrow 0$
1-500 { Do $i = 0$ to 499
 $y \leftarrow y + A(i)$
 End
501 $y \leftarrow x + y$
502 OUTPUT y
503 STOP

0 $x \leftarrow 0$
1-500 { Do $i = 500$ to 999
 $x \leftarrow x + A(i)$
 End
501 STOP

Asynchronous (e.g. Pthreads)

$y \leftarrow 0$
Do $i = 0$ to 499
 $y \leftarrow y + A(i)$
End
Barrier
 $y \leftarrow x + y$

$x \leftarrow 0$
Do $i = 500$ to 999
 $x \leftarrow x + A(i)$
End
Barrier



Addition: Pthreads model?

Instruction execution NOT synchronized

Thread j

1. Do $i = 0$ to $\log_2 n - 1$
 2. If $j = k \cdot 2^{i+1}$, for same $k \in N$
then $A(j) \leftarrow A(j) + A(j + 2^i)$
 3. **BARRIER**
 4. end
- Number of active threads
in iteration $i = \frac{2^n}{2^{i+1}}$
- Complete each iteration before
proceeding to next iteration

Within each iteration threads may execute asynchronously

Correct output?

How to measure time?




Addition using Pthreads

Example: $n = 8$

Total # of threads = n

Active threads (threads doing arithmetic operations)





Addition using Pthreads, $p < n$ threads (1)

Each thread is allocated n/p data values

Thread j $0 \leq j < p$

Compute $B(j) \leftarrow$ sum of values allocated to it

BARRIER


i is a local variable
to each thread

Do $i = 0$ to $\log_2 p - 1$

If $j = k \cdot 2^{i+1}$, for some $k \in N$
then $B(j) \leftarrow B(j) + B(j + 2^i)$

BARRIER

end





Addition using Pthreads, $p < n$ threads (2)

- There are p threads
- All threads execute in parallel
- Correct output is produced for all inputs independent of thread execution speed

Note: Threads can execute asynchronously

Variable access latency to shared variables okay

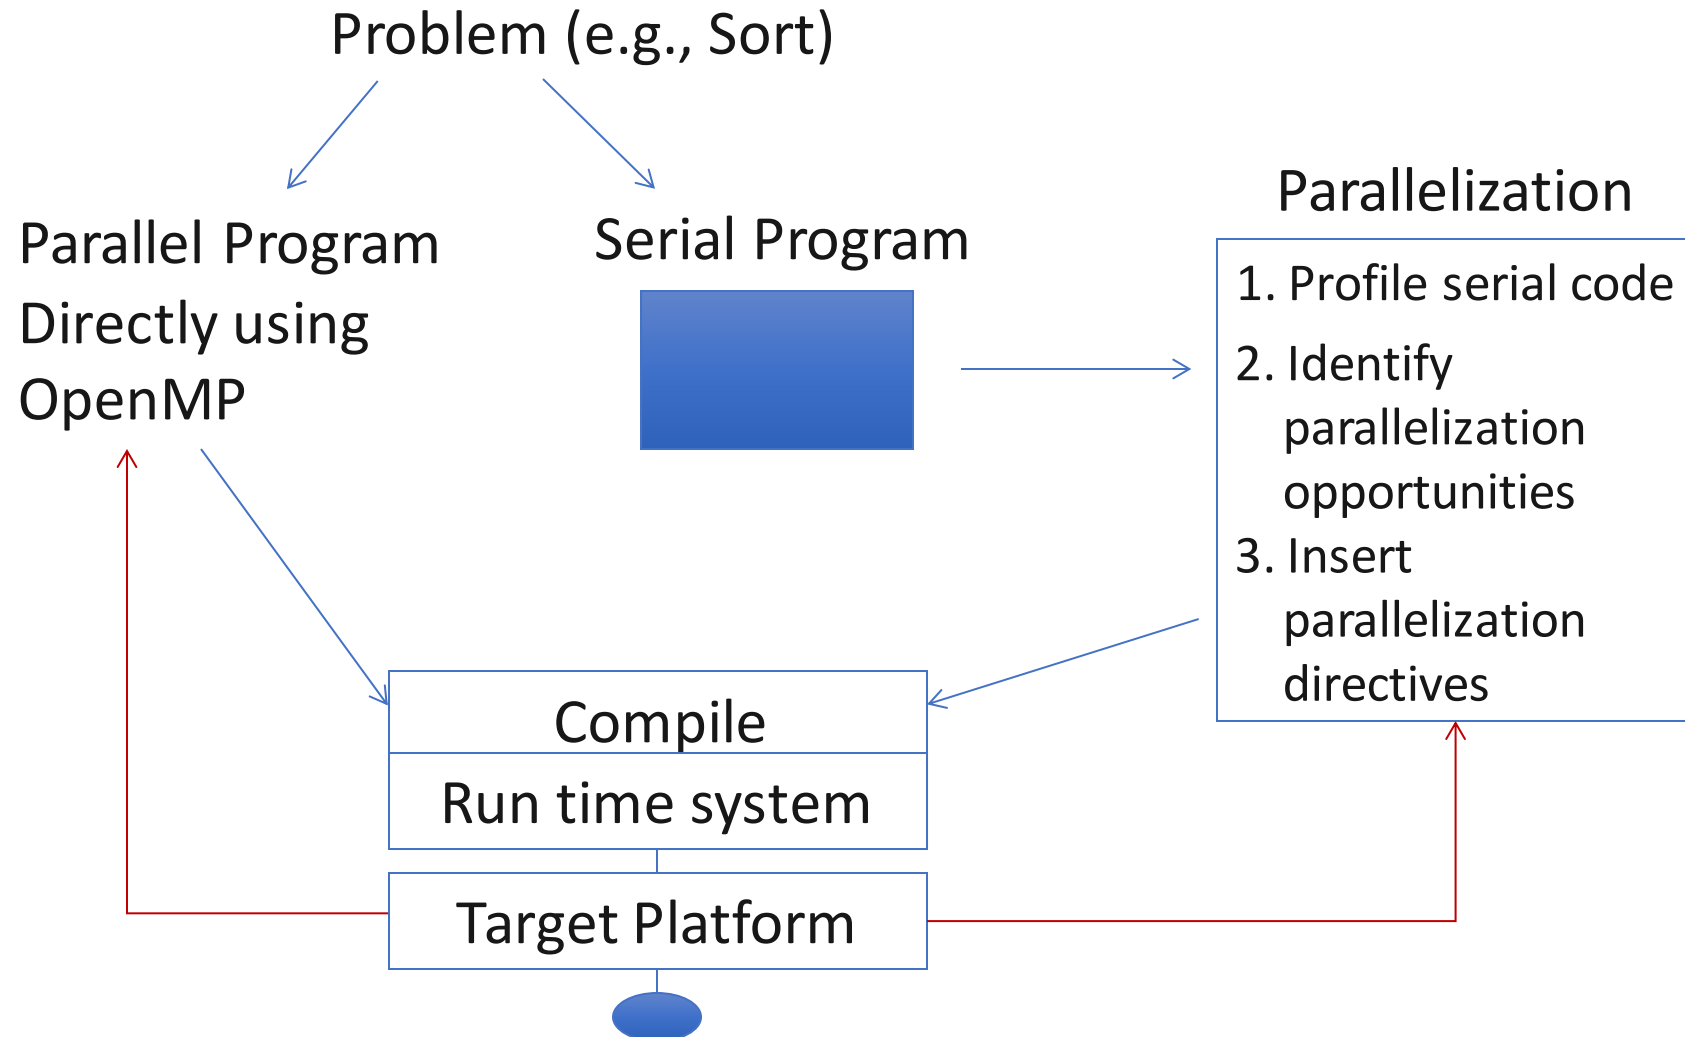


OpenMP (Open Multi-Processing)

- Application programming interface (API) for **shared memory** parallel programming
- A portable (and scalable) programming model (can be used with C, C++, Fortran and various parallel platforms)
- Outcome of standardization efforts
- Consists of compiler directives, library routines, and environment variables
- Higher level of abstraction than Pthreads

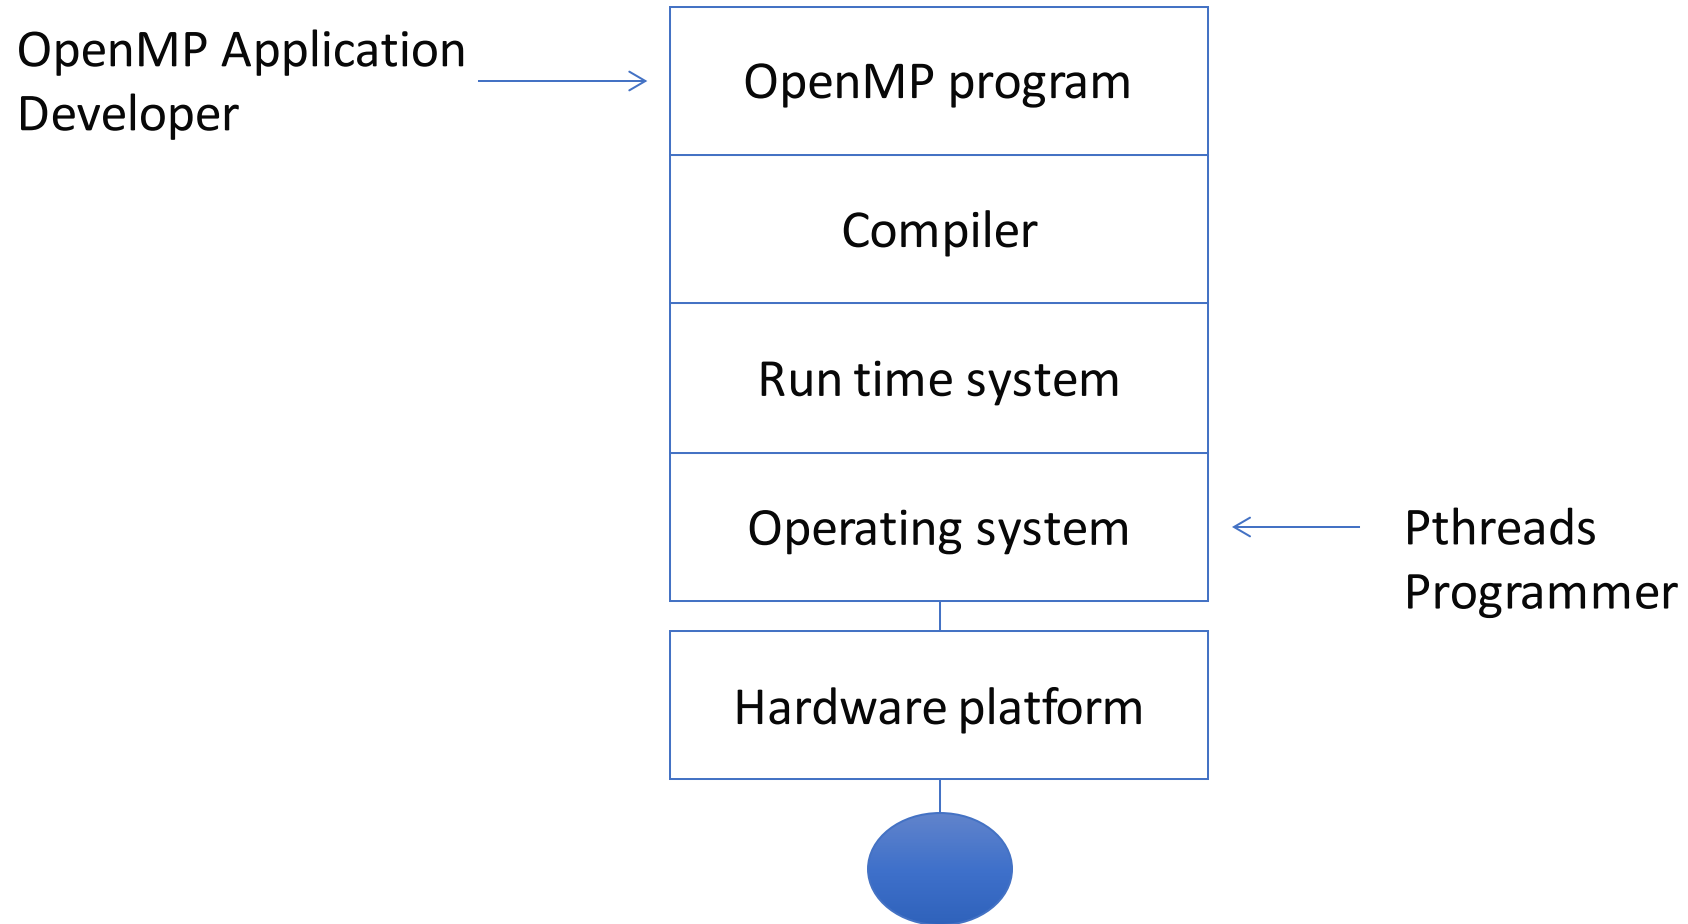


Two Uses of OpenMP





OpenMP vs. Pthreads





OpenMP Programming Model (1)

- Shared memory, threads based parallelism
- Explicit parallelism
 - Explicit (not automatic) programming model, offering the programmer full control over parallelization
 - Parallelization can be taking a serial program and inserting compiler directives
- Underlying hardware may or may not provide hardware support for shared memory



OpenMP Programming Model (2)

- Directive based parallel programming

- 1 `#pragma omp directive [clause list]`
- 2 `/* structured block */`


Structured block = a section of code which is grouped together

- OpenMP directives are not part of a programming language
 - Can be included in various programming languages (e.g., C, C++, Fortran)

OpenMP Programming Model (3)



- Fork - Join model:
 - **Fork:** The master thread creates a team of parallel threads
 - **Join:** When the team threads complete the statements in the parallel region, they synchronize and terminate, leaving only the master thread






OpenMP Directives (1)

- OpenMP executes serially until **parallel** directive
- **parallel** region construct
 - Same code will be executed by multiple threads
 - Fundamental OpenMP parallel construct
 - Example

```
#pragma omp parallel [clause list]
{
    task();
}
```



- Clause list specifies parameters of parallel execution
 - # of threads, private variables, shared variables



OpenMP Directives (2)

- Work-Sharing Constructs
 - **Divide** the execution of the enclosed code region among the members
 - Implied barrier at the end of a work sharing construct
 - Types of Work-Sharing Constructs:
 - **for** - shares iterations of a loop across the team. Represents a type of "data parallelism".
 - **sections** - breaks work into separate, discrete sections. Each section is executed by a thread. Represents a type of "functional parallelism".



OpenMP Directives (3)

Loop parallelism

- Most programs have loops
- Rich source for parallelization/optimization
- 90/10 rule
- Many scheduling strategies and compile time and run time optimizations



OpenMP Directives (4)

for directive example: $N \times N$ matrix multiplication, $C = A \times B$

Serial

```
for (i=0; i<N; i++) {  
    for (j=0; j<N; j++) {  
        c(i,j) = 0;  
        for (k=0; k<N; k++)  
            c(i, j) += a(i, k) * b(k, j);  
    }  
}
```

Static scheduling of loops

```
#pragma omp parallel num_threads (4)  
#pragma omp for schedule(static)  
for (i=0; i<N; i++) {  
    for (j=0; j<N; j++) {  
        c(i,j) = 0;  
        for (k=0; k<N; k++)  
            c(i, j) += a(i, k) * b(k, j);  
    }  
}
```

Static scheduling splits the iteration space into equal chunks of size (specified in the clause list) and assigns them to threads in a round-robin fashion




OpenMP Directives (5)

Iteration Space

3 level nested loop for matrix multiplication

dim = 128

```
for (i = 0; i < dim; i++) {  
    for (j = 0; j < dim; j++) {  
        c(i,j) = 0;  
        for (k = 0; k < dim; k++)  
            c(i, j) += a(i, k) * b(k, j);  
    }  
}
```





OpenMP Directives (5)

Iteration Space


3 level nested loop for matrix multiplication

dim = 128, 4 threads

```
#pragma omp parallel num_threads (4)
```

```
#pragma omp for schedule(static)
```

```
for (i = 0; i < dim; i++) {  
    for (j = 0; j < dim; j++) {  
        c(i,j) = 0;  
        for (k = 0; k < dim; k++)  
            c(i, j) += a(i, k) * b(k, j);  
    }  
}
```



Note: Iteration space can be partitioned in many ways ← compiler optimizations



OpenMP Directives (7)

Scheduling iteration space



- **Static:** partition into equal sized chunks, map at compile time
- **Dynamic:** partition into equal sized chunks, run time system assigns them to threads (for ex. as the threads become idle)
- **Guided:** Vary the chunk size as the (loop) computation proceeds to ensure load balance among the threads
- **Runtime:** scheduling decisions left to the run time system

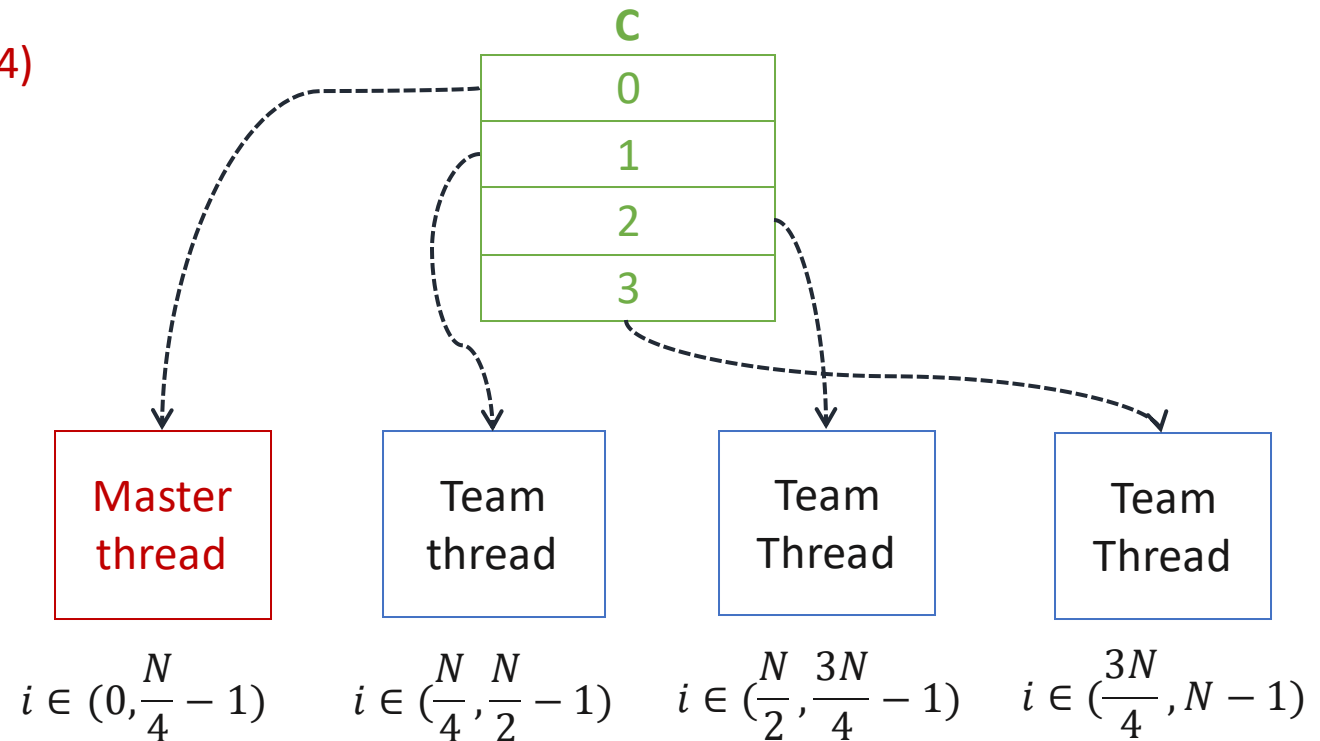


OpenMP Directives (8)

- **for** directive example (Data/Loop Parallelism):

$N \times N$ matrix multiplication: $C = A \times B$

```
#pragma omp parallel num_threads(4)
{
    #pragma omp for
    for (i=0; i< N; i++)
        Calculate_block()
}
```






OpenMP Directives (9)

- **sections** example (Task Parallelism):

$N \times N$ matrix multiplication: $C = A \times B$


```
#pragma omp parallel num_threads(4)
{
    #pragma omp section
    Caculate_block0;
    #pragma omp section
    Caculate_block1;
    #pragma omp section
    Caculate_block2;
    #pragma omp section
    Caculate_block3;
}
```





Example (1)

5-field classification



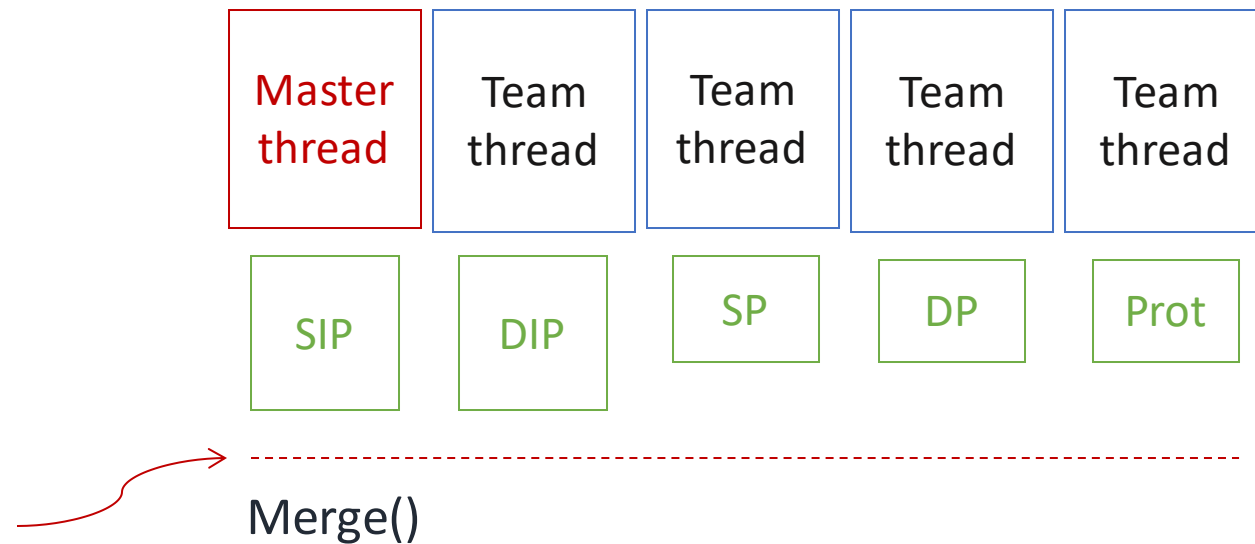
Source IP	Destination IP	Source Port	Destination Port	Protocol	Action
175.77.88.155/32	119.16.158.230/32	123 - 123	0 - 65535	TCP	Act 0
152.175.65.32/28	39.240.26.229/32	0 - 65535	750 - 760	TCP	Act 1
17.21.12.0/23	224.0.0.0/5	0 - 65535	0 - 65535	UDP	Act 2



Example (2)

```
#pragma omp parallel num_threads (5)
{
    #pragma omp section
    search_SIP();
    #pragma omp section
    search_DIP();
    #pragma omp section
    search_SP();
    #pragma omp section
    search_DP();
    #pragma omp section
    search_Prot();
    #pragma omp barrier
    #pragma omp single
    Merge();
}
```

Example: 5-field packet classification

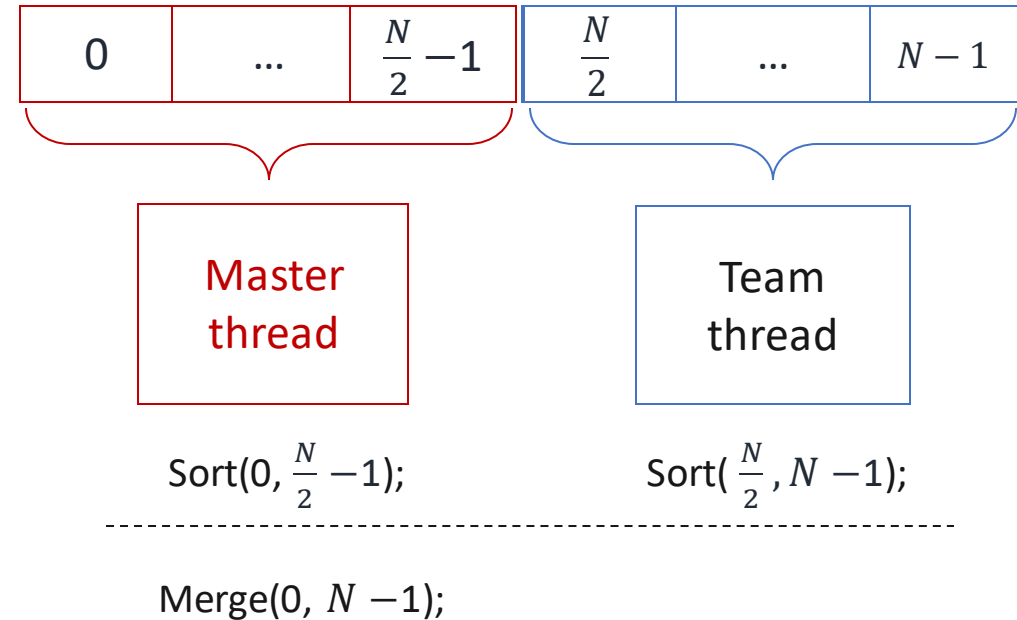




More Examples (1)

Sort an N -element Array

```
#pragma omp parallel num_threads (2)
{
    #pragma omp section
        Sort(0,  $\frac{N}{2} - 1$ );
    #pragma omp section
        Sort( $\frac{N}{2}$ ,  $N - 1$ );
    #pragma omp barrier
    #pragma omp single
        Merge(0,  $N - 1$ );
}
```





More Examples (2)

Find min. element in an N -element Array

Serial Program

```
min = A[0];  
For  $i = 1$  to  $N - 1$   
    if ( min > A[ $i$ ] )  
        min = A[ $i$ ];  
End
```



OpenMP Program

```
#pragma omp parallel num_threads(2)  
{  
    #pragma omp section  
        min1 = Min (0,  $\frac{N}{2} - 1$ );  
    #pragma omp section  
        min2 = Min (  $\frac{N}{2}$ ,  $N - 1$ );  
    #pragma omp barrier  
    #pragma omp single  
        min = Min (min1, min2);  
}
```



Summary

OpenMP

- Simple high level abstraction for shared address space programming
- Directives to the compiler
- Specify: # of threads, conditions for parallelization, local/global variables
- Serial code → incremental parallelization
- Loop parallelization, task parallelization
- Static analysis → parallel code that executes with the aid of a run time system
- Run time system: data management, layout, communication, mapping, optimizations, ...



Backup Slides



Resources

- Slides from CMU introducing multi-core architecture, thread-level parallelism, cache & cache coherence, etc.
 - <https://www.cs.cmu.edu/~fp/courses/15213-s06/lectures/27-multicore.pdf>
- A recent publication, light reading on single-core and multi-core machines in intro section, comparisons and measurements are made on two real processors
 - <https://airconline.com/ijcsit/V10N1/10118ijcsit01.pdf>
- A Professor's webpage talking about multi-core, parallelism, perf. model (very high-level)
 - <https://www.cse.wustl.edu/~jain/cse567-11/ftp/multicore/#sec1>
- A youtube video, talking about single & multi core architecture, Intel sandy bridge & AMD interlagos
 - <https://www.youtube.com/watch?v=crZwPhNjNiU>