



EE/CSCI 451: Parallel and Distributed Computation

Lecture #7

9/8/2020

Viktor Prasanna

prasanna@usc.edu

ceng.usc.edu/~prasanna

University of Southern California



Announcement

- Midterm 1 date: 9/25
 - Discussion session - 2hours
- HW2 due 9/10 (upcoming Thu.)
- PHW2 due 9/14 (+3 free late-days)
- HW1 and PHW1 grades are out

PHW1 Statistics	
Average	93.51
Median	98.00
Standard Deviation	18.86

HW1 Statistics	
Average	91.10
Median	96.00
Standard Deviation	15.39



Announcement: Midterm1 Logistics

- **Online Proctored Exam:**
 - Time: Week 6 discussion session – 2 hours: 3:30-5:30PM (Los Angeles time)
 - Format: Open-book, open-notes [**Attendance is required, no make-up given**]
 - Proctoring: 2 proctors watching different subgroups of students in separate Zoom meetings, links will be sent to students in advance
 - Require **camera-enabled** device
- Receiving and returning your exam:
 - Exam will be released on Piazza under resource page at/around 3:26 PM
 - You will submit the completed exam on Blackboard (a submission portal will be created in advance)
- Completing your exam:
 - Download the assignment pages (exam pages) as pdf files on to your tablet and annotate it with your answers. Only hand-annotated pdf files are acceptable.
 - Require a **writable tablet** device
- Coverage: **Week 1-Week 5 contents** (Week 6 contents - analytical modeling & communication primitives - not covered)
- Special note 1: **Important - discussion attendance is required on Week 5 (Sept 18)!**
 - 10-min midterm trial run to make sure all students are prepared for and comfortable with the exam process
- Special note 2:
 - We have created a Piazza poll [link] to collect info regarding your available resources/capabilities to complete the exam with writable tablet. Everyone is **required to participate in the poll**



Announcement

- Lecture slides will be released **before** each class for your reference
- Occasionally, lecture slides will be **updated** after the lecture
 - The updated slides will over ride the slides released before the class
- Lecture slides are available at
<https://piazza.com/usc/fall2020/eecsci451/resources>



Outline


- From last class
 - Message Passing Programming Model
 - Asynchronous/Loosely Synchronous/SPMD
 - Send/Receive
 - Blocking/Non-blocking
 - Cannon's Algorithm for Matrix Multiplication
- Today
 - Interconnection networks (Chap. 2.4.3-2.4.5)
 - Crossbar network
 - Shuffle exchange network
 - Multistage network




Interconnection Networks

Data communication among processors and memory

Static (direct) networks



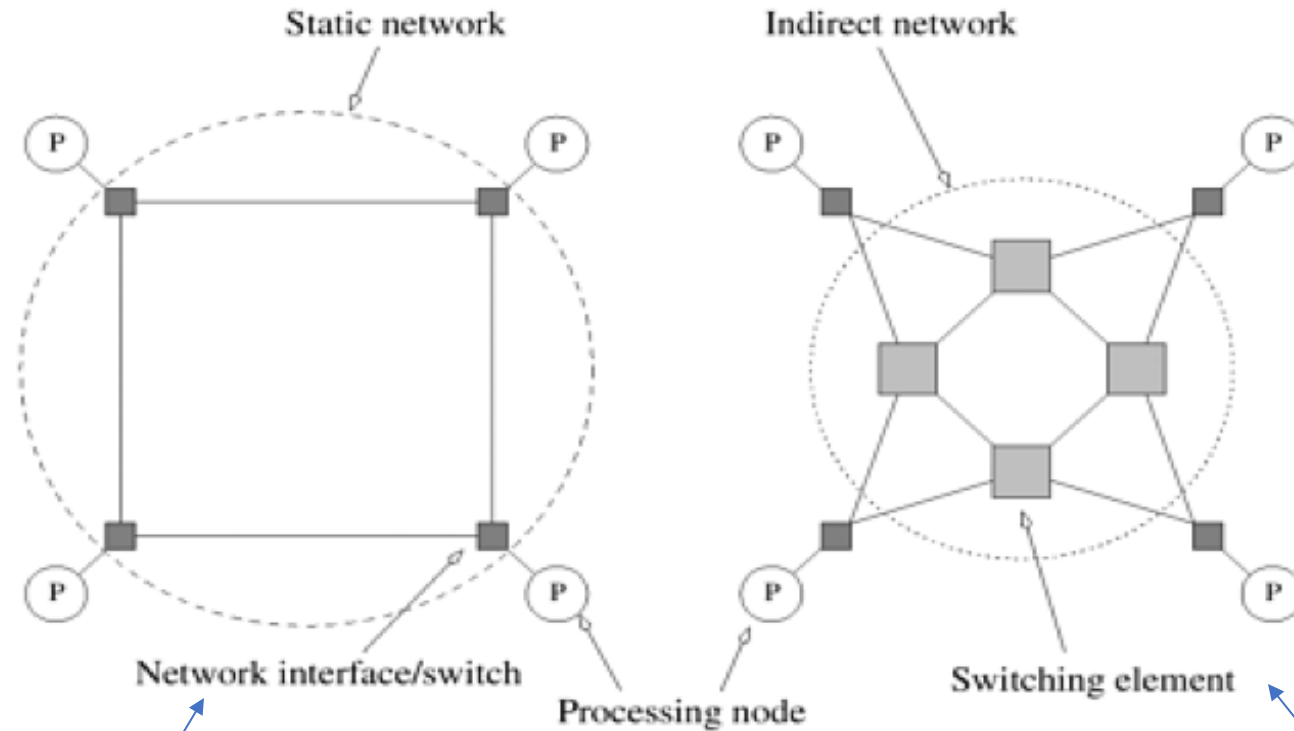
Dynamic (indirect) networks
(Multistage network)





Network Topologies

Examples:




A static network of four processing elements or nodes

A dynamic network of four nodes connected via a network of switches to other nodes



Bus-based Network




- Distance: $O(1)$
- Low cost
- Scalable?
- Bandwidth = bus width (bits) \times clock rate (independent of n)
- Traffic on the bus can be reduced by using a cache in each node

One transaction at
any time



Crossbar Network



- $p \times p$ crossbar: cost $\sim O(p^2)$ Number of switches
- To make a connection from P_i to M_j (permutation on p items):
 P_i broadcast $j \rightarrow$ switch $(i, j) \rightarrow$ close connection




Shuffle Network

Perfect Shuffle (PS) connection

- A link exists between input i and output j if:

$$j = \begin{cases} 2i, & 0 \leq i < \frac{p}{2} \\ 2i + 1 - p, & \frac{p}{2} \leq i < p \end{cases}$$

Left rotation (circular left shift) of binary representation of i
 $p = \text{power of } 2$





Shuffle Exchange Network

Perfect Shuffle (Shuffle) for $n = 8$

0	→	0
1	→	2
2	→	4
3	→	6
4	→	1
5	→	3
6	→	5
7	→	7


$$i \rightarrow 2i \bmod n \quad 0 \leq i < \frac{n}{2}$$

$$i \rightarrow (2i + 1) \bmod n \quad \frac{n}{2} \leq i < n$$

$$2i \leftrightarrow 2i + 1$$

$$0 \leq i < \frac{n}{2}$$


Exchange





Shuffle Connection

Example: $n=8$ 3 bit index



- $100 \rightarrow 001$
- Exchange connection
- Diameter:
(discussed later)

Circular left shift

$$(4 \rightarrow 1)$$

$$2i \leftrightarrow 2i + 1$$

Complement lsb

$$O(\log n) \quad n = 2^k$$



Routing in Shuffle Exchange Network (1)

Source $x = x_{k-1} \cdots x_0 \longrightarrow$ Destination $d = d_{k-1} \cdots d_0$

$y \leftarrow x$ {current location}

$i \leftarrow 1$

While $i \leq k$

 Shuffle y {Rotate left}


 Compare LSB of y with bit $(k - i)$ of destination (d)

 If bits are the same, then do not Exchange;

 else Exchange {Complement y_0 }

$i \leftarrow i + 1$

End



Total # of hops $\leq 2k$ ($2\log_2 n$)



Routing in Shuffle Exchange Network (2)

Source $x_2x_1x_0$ (000) \rightarrow Destination $d_2d_1d_0$ (110) $k = 3$

Example: $i = 1$

- Shuffle

$$000 \leftarrow 000$$

- Compare LSB of y with bit 2 of destination

$$y_0 = d_2?$$

Same as

$$x_2 = d_2?$$

Position at the end of first iteration: 001

- End of i^{th} iteration: $y = x_{k-1-i} \dots x_0 d_{k-1} \dots d_{k-i}$


$$x = 000 \quad d = 110$$

$i = 1$	000	S
	001	E
$i = 2$	010	S
	011	E
$i = 3$	110	S
	110	No E



Routing in Shuffle Exchange Network (3)

Source $x_2x_1x_0$ (000) \rightarrow Destination $d_2d_1d_0$ (110) $k = 3$






Routing in Shuffle Exchange Network (4)

- Theorem: In a shuffle exchange network with $n = 2^k$ nodes, data from **any** source to **any** destination can be routed in at most $2\log_2 n$ steps.



Multistage Network (1)

- Can realize rich set of connecting patterns from input to output



- Dynamic networks
- Multiple stages of switches and connections



Multistage Network (2)

Connecting Pattern (Data Communication Pattern)

n inputs / n outputs


For each i , $0 \leq i < n$, connecting pattern specifies

output(s) j to which data from i is to be routed to

- Example:
 - Connecting Pattern is a permutation
 - Given n inputs and n outputs,

total # of connecting patterns = $n!$

Example







Multistage Network (3)

Building Blocks

$2 \rightarrow 1$ MUX



2×2 switch

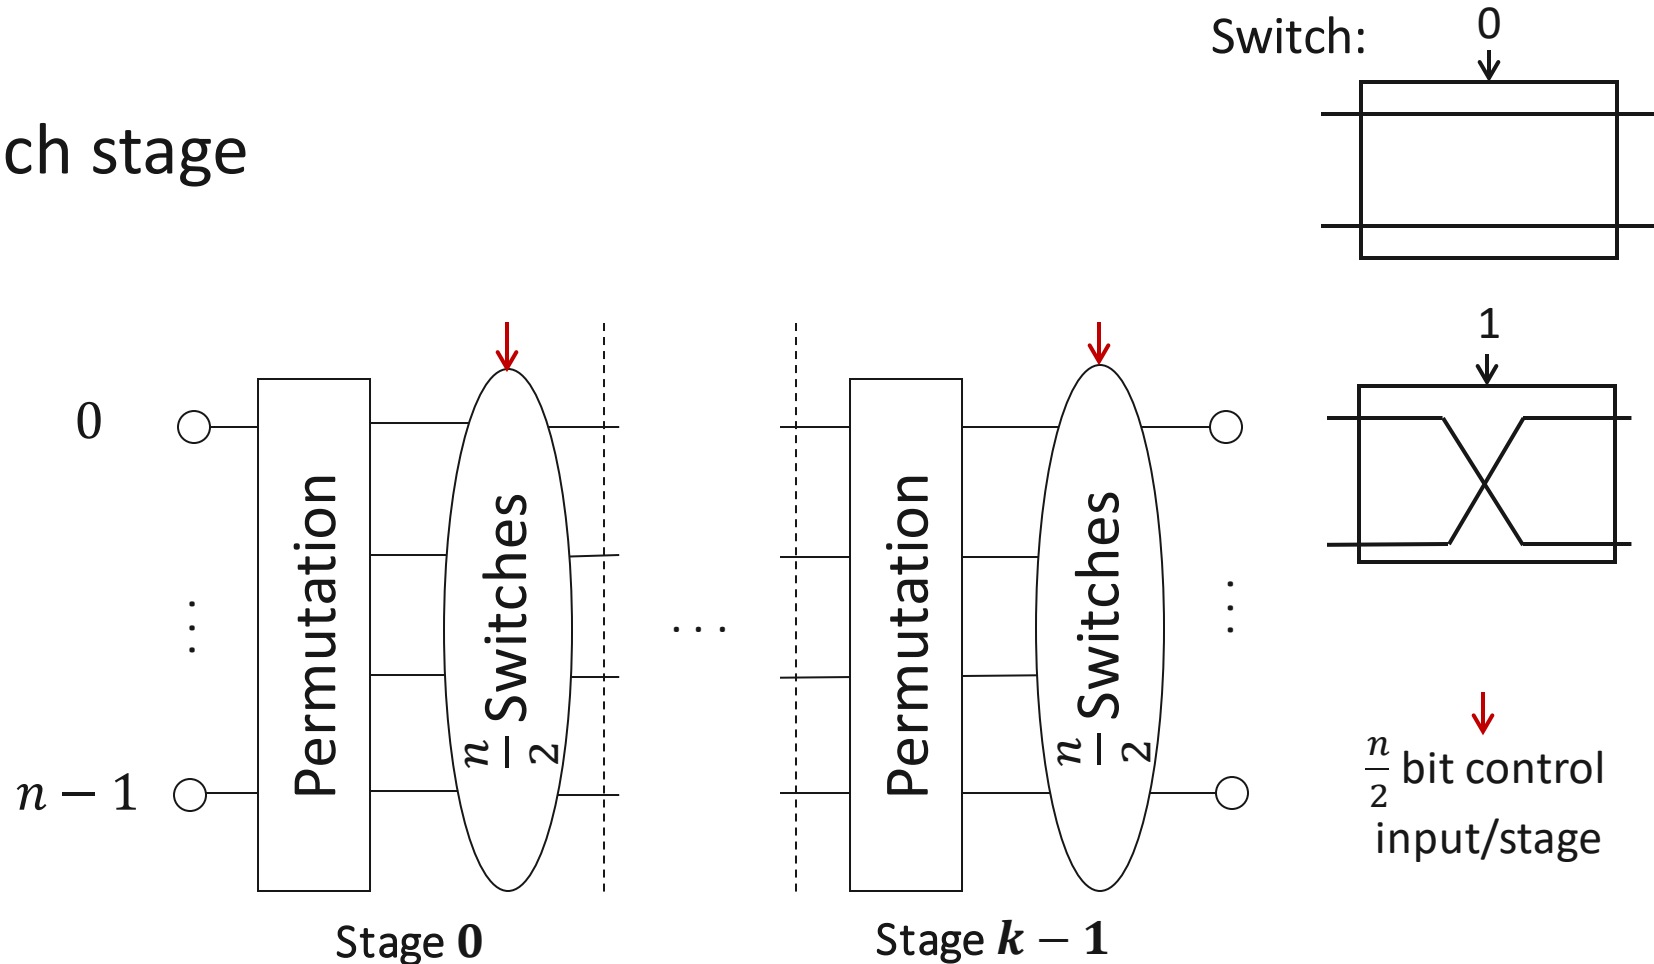




Multistage Network (4)

Multistage network

- k stages, $k \geq 1$
- $\frac{n}{2}$ switches in each stage
- delay = k





Performance of a Multistage Network

1 Switch = unit cost

unit delay

- Cost of a network: Total no. of 2×2 and 2×1 switches

Note: **Wiring cost ignored**


- Traditional metric
- Important metric
- Delay = Total # of stages

Note: **Wire delay ignored**



Example Multistage Network (1)

CROSS BAR Switch

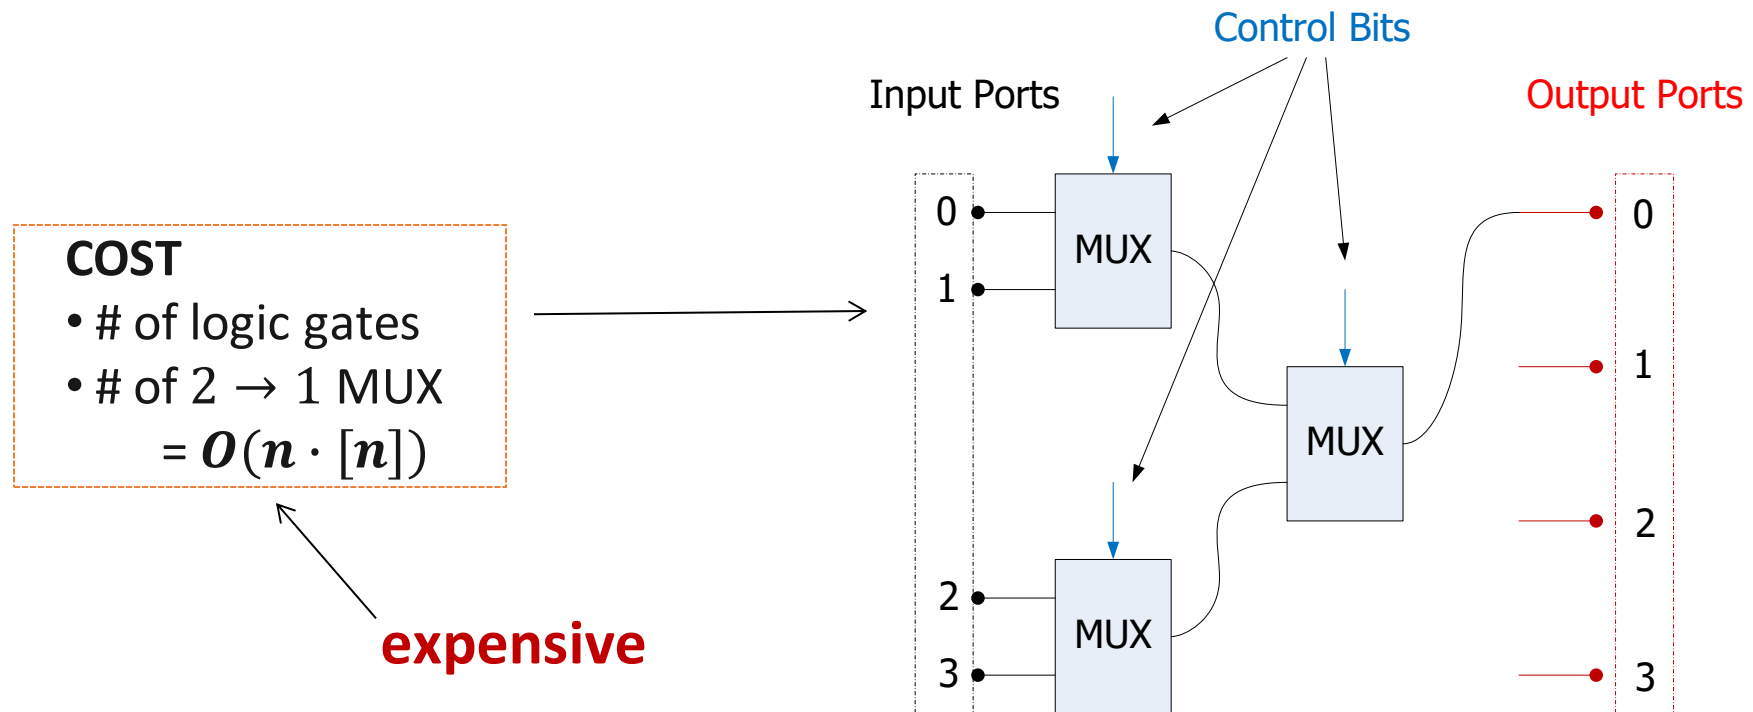


All $n!$ permutations can be realized




Example Multistage Network (2)

Example implementation using MUX






Total Number of Nodes in a Complete Binary Tree



$$4 + 2 + 1 = 7 = 2 \cdot n - 1$$
$$= 2 \cdot (\# \text{ of leaf nodes}) - 1$$




$$T(n) = 2T\left(\frac{n}{2}\right) + 1 = 15$$



Routing (1)

Example of Multistage Network



Permutation

$0 \rightarrow 2$

$1 \rightarrow 3$

$2 \rightarrow 1$

$3 \rightarrow 0$

- **Delay** = No. of stages
= 3 stages
- **Cost** = No. of switches
= $\frac{n}{2} \times (\text{number of stages})$
= 2×3

2^6 combinations
of switch settings

Note: All 4!
Permutations can be
realized by this network



Routing (2)

k stage, n input network



- Total number of switches: $\frac{n}{2} \cdot k$
- Total number of control bits: $\frac{n}{2} \cdot k$
 - Control bits specify a configuration of the network
 - Configuration \rightarrow permutation from input to output
- Total number of permutations that can be realized: $\leq 2^{nk/2}$
- If we want all $n!$ permutations to be realized: $2^{nk/2} \geq n!$

$$k = \text{no. of stages} = \Omega(\log n)$$



Omega Network (1)

- p inputs, p outputs
- $\log_2 p$ stages, each stage having $\frac{p}{2}$ 2×2 switches





Omega Network (2)

Multistage network

Omega network properties


- Cost = $\frac{P}{2} \log_2 p$ (number of switches)
- Note: in actual hardware design, routing cost dominates !
- Omega network can do $2^{(\frac{p}{2} \log_2 p)} < p!$ Permutations
 - All $p!$ permutations **can not** be realized
- Unique (only one) path from any input to any output



Omega Network (3)

Example of Blocking

- one of the messages (010 to 111 or 110 to 100) is blocked at link AB



Blocking: To realize a given connecting pattern, two inputs to a switch need to go through the same output port of the switch



Omega Network and Shuffle Exchange Network

n node SE network
(S+E) $\log_2 n$ times

\equiv

Omega network
with n inputs

Specify $\frac{n}{2}$ bits for
each exchange step

(We can pipeline here)


(Both networks realize the same set of permutations)



Congestion in a Network (1)

Given a routing protocol and data communication pattern (ex. a permutation)

Congestion at a node = $\text{Max. } \{ \# \text{ of paths passing through the node} \}$






Congestion in a Network (2)

Interconnection Network = Graph + Routing algorithm

Assume the routing algorithm provides unique (exactly one path) communication from source i to destination j for all i, j

For a given permutation:

- Congestion at node k = # of paths (based on the routing protocol) that pass through k





Congestion in a Network (3)

- Congestion for a given permutation in the network = $\text{Max}_{\text{over nodes } k} \{ \# \text{ of paths that pass through } k \}$
- Congestion in the network = $\text{Max}_{\text{all permutations}} \{ \text{congestion in the network for a given permutation} \}$




Example Computation of Congestion (1)

2-D Mesh Architecture

$n \times n$ mesh

Data $A(i,j)$ in $PE(i,j)$, $0 \leq i,j < n$






Example Computation of Congestion (2)

Permutation: Matrix Transpose

Data in $PE(i, j)$ [$A(i, j)$] $\rightarrow PE(j, i)$ $0 \leq i, j < n$

Row $i \rightarrow$ Col i $0 \leq i < n$





Example Computation of Congestion (3)

Routing

- Dimension ordered routing
- Widely used technique

For any permutation $(i, j) \rightarrow (d_i, d_j)$ $0 \leq i, j < n$

- 1) Go along row i to destination column d_j \leftarrow dimension 0
- 2) Go along column d_j to destination row d_i \leftarrow dimension 1



Example Computation of Congestion (4)

Row 0 \rightarrow Col 0


$(0,0) \rightarrow (0,0)$

$(0,1) \rightarrow (1,0)$

$(0,2) \rightarrow (2,0)$

$(0,3) \rightarrow (3,0)$

$n = 4$





Example Computation of Congestion (5)

Using dimension ordered routing, $n \times n$ mesh

- For any permutation
 - Congestion at any node $\leq n$
- Matrix Transpose results in worst case congestion
 - Congestion at $PE(i, i) = n, \quad 0 \leq i < n$
 - Congestion in the network $= n$



Summary

- Interconnection networks
 - Static/Dynamic network
 - Shuffle exchange network
 - Multistage network
 - Routing
 - Congestion