



EE/CSCI 451: Parallel and Distributed Computation

Lecture #11

9/24/2020

Viktor Prasanna

prasanna@usc.edu

ceng.usc.edu/~prasanna

University of Southern California



Announcement

- Midterm 1 on 9/25
- PHW3 due 9/24
- HW4 due 9/24
- HW3 grades are out

HW3 Statistics	
Average	90.4
Median	94
Standard Deviation	13.9



Announcement: Midterm 1 Logistics

- **Open Book, Open Notes**
- Time: 9/25 3:30-5:30 PM (Los Angeles time)
- Exam will be released on **Piazza**
- Upload your answer (pdf file) on **Blackboard**
- Completing your exam:
 - Option 1 - Download the exam as a pdf file onto your tablet and annotate it with your answers.
 - Option 2 – Download and print the exam and write your answers on the (printed) paper. Scan the paper and save into PDF format.
- Turn on your camera on Zoom: We will be proctoring!



Midterm 1 Details

- Materials covered till the end of last week
- 5 problems in total
 - Memory System Performance Modeling
 - Shared Memory Programming
 - Shared Memory Programming
 - Message Passing Protocols & Programming
 - Interconnection Networks

Course Info.



- Academic Integrity
 - Cheating will not be tolerated
 - Grade of F will be assigned
 - Cheaters will be reported to USC Student Judicial Affairs and Community Standards (SJACS)



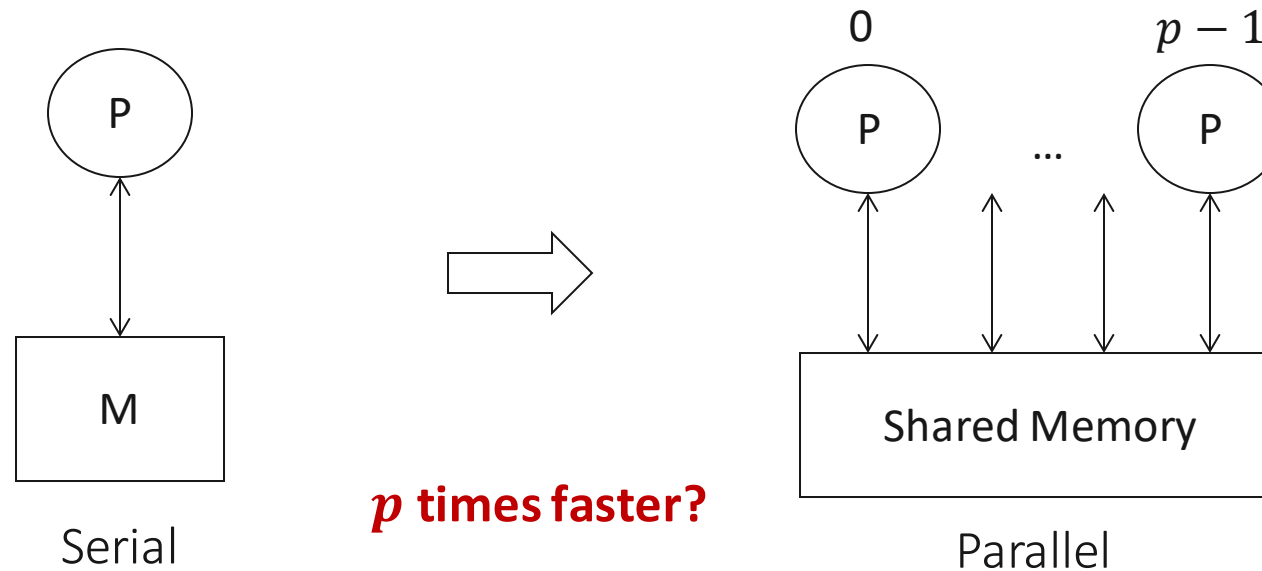
Outline

- From last class
 - Program and Data Mapping
 - Graph Embedding
 - Dilation
 - Expansion
 - Congestion
 - Network Model
 - Simulations
 - Hypercube on 1-D mesh (dilation and congestion)
- Today
 - Analytical Modeling of Parallel Systems (Chapter 5.2, 5.4.1)
 - Scalability
 - Achievable Speedup
 - Amdahl's Law
 - Gustafson's Law
 - Efficiency
 - Work Optimal parallel solution
 - Performance analysis
 - Big O notation



Scalability (1)

Does performance (execution time) improve as we use more resources (processors)





Scalability (2)

$$\text{Speedup} = \frac{\text{Serial time (on a uniprocessor system)}}{\text{Parallel time using } p \text{ processors}}$$

If speedup = $O(p)$, then it is a **scalable** solution



Overheads in Parallel Computation


- Communication
- Coordination
- Load balance (processors may idle)




Amdahl's Law (1)

- Amdahl's Law — Limit on speedup achievable when a program is run on a parallel machine
- Given an input program

- Total execution time: $S + P$



- Time on a uni-processor machine: $1 = S + P$
 - Time on a parallel machine: $S + P/f$
 f = speedup factor





Amdahl's Law (2)

Example:

1 unit of time = 1 arithmetic operation

$S = \frac{n}{2n}$	$A(0) \leftarrow 0$	}	n arithmetic operations Serial operations
$P = \frac{n}{2n}$	Do $i = 1$ to n $A(i) \leftarrow A(i) + A(i - 1)$		
$S + P = 1$	Do $i = 1$ to n $A(i) \leftarrow A(i) * A(i)$	}	n arithmetic operations Parallelizable operations



Amdahl's Law (3)

Overall speedup


$$= \frac{\text{Serial Time}}{\text{Parallel Time}}$$

$$= \frac{S + P}{S + P/f}$$

$$= \frac{1}{S + P/f}$$

$$\leq \frac{1}{S}$$

f = speedup factor




If serial portion is 50 %, then the **overall speedup ≤ 2**

Note: Speedup factor (f) \leq # of processors



Amdahl's Law (4)


$$\text{Overall Speedup} = \frac{1}{S + P/f} = \frac{1}{S + (1 - S)/f} \rightarrow \frac{1}{S}$$





Amdahl's Law (5)

(Overall) Speedup is upper bounded by $\frac{1}{\text{Serial portion}}$ $\left(\frac{1}{S}\right)$






Amdahl's Law (6)

An alternate interpretation

- **Amdahl's Law for Energy:** Limit on energy improvement when only part of the energy consumption can be improved
- Given a program
 - Total energy consumed by a code: 1
 - The portion of energy that can not be improved: e (ex. Static Power)
 - The portion of energy that can be improved: $1 - e$ (ex. Dynamic Power)
 - Energy dissipation improvement factor = f






Amdahl's Law (7)

Overall Energy Improvement

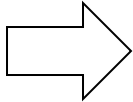
$$= \frac{1}{e + (1 - e)/f}$$
$$\leq \frac{1}{e} \quad (\text{When } f \text{ is large})$$



If e is 50 %, the overall energy improvement ≤ 2



Scaled Speedup (Gustafson's Law) (1)


- Amdahl's Law — Serial portion of code limits performance
(As we use more processors)
- As we use more processors
 - we use more data
 - e.g., more fine grained model

more opportunities
for parallelism
- E.g. Processing $N \times N$ image
Using $p \times p$ processor array
As we increase p we usually increase image size




Scaled Speedup (Gustafson's Law) (2)

Amdahl's Law: Fixed amount of computations



Gustafson's Law: Increase p and amount of computations



If parallelism scales linearly with p , number of processors

$$\text{Scaled Speedup} = \frac{\text{Serial time}}{\text{Parallel time}} = \frac{S' + (1 - S')p}{S' + P'} = \frac{S' + (1 - S')p}{1}$$



Scaled Speedup (Gustafson's Law) (3)


Gustafson's Law

$$\begin{aligned}\text{Scaled Speedup} &= S' + (1 - S')p \\ &\approx (1 - S')p\end{aligned}$$

$S' = 0 \rightarrow$ Scaled Speedup = Ideal speedup (p)

$S' = 0.5 \rightarrow$ Scaled Speedup = $0.5 p$, 50% of Ideal speedup

Example: Cloud



$$\text{Scaled Speedup} \propto (1 - S') \times \# \text{ of threads}$$




Scaled Speedup (Gustafson's Law) (4)

Gustafson's Law

- If we increase
 - the number of processors (p)
 - the amount of parallelizable portion of computations
- Scaled speedup is limited by the fraction of program that can be parallelized (higher the fraction, higher the speedup).



Scaled Speedup (Gustafson's Law) (5)






Strong or Weak Scaling

- Strong Scaling
 - Governed by Amdahl's Law
 - The number of processors is **increased** while the problem size **remains constant**
 - Results in a **reduced** workload per processor
- Weak Scaling
 - Governed by Gustafson's Law
 - **Both** the number of processors and the problem size are **increased**
 - Results in a **constant** workload per processor



Superlinear Speedup (1)


- Achieved speedup $>$ number of processors
 - Hardware features (ex. cache effect)





Superlinear Speedup (2)

- Achieved speedup > number of processors
 - Work performed by serial algorithm is greater than its parallel formulation






Performance (1)

Efficiency


- Question: If we use p processors, is speedup = p ?
- Efficiency \triangleq Fraction of time a processor is usefully employed during the computation
- $E = \text{Speedup} / \# \text{ of processors used}$
 - E is the average efficiency over all the processors
 - Efficiency of each processor can be different from the average value





Performance (2)

Ex. $S = 0.5$
 $P = 0.5$ 2 processor system



$$\text{Speedup} = \frac{1}{0.75} = 4/3$$

$$(\text{Average}) \text{ Efficiency} = \frac{4/3}{2} = 2/3$$

$$\text{Efficiency of } P_1 = \frac{0.25}{0.75} = 1/3$$

$$\text{Efficiency of } P_2 = \frac{0.75}{0.75} = 1$$

$$(\text{Average}) \text{ Efficiency} = \frac{1/3 + 1}{2} = 2/3$$



Performance (3)

- Cost = Total amount of work done by a parallel system
= Parallel Execution Time x Number of Processors
= $T_p \times p$
- Cost is also called **Processor Time Product**
- COST OPTIMAL (or WORK OPTIMAL) Parallel Algorithm
 - Total work done = Serial Complexity of the problem



Performance (4)

- Example: addition on PRAM
 - n processor PRAM
 - n input data
 - Add n numbers



Performance (5)

- Algorithm

Program in processor j , $0 \leq j \leq n - 1$

1. Do $i = 0$ to $\log_2 n - 1$
2. If $j = k \cdot 2^{i+1}$, for some $k \in N$
 then $A(j) \leftarrow A(j) + A(j + 2^i)$
3. end


Note:

A is shared among all the processors

Synchronous operation [For ex. all the processors execute instruction 2 during the same cycle, $\log_2 n$ time]

N = set of natural numbers = $\{0, 1, \dots\}$

Parallel time = $O(\log n)$





Performance (6)

Serial time = $O(n)$ (Serial complexity)

Parallel time = $O(\log n)$

Speedup = $O(n/\log n)$

of processors = n

$$E = \frac{O(n/\log n)}{n} = O(1/\log n)$$

NOT WORK OPTIMAL



Performance Analysis (1)

Asymptotic Analysis

Big O Notation or Order Notation

Worst case execution time of an algorithm

Upper bound on the growth rate of the execution time

Example: $n \times n$ matrix multiplication

```
1. Do  $i$ 
2.   Do  $j$ 
3.      $C(i,j) \leftarrow 0$ 
4.     Do  $k = 1$  to  $n$ 
5.        $C(i,j) \leftarrow C(i,j) + A(i,k) * B(k,j)$ 
6.     End
7.   End
8. End
```

$T(n)$ = time complexity function = $n^2 + n^3 + n^3$



Performance Analysis (2)

- Actual execution time depends on the processor infrastructure, compiler, etc.

Number of computation steps is upper bounded by cn^3

For some constant c , c **does not** depend on n .

We say $T(n) = O(n^3)$

- Definition:


$f(n)$ is $O(g(n))$ if there is a constant c such that
 $f(n) \leq c \cdot g(n)$ for sufficiently large n , i.e. there
exists n_0 such that for all $n \geq n_0$, $f(n) \leq c \cdot g(n)$

- Ex: $f(n) = 2n^3 + n^2 + n$
 $f(n) = O(n^3)$

Performance Analysis (3)



$$f(n) = O(g(n))$$






Performance Analysis (4)

- Lower bound on execution time
 - Definition: $f(n) = \Omega(g(n))$
if there exist constants c and n_0 such that
for all $n \geq n_0$, $f(n) \geq c \cdot g(n)$
ex: $f(n) = n^3 + n^2 + n$ then, $f(n) = \Omega(n^3)$
- Tight bound on execution time
 - Definition: $f(n) = \Theta(g(n))$
if $f(n) = O(g(n))$ and $f(n) = \Omega(g(n))$

Performance Analysis (5)



$$f(n) = \Omega(g(n))$$





Performance Analysis (6)

Ex. Execution time of two algorithms

$$A_1: T_1(n) = 100n$$

$$A_2: T_2(n) = n^2$$

A_1 is asymptotically superior to A_2

Ex.

$$A_1: T_1(n) = 5n^2$$

$$A_2: T_2(n) = 100n^2$$

A_1 and A_2 are asymptotically of the same complexity

$$T_1(n) = T_2(n) = \theta(n^2)$$



Performance Analysis (7)

$$f_1(n) = n$$

$$f_3(n) = n^2$$

$$f_2(n) = n \log n$$

$$f_4(n) = n^{1+\varepsilon}, \quad 0 < \varepsilon < 1$$

$$f_1(n) = O(f_2(n))$$

$$f_2(n) = O(f_3(n))$$

$$f_2(n) = O(f_4(n)) \quad \checkmark \qquad f_2(n) = \Omega(f_3(n)) \quad \times$$

$$f_4(n) = O(f_3(n)) \quad \checkmark \qquad f_2(n) = \Omega(f_4(n)) \quad \times$$



Summary

- Scalability
- Achievable Speedup
 - Amdahl's Law
 - Gustafson's Law
- Efficiency
 - Processor Time Product
 - COST OPTIMAL (work optimal)
- Performance Analysis