



EE/CSCI 451: Parallel and Distributed Computation

Lecture #15

10/8/2020

Viktor Prasanna

prasanna@usc.edu

ceng.usc.edu/~prasanna

University of Southern California



Announcement

- PHW4 due on 10/9 (Friday)
- HW6 due on 10/9 (Friday)
- PHW5 and HW7 will be out tomorrow
 - PHW5 due on 10/22
 - HW7 due on 10/16
- Project Proposal due on 10/16
 - Remember to fill your group info [here](#)
- Midterm 2
 - Time: Oct. 23 (Friday) 3:30-5:30 PM
 - Covers material from Sept. 22 to Oct. 16
 - Logistics: same as Midterm 1



Course Project

- Project Proposal Template
 - Introduction
 - Background on the problem
 - Contributions/hypotheses of the project -- **parallelization, scalability** should be the focus
 - Description of the approach and implementation
 - Parallelization techniques if proposing a new technique
 - Selection criteria for benchmarks if comparing parallel algorithms
 - Description of the application if developing an application
 - Implementation platform and resource requirements
 - Evaluation methodology/Criteria for success
 - Datasets, evaluation metrics, and baselines for comparison
- We will evaluate the proposals and give feedback

Course Project



- Proposal Format
 - 12 pt font size
 - Single-column
 - Single space
 - Cite relevant papers
 - Length: 2-4 pages including references & appendix(if any)
 - A template proposal can be found Piazza under Resources-Course Project



Outline

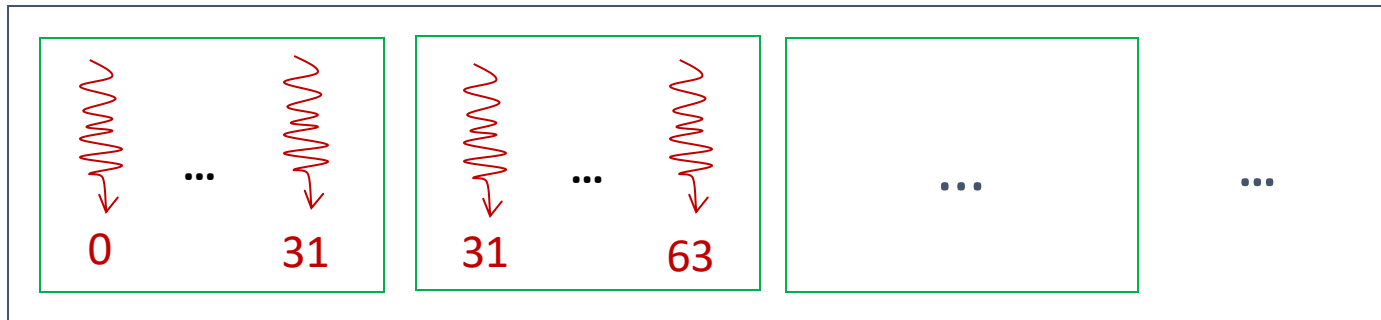
- Architecture, Programming and Execution Model of (GP) GPU
 - GPU and data parallel programming
 - CUDA Programming Model
 - GPU Architecture
 - **Execution Model**
 - Performance Model
 - Examples
 - CUDA, OpenCL, ...
 - APU, Zynq...



Warp

- Basic unit of execution on GPU
- 32 threads form a Warp
 - share the same program counter (PC)
- Threads within a Warp execute the same instruction in each clock cycle (SIMT)

Thread block





Basic Organization

Typical organization

16 SM (Streaming Multiprocessors)

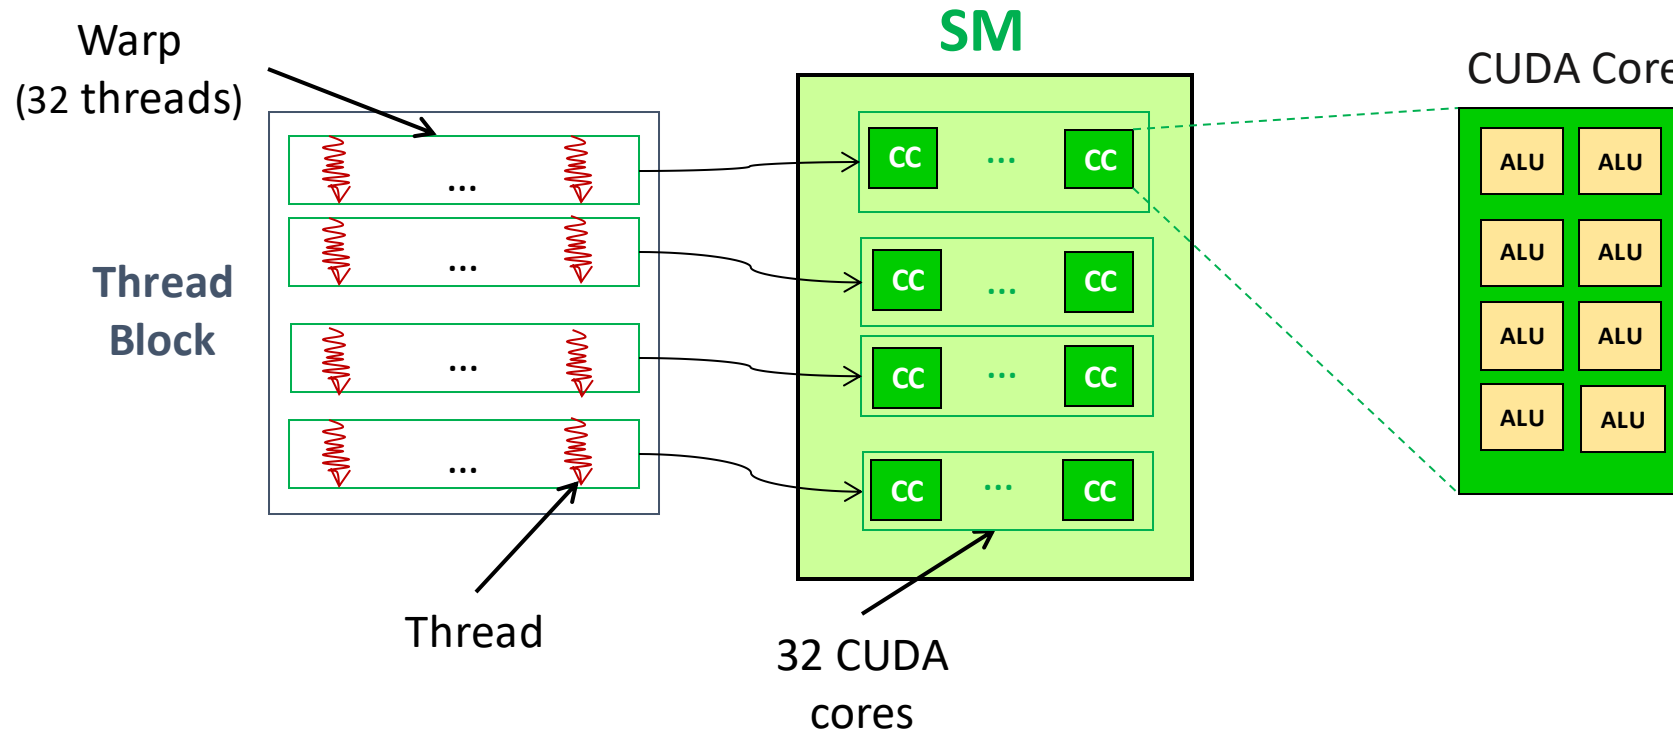
128 CC/SM → Up to 4 Warps can be concurrently executed on each SM

8 ALU/CC SIMT – (single instruction, multiple threads)



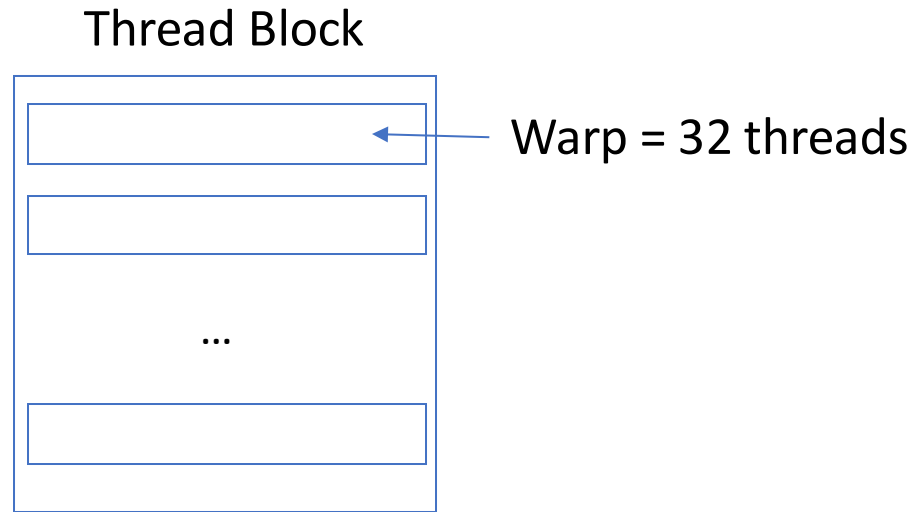
SIMT Execution Model (1)

- Threads within the same thread block are mapped onto the same streaming multiprocessor
- Each Warp is executed by set of CUDA cores in SIMD fashion





SIMT Execution Model (2)



- SIMD within each Warp
- Each Warp is executed by collection of CC
- No guarantee all the Warps are executing the same instruction at the same time
- SIMT programming model does NOT guarantee SIMD execution across all the threads in the block



SIMT Execution Model (3)

- Single Instruction, Multiple Thread
- Thread block is partitioned into groups of 32 threads (Warp)
 - Always executing the **same** instruction
 - Shared instruction fetch/dispatch
 - Some threads become inactive when code path diverges
 - Hardware automatically handles divergence
- Warps are the primitive unit of scheduling
 - Largely invisible to the programmer
 - Must understand for performance [not correctness ?]



SIMT Execution Model (4)

- SIMT is a virtualized execution of SIMD
- # of threads can be $>$ # of hardware units
 - For example:
32 threads in Warp and 16 CCs assigned to it
- Within Warp it is SIMD but no guarantee outside each Warp
 - Warp #0 of a thread block may be executing instruction i while Warp #1 of the same thread block may be executing instruction j , $i \neq j$



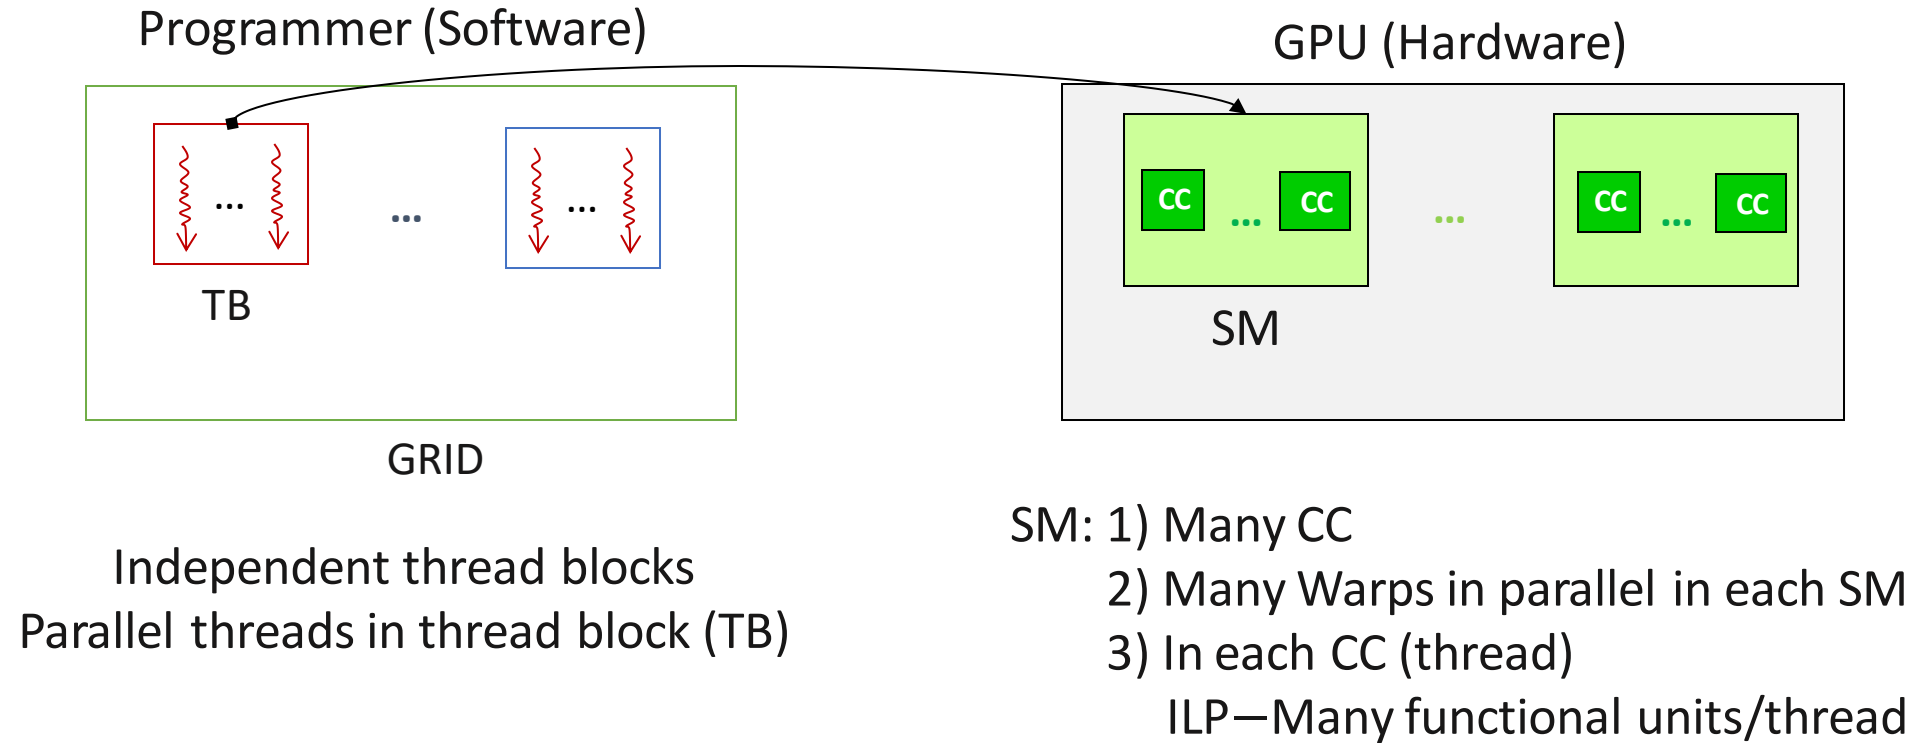
SIMT Execution Model (5)

- A Warp (32 threads) is executed by a set of CC (ex. 32 CC in 1 clock cycle or 16 CC in 2 clock cycles)





Parallelism (1)





Parallelism (2)

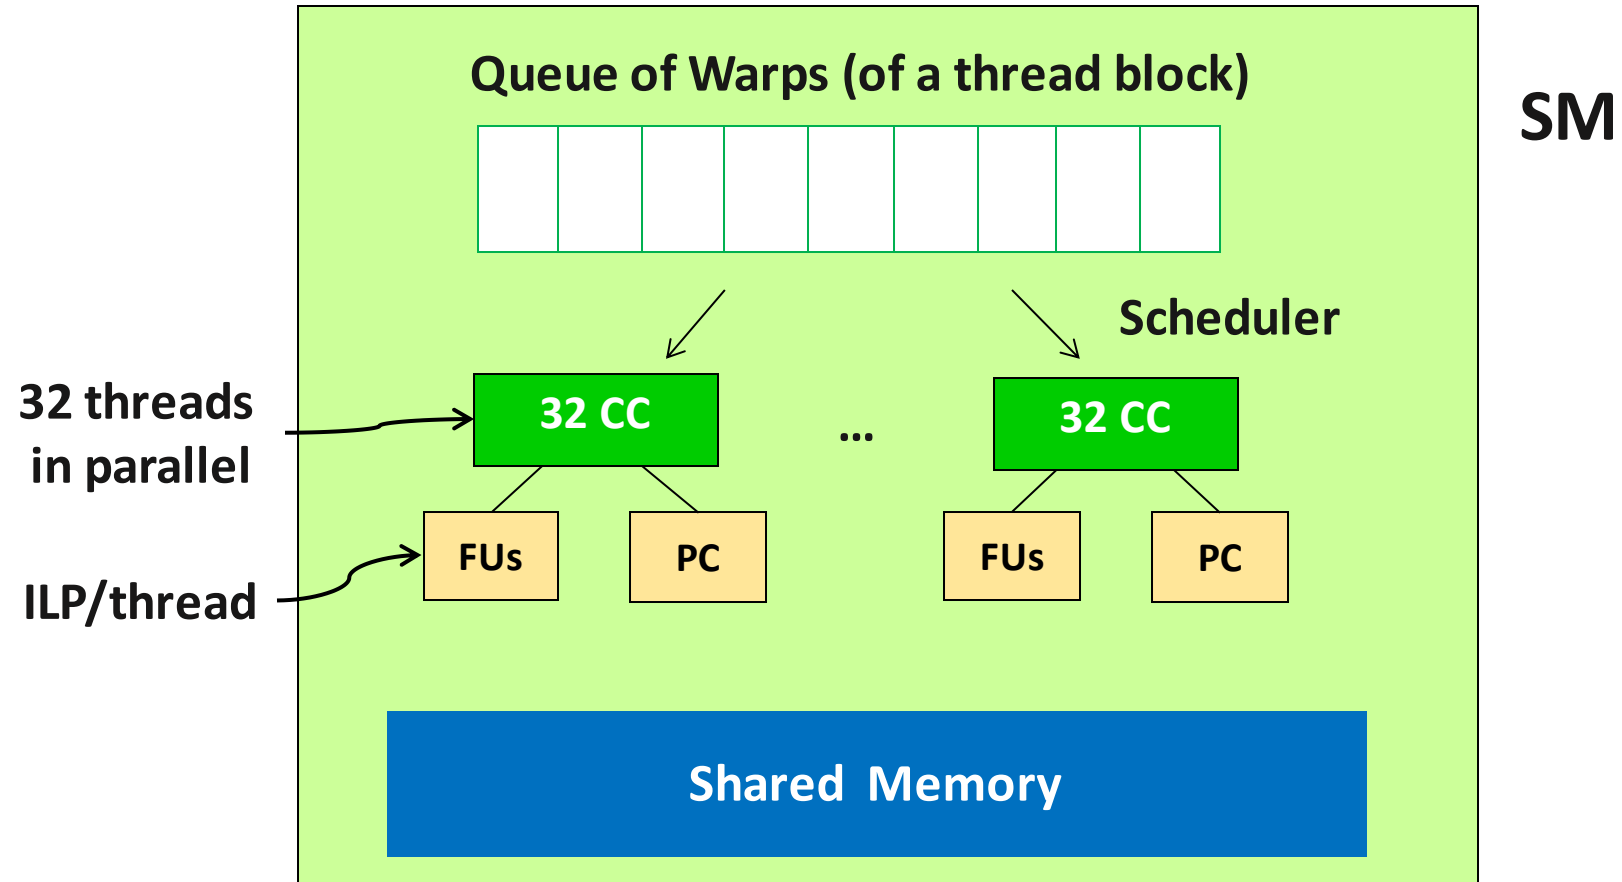
How many program counters (PC) in each SM?

For every Warp, we need a PC

$$\# \text{ of Warps can be } > \frac{\# \text{ of CC}}{32}$$



Parallelism (3)





Parallelism (4)

Restriction

In each thread block — only one program!

Not **all** applications have such parallelism.



Synchronization (1)

- At run time, Warps in the same thread block may execute different instructions in the same clock cycle

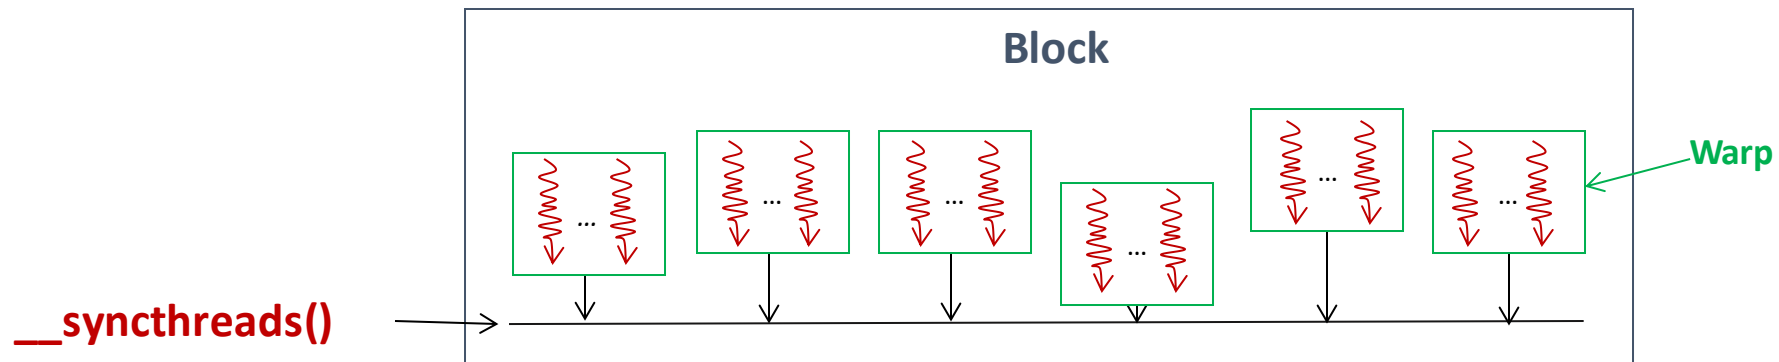
	Warp 0	Warp 1	Warp 2	Warp 3
Cycle 0	Execute Ins. 0	Execute Ins. 0	×	×
Cycle 1	Execute Ins. 1	×	×	×
Cycle 2	Execute Ins. 2	×	×	Execute Ins. 0
	Execute Ins. 3	×	×	Execute Ins. 1
Time	↓			

×: waiting to be scheduled, blocked due to I/O, global memory access, ...



Synchronization (2)

- **Need for Synchronization:** When some threads within a block access the same addresses in shared or global memory, there are potential read-after-write, write-after-read, or write-after-write hazards
- **__syncthreads()** is a block level synchronization barrier;
wait until all threads in the thread block have reached this point





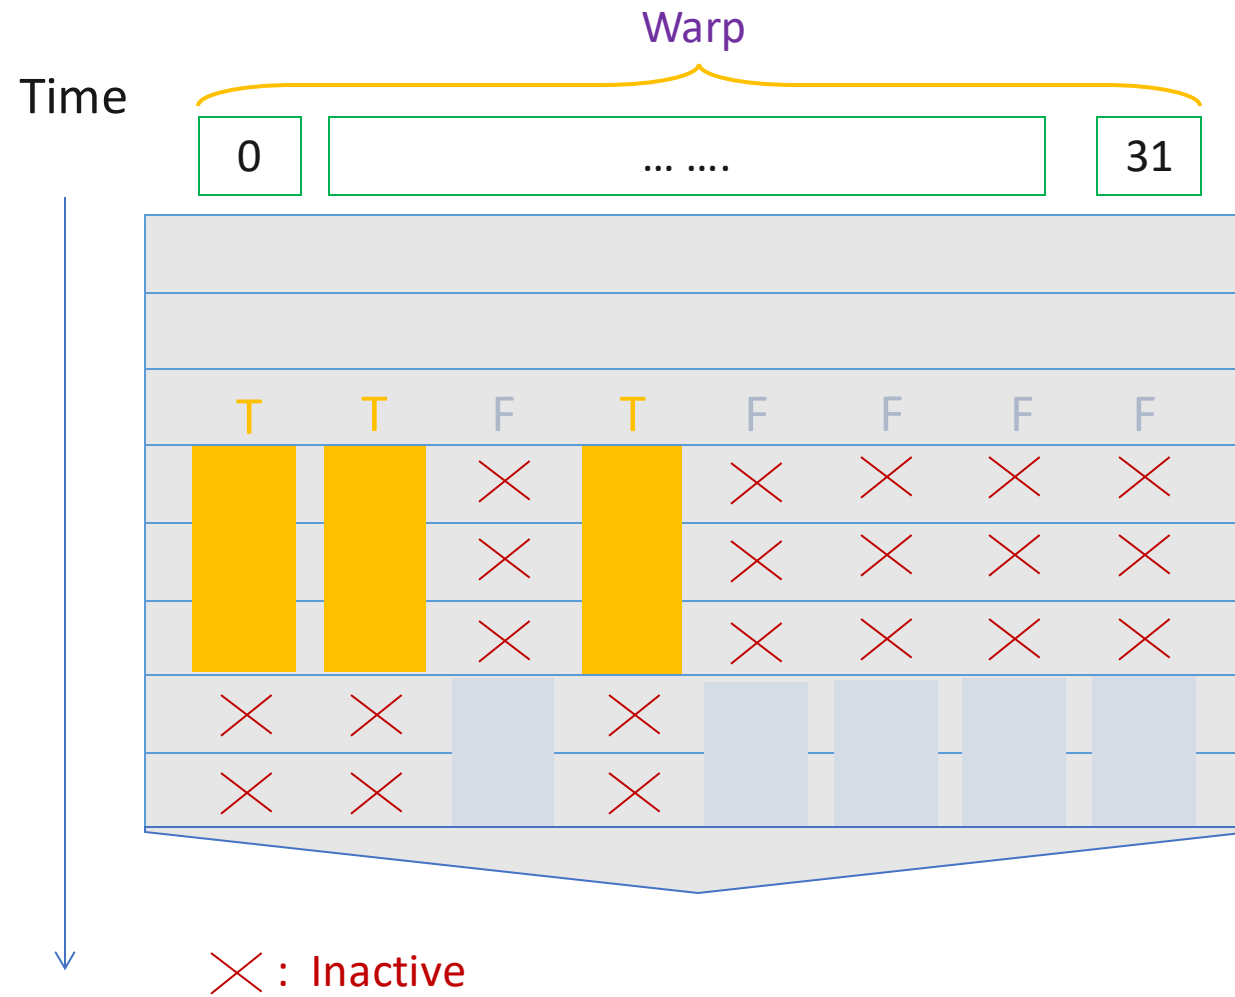
Branching (1)

- Branching (e.g., if-else, for, while, do, switch, etc.) causes a **divergence** in the flow of execution
- GPU (SM) simply executes one path of the branch and then the other (i.e., serial execution)
- Those threads that take the branch are executed and those that do not are marked as inactive
- Once the taken branch is resolved, the other side of the branch is executed, until the threads converge once more

```
if (some_condition) then {  
    Action_a();  
}  
else {  
    Action_b();  
}
```



Branching (2)



<unconditional code>

```
if (x > 0) {
```

Action_a()

```
} else {
```

Action_b()

```
}
```

<unconditional code>



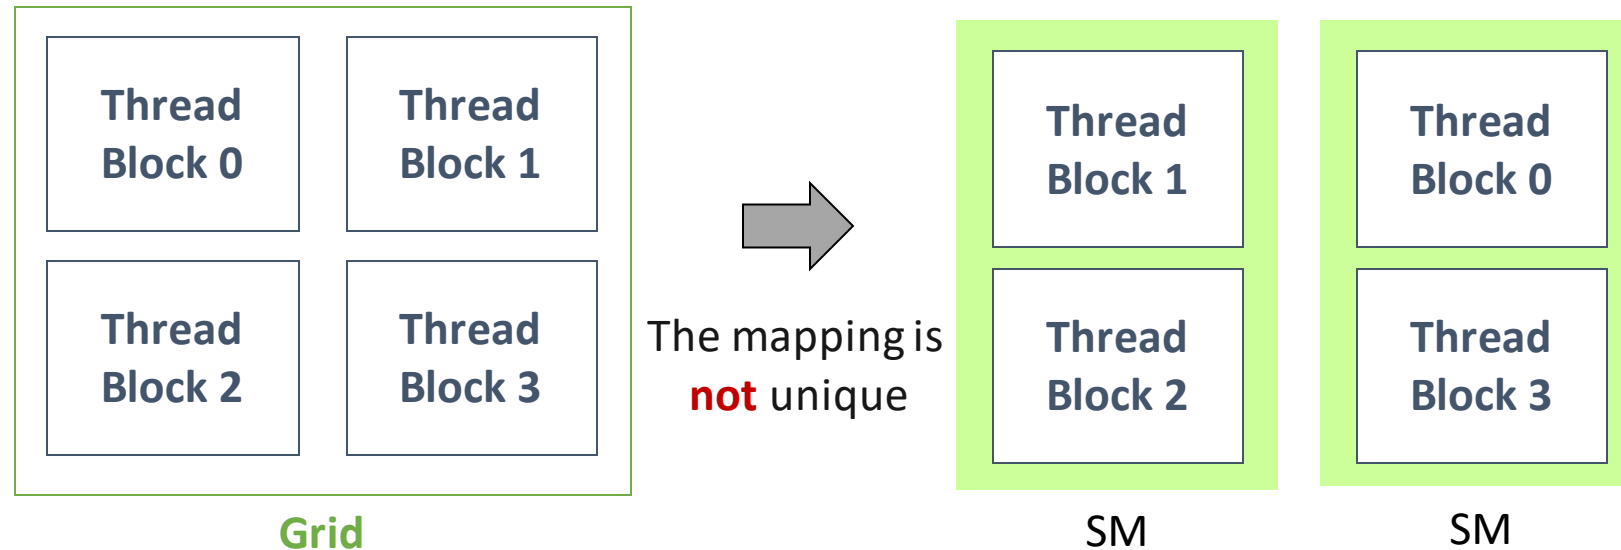
Impact of Divergence (branching)

- Serial execution of the two codes (then, else portions)
- Severe performance degradation



Global Execution Model

- Kernel = Collection of thread blocks
- Concurrent blocks processed using SM

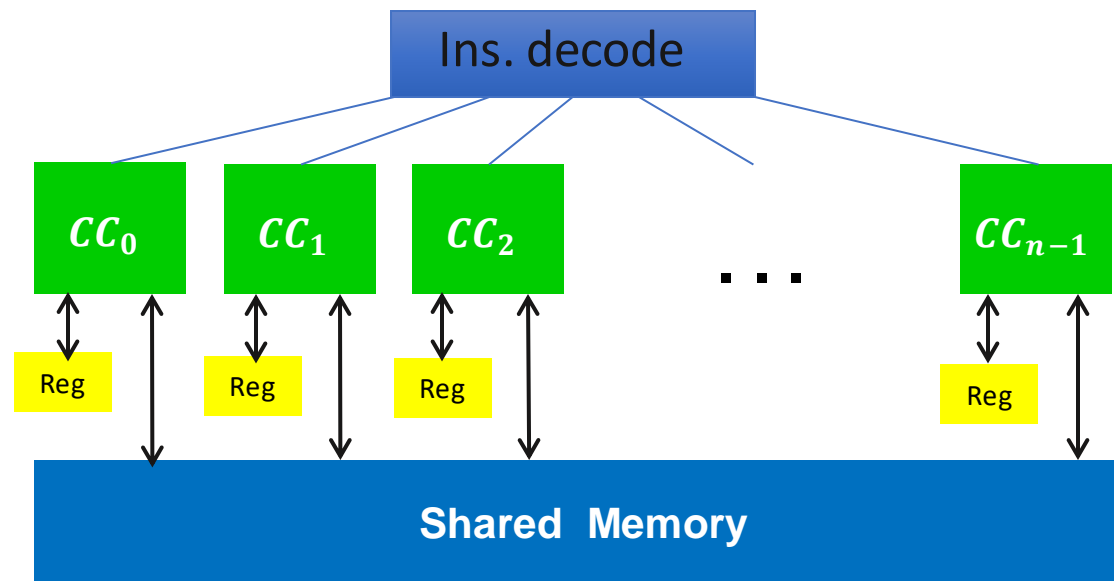


Note: All threads in a block execute the same code



Local Execution Model

- SIMT (Single Ins. Multiple Threads) (in SM)
- Parallel execution across CUDA Cores (CC)



$n=16,32$



Outline

- **Architecture, Programming and Execution Model of (GP) GPU**
 - GPU and data parallel programming
 - Programming Model
 - GPU Architecture
 - Execution Model
 - **Performance Model**
 - Examples
 - CUDA, OpenCL, ...
 - APU, Zynq...

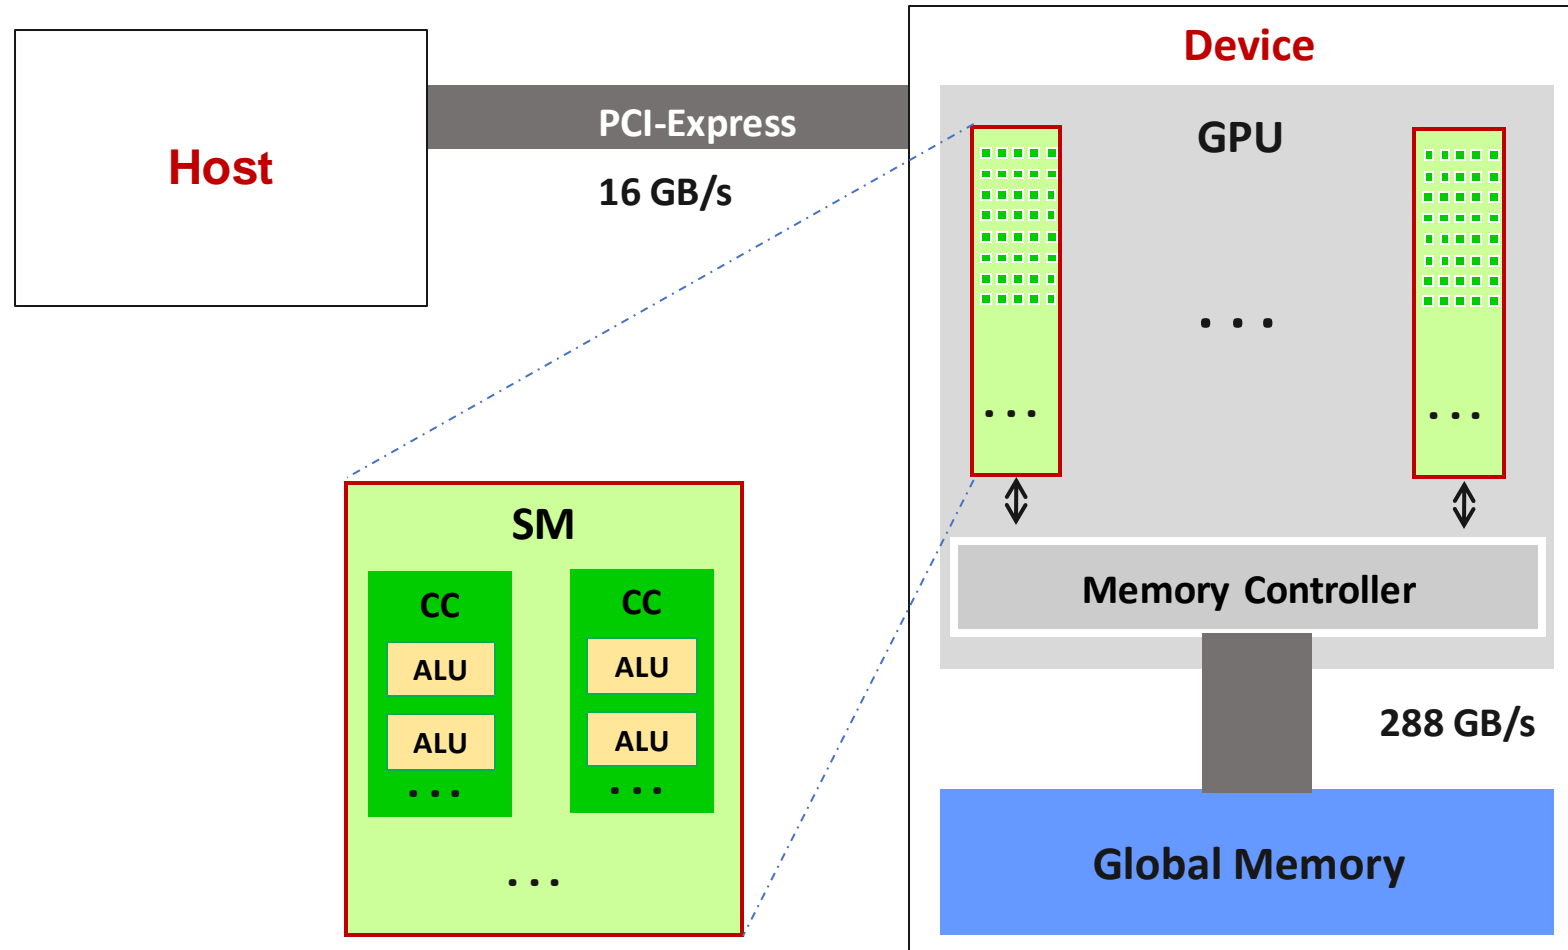


A Simple Performance Model (1)

- ***N*** No. of threads in a grid
- ***D*** GB of input data used for a kernel
- ***M*** GB of data accessed from global memory to the device for each thread
- ***C*** Tflops executed in each thread
- Example Device Configuration
 - 16 SM in a Device, 128 CC/SM, 8 ALU/CC, 2 flops/ALU/cycle, 1 GHz
 - Peak performance
 - $16 * 128 * 8 * 2 * 1 \text{ G} = 32 \text{ TFLOPS}$
 - Bandwidth
 - PCI-e 3.0 (x16) : 16 GB/s
 - GDDR5 memory controller : 288 GB/s



A Simple Performance Model (2)





A Simple Performance Model (3)

- Assume:
 - no diverging conditional branch
 - No memory bank conflicts

- Total Execution time for a kernel

$$T(\text{DataTransferCPU} \leftrightarrow \text{Device}) + (\# \text{ Threads} / \# \text{ CC}) \times T(\text{Thread})$$

$$= 2 \times (D/16) + N/(16 \times 128) \times [T(\text{MemOpsglobal}) + T(\text{Computation Ops})]$$

$$= 2 \times (D/16) + N/(16 \times 128) \times [M/288 + C/32 \text{ TFLOPS}]$$



Outline

- **Architecture, Programming and Execution Model of (GP) GPU**
 - GPU and data parallel programming
 - Programming Model
 - GPU Architecture
 - Execution Model
 - **Examples**
 - CUDA, OpenCL, ...
 - APU, Zynq...



Matrix Multiplication (1)

$$\text{Example: } C = A \times B$$

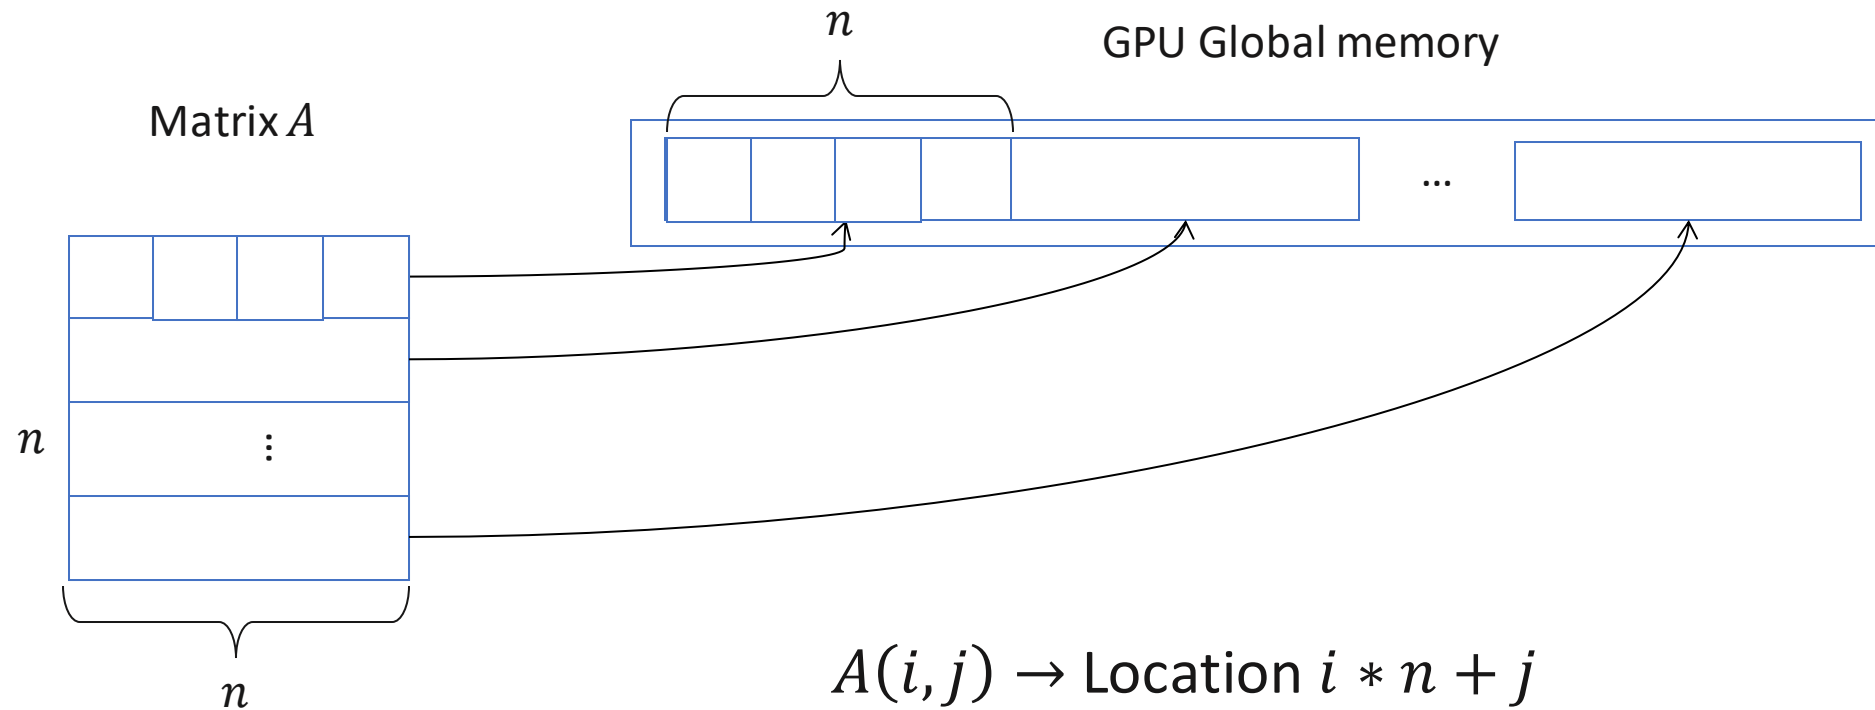
Matrix size = 1024×1024

- Allocate device memory for A, B, C
- Load A and B to GPU memory
- Invoke kernel function
- Copy result from device to host



Matrix Multiplication (2)

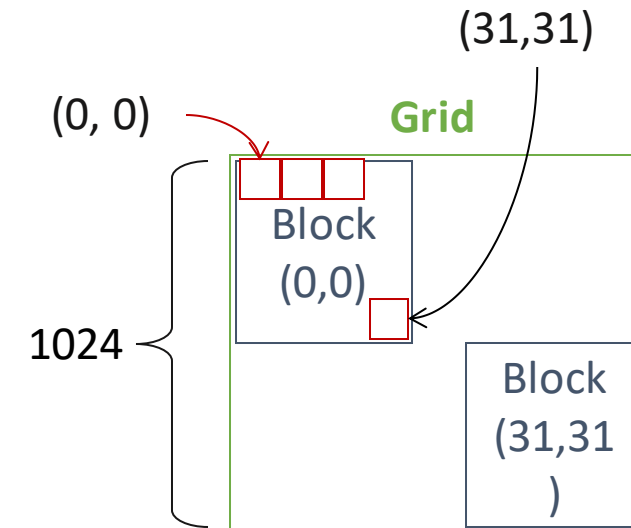
- Example Data layout: All matrices are stored in **row major** order





Matrix Multiplication (3)

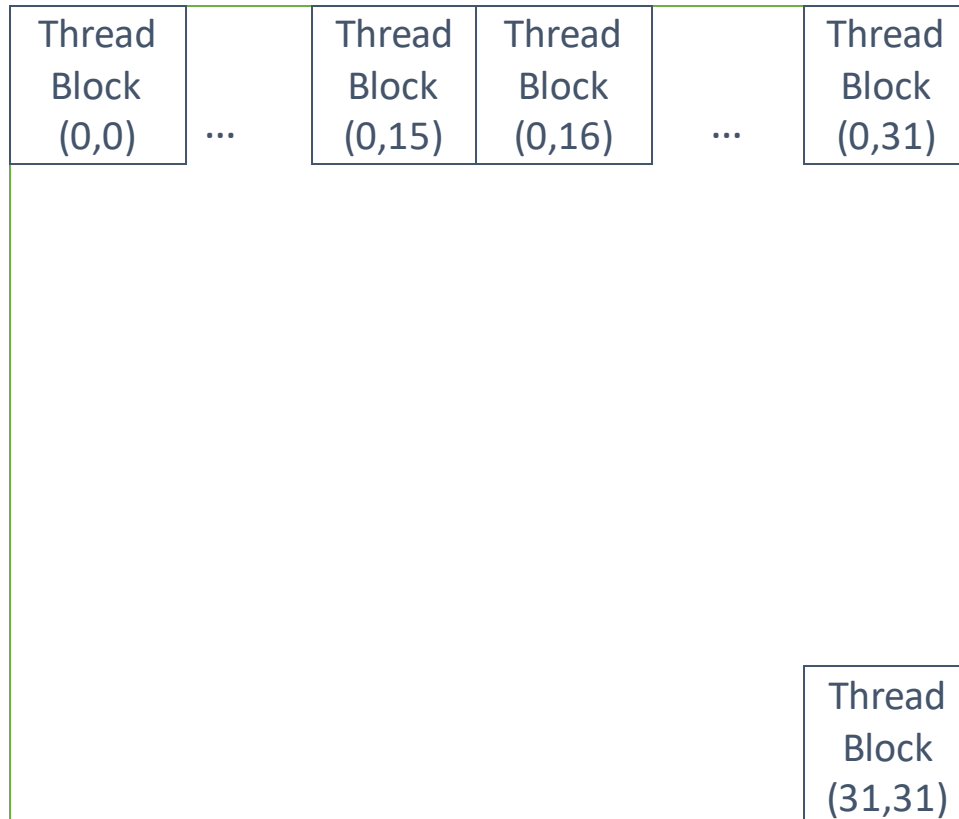
- Total # of threads = $1K \times 1K$
 - 1 thread for each element of C
- Example grid, thread block, thread definition
 - 2-D grid: 32×32 – Thread blocks
 - 2-D block: 32×32 – Threads/Block
 - Each thread computes one element of C





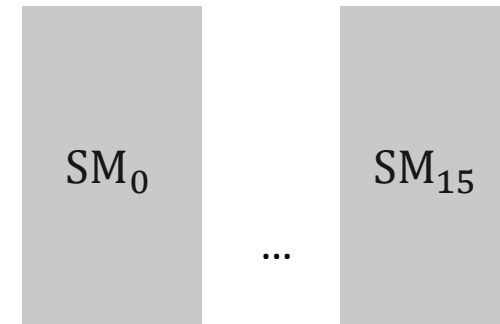
Matrix Multiplication (4)

Grid and Blocks and mapping



blockID.x = 31
blockID.y = 31

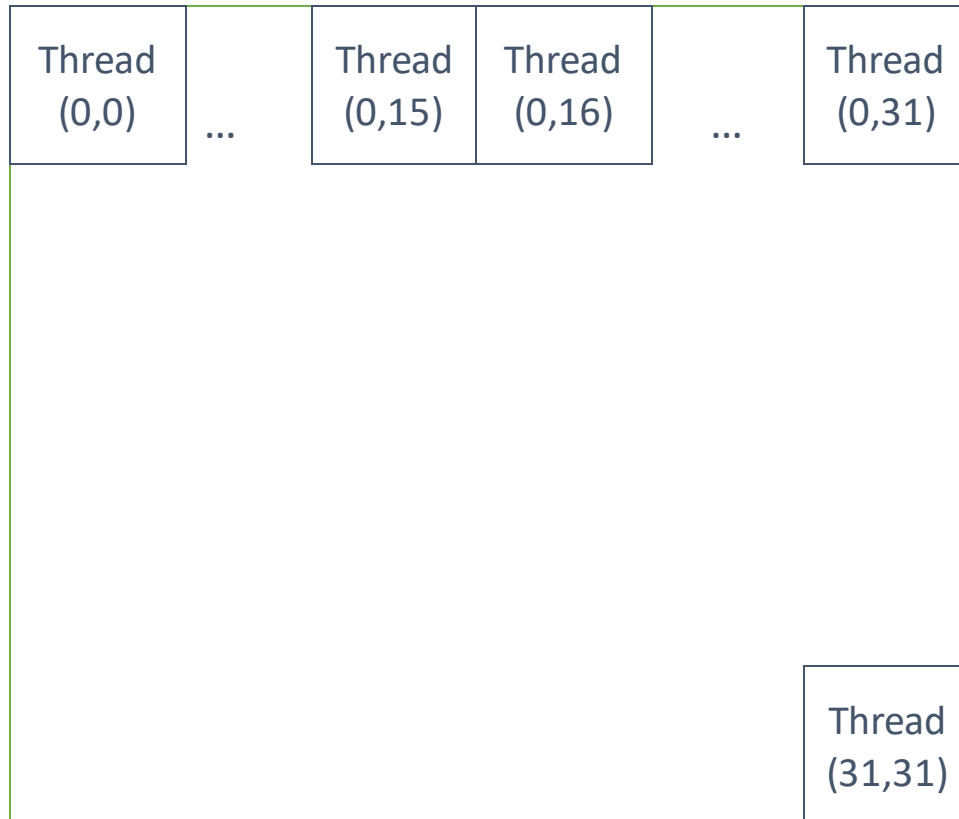
Example GPU organization



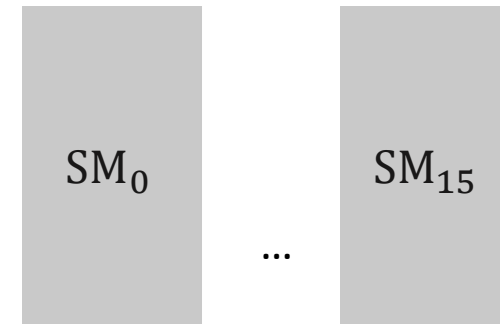


Matrix Multiplication (4)

Thread Block



Example GPU organization



threadID.x = 31
threadID.y = 31



Matrix Multiplication (5)

Algorithm 1 (straight forward inner product, no optimizations)

Grid configuration:
1-D or 2-D

Block configuration:
1-D or 2-D or 3-D

- Kernel function $\ll \text{dim_grid}, \text{dim_block} \gg (A, B, C) \{$

`Local_C = 0;`

Local register

`row` = row index of the thread in the matrix

`col` = column index of the thread in the matrix

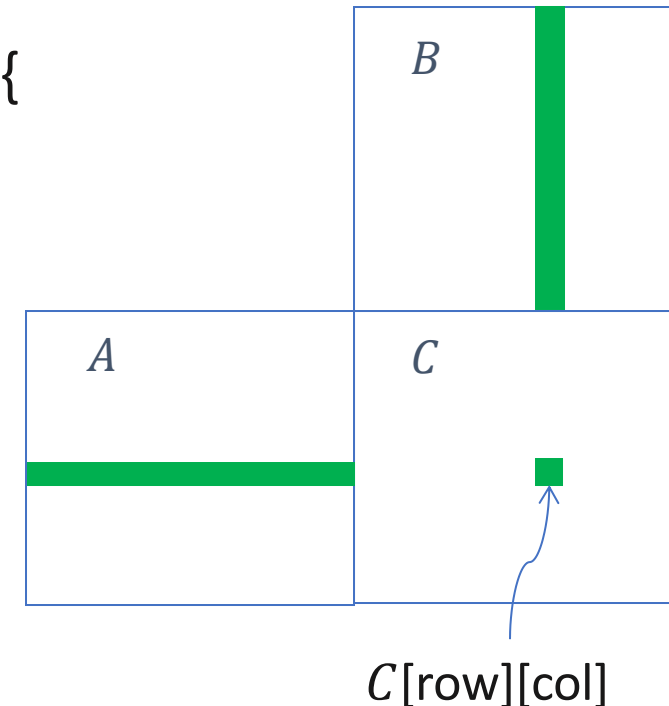
`for (k=0; k<1024; k++)`

`Local_C += A[row][k] * B[k][col];`

`End for`

`C[row][col] = Local_C`

`}`



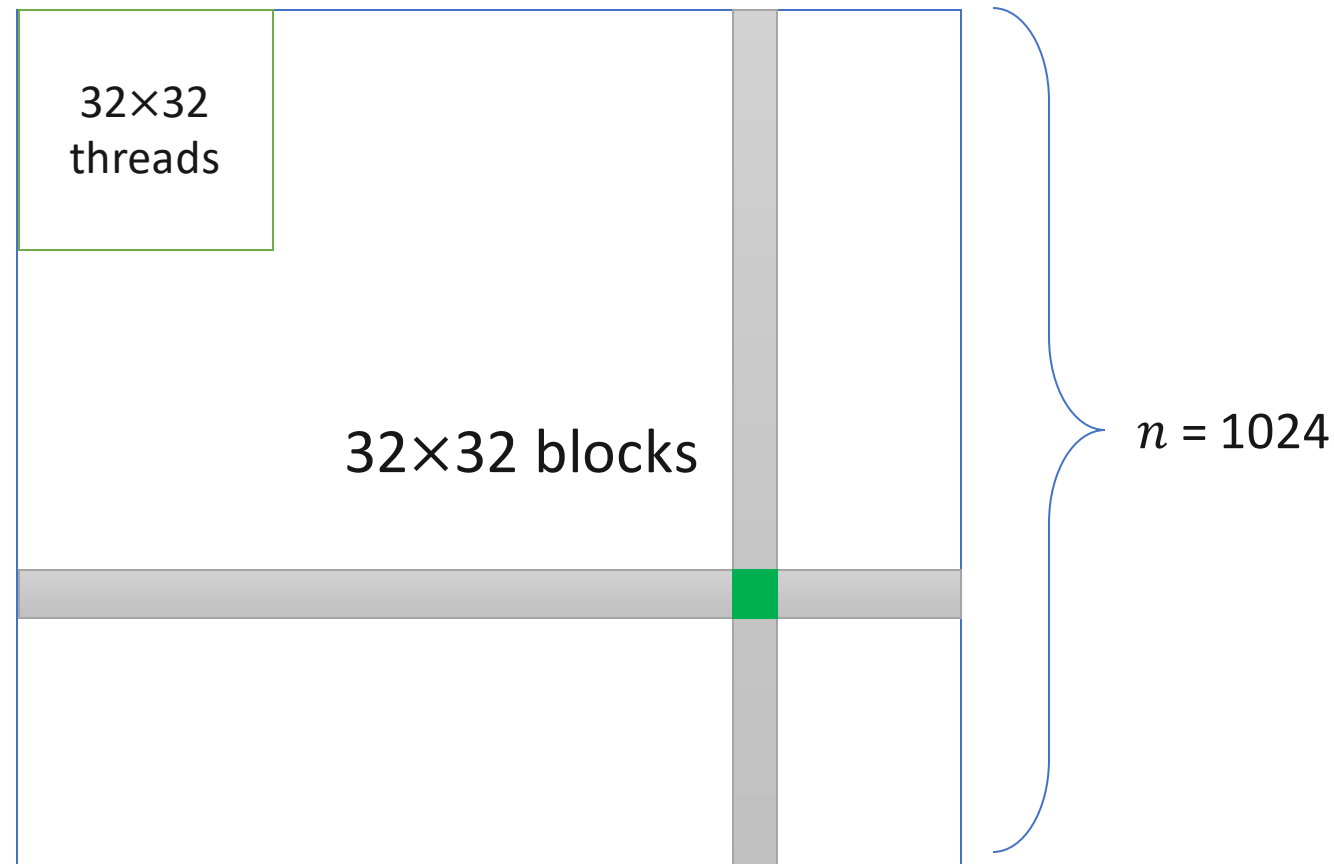
Global memory (data in row major order)



Matrix Multiplication (6)

$$\begin{aligned}\text{row} &= 32 \times \text{blockID.x} + \text{threadID.x} \\ \text{col} &= 32 \times \text{blockID.y} + \text{threadID.y}\end{aligned}$$

Layout:





Matrix Multiplication (7)

1K × 1K MM

- Each thread loads 2×1024 elements from global memory
 - 1024 elements from A
 - 1024 elements from B
- Note: threads may idle for hundreds of clock cycles for each access
- Total data communicated from global memory to device = $n^2[2n]$
- Algorithm does not exploit data sharing (data reuse) among the threads



Block Matrix Multiplication

- **Optimization:** load portions of A and B into **shared memory**
 - 64 KB shared memory can not hold two $1K \times 1K$ matrices
 - Decompose matrices A and B into non-overlapping sub-matrices, eg. 32×32 sub-matrix

(See Discussion Session)



Outline

- Architecture, Programming and Execution Model of (GP) GPU
 - GPU and data parallel programming
 - Programming Model
 - GPU Architecture
 - Execution Model
 - Examples
 - **CUDA, OpenCL, ...**
 - APU, Zynq...



CUDA, OpenCL, ...

- OpenCL
 - Cross-vendor standard for heterogeneous computing
 - Maintained by Khronos Group
 - Provides C/C++ extension and API
- CUDA
 - NVIDIA's parallel computing architecture using GPU
 - Provides CUDA C (C extension) and API
 - Supports OpenCL
- OpenACC
 - Designed to simplify parallel programming of heterogeneous CPU/GPU system



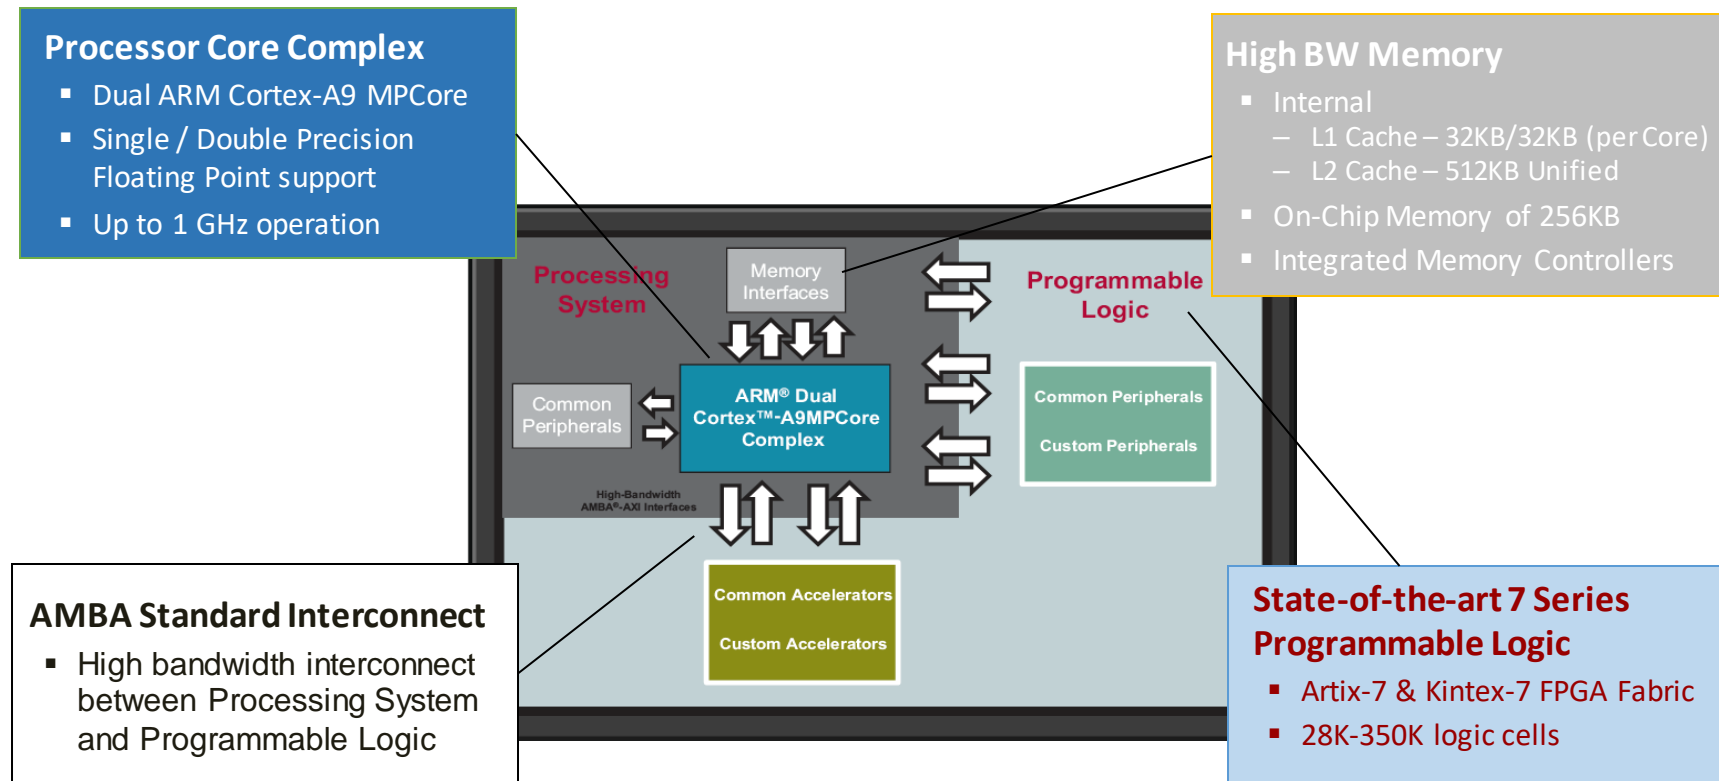
CUDA vs. OpenCL Terms

CUDA	OpenCL
Global Memory	Global Memory
Per-Block Shared Memory	Local Memory
Local Memory	Private Memory
Kernel	Program
Block	Work-group
Thread	Work-item
NVIDIA	NVIDIA, AMD, Intel, ARM ...



Zynq

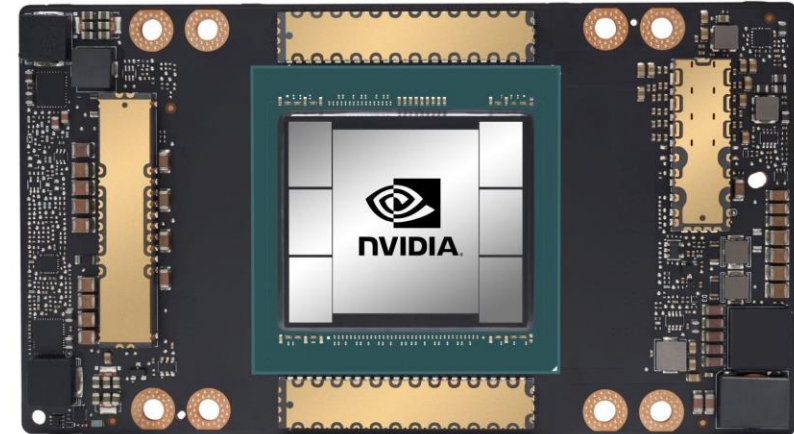
- Zynq-7000 Programmable SoC family
 - Efficient ARM + FPGA





State-of-the-art GPU Example

- GPU family – Nvidia
 - A100 Tensor Core GPU
 - 108 Streaming Multiprocessors
 - 40 MB L2 cache
 - 40 GB high bandwidth memory
 - 1.6 TB/s peak bandwidth
 - 600 GB/s inter-GPU communication bandwidth
 - 54.2 billion transistors





Comparison

AMD



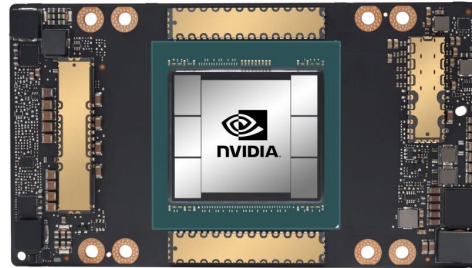
16-64 heavy cores

Up to 40 B transistors

3.4 GHz

225 W

Nvidia



~7 K light cores

54 B transistors

1200 MHz

400 W

Xilinx MPSoC



1-8 M fine grained cores

35 B transistors

Up to 600 MHz

30 W



Summary

- GPU and data parallel programming
- CUDA model
- SIMT
- Kernel → grid, thread blocks, threads
- Explicit parallelism (data parallelism/Block, task parallelism across Blocks)