



EE/CSCI 451: Parallel and Distributed Computation

Lecture #16

10/13/2020

Viktor Prasanna

prasanna@usc.edu

ceng.usc.edu/~prasanna

University of Southern California



Course Project

- Project timeline
 - Week 7-8: Identify team members and project topic, discuss with instructor or TA
 - End of Week 9: Project proposal due (Oct. 16)
 - Week 12-13: Presentation
 - End of Week 13: Project report due (Nov. 13)
- Presentation order would be randomly assigned



Announcement

- PHW5 due on 10/22 (Thursday)
- HW7 due on 10/16 (Friday)
- HW8 will be out on 10/17 and due on 10/22 (1 day before midterm2)
- Final exam date: 2-4 PM Thursday, November 19



Announcement: Midterm 2

- Time: Oct. 23 (Friday) 3:30-5:30 PM
- A sample exam is posted on Piazza!
- Covers material from Sept. 22 to Oct. 16
 - Program mapping questions (covered in Week 6) will be on midterm 2!
- Logistics: same as Midterm 1



Outline

Last class

- GPU and Data Parallel (CUDA) programming model
 - SIMT
 - Streaming multiprocessor
 - Warp
 - CUDA Core


Today

- Parallel algorithm design (Chapter 3.1)
 - Task and dependency
 - Critical path
 - Task dependency graph
 - Mapping



Parallel Algorithm Design

Parallel Algorithm = Concurrency + Coordination



Design issues to be addressed:

- Synchronizing the various activities (processors)
- Identifying tasks that can be performed concurrently
- Mapping Work → processors
- Scheduling When to execute
- Data partitioning and distribution Data → storage (layout)
- Data access management

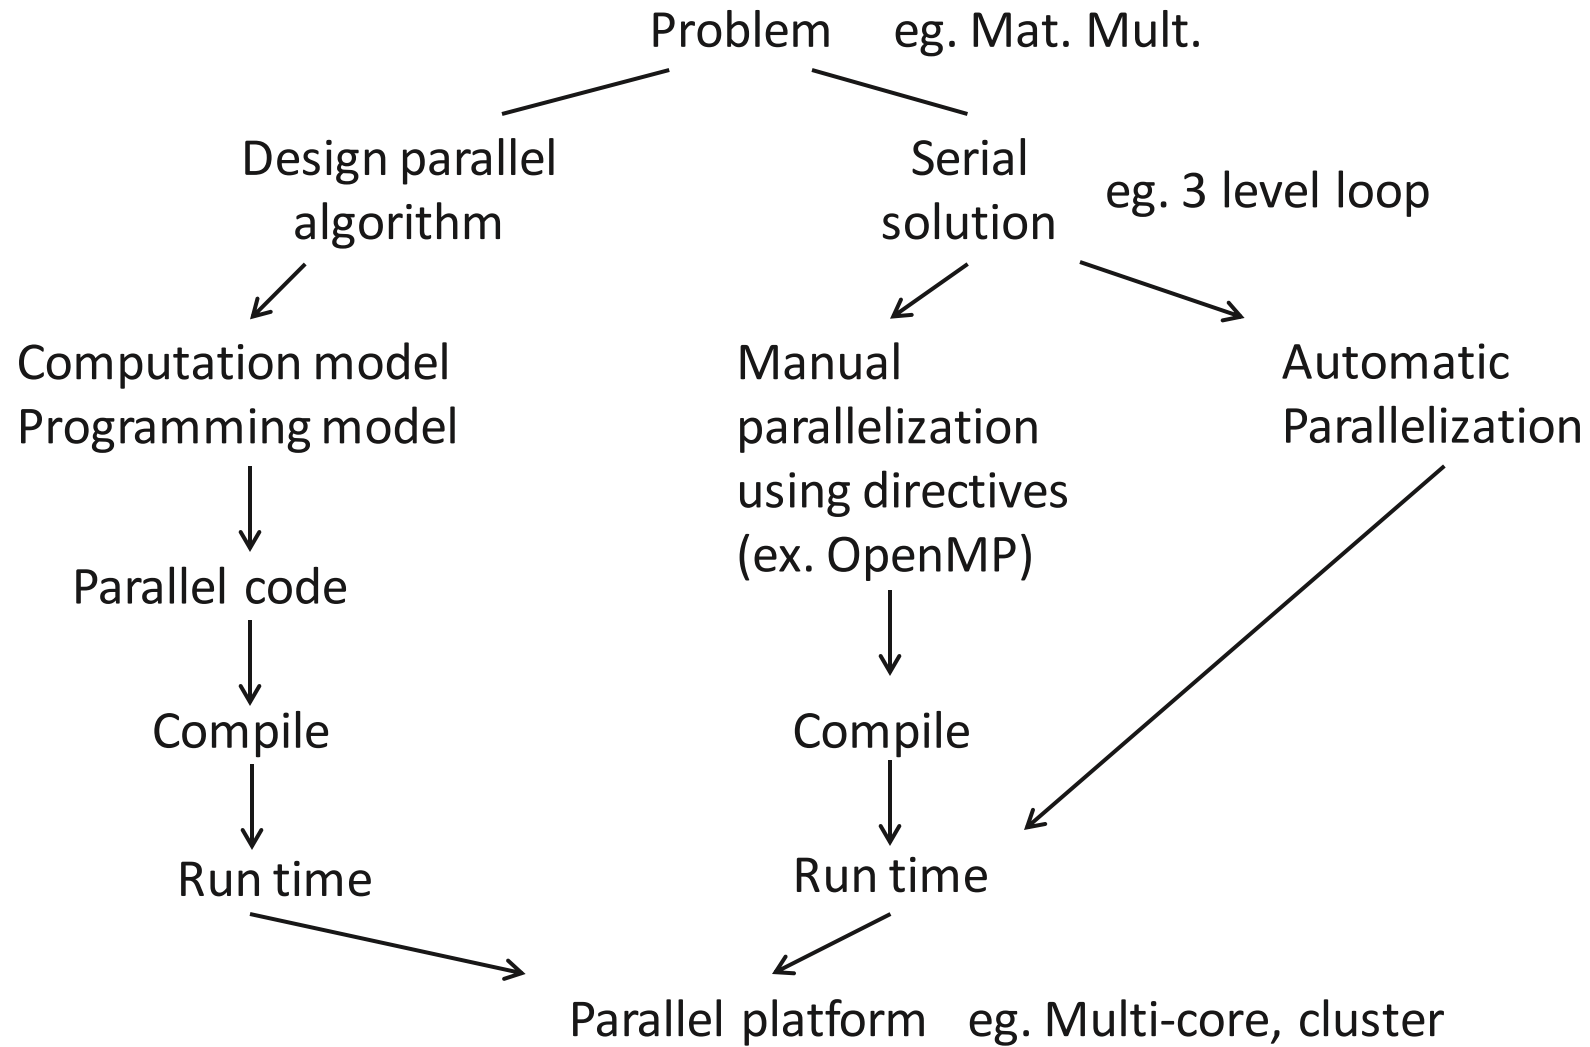


Metrics and Constraints

- Memory footprint
 - # of processors used
 - Cost
-
- Latency, throughput
 - Speed-up
 - Scalability
 - Efficiency
 - Energy efficiency



Approaches






Tasks and Dependencies

Computation = decompose into tasks

Task = Set of instructions (program segment)

Inputs & outputs



Begin execution
once **all** inputs
are available

Task size?

= weight of the node
(e.g., # of instructions executed)

Fine grain

Coarse grain

Note: tasks need not be of the same size



Task Dependency Graph


Directed Acyclic Graph

Task_{*j*} cannot start until Task_{*i*} completes

Data x_1 : output of Task_{*i*}

Input to Task_{*j*}

Weight of a node: task size



Task dependency graph need not be connected
Graph is acyclic




Example (1)

Code

$A[2] = A[0] + 1$

$B[0] = A[2] + 1$



Instructions

Load $R0 \leftarrow A[0]$

Add $R1 \leftarrow R0 + 1$

Add $R2 \leftarrow R1 + 1$

Store $A[2] \leftarrow R1$

Store $B[0] \leftarrow R2$

Tasks

T_0


T_1

T_2

T_3

T_4

Task Dependency Graph





Maximum Degree of Concurrency (1)

Given a task dependency graph,
maximum number of tasks that can be executed
concurrently

Note: maximum degree of concurrency depends
on scheduling strategy



Maximum Degree of Concurrency (2)

Example: **level by level ordering (topological sort)**

Task dependency graph

Order tasks level by level

- Any level i task has dependency with some task in level $i - 1$ (and possibly with other lower levels) but no dependency with level i
- All tasks in any level i are independent

Execute level i tasks (in parallel), and then level $i + 1$ tasks



Maximum Degree of Concurrency (3)

Example: **level by level ordering (cont.)**


Find the number of tasks in each level

Take maximum over all levels = maximum degree of concurrency
(if scheduled using level by level ordering)



Maximum Degree of Concurrency (4)


Example





Maximum Degree of Concurrency (5)

Example (cont.)



Maximum degree of concurrency = 4



Critical Path (1)

Dependency graph

Start nodes ($\text{indeg} = 0$)

Finish nodes ($\text{outdeg} = 0$)



Critical path = A longest path from a start node
to a finish node (# of edges)

Critical path length = Sum of the task weights of the
nodes along the critical path



Critical Path (2)

Note: Critical path length considers critical path only (long paths)






Critical Path (3)

For a given number of tasks,


Longer critical path \Rightarrow Longer execution time
(may also mean less concurrency)



Critical Path (4)




Total no. of tasks = n
Critical path length = n



Total no. of tasks = $n - 1$
Critical path length = $\log_2 n$



Critical Path (5)



Critical path length = $10+6+11+7 = 34$



Task dependency graph to Parallel Program (1)

Given a task dependency graph

Assume weight of each node = 1

$c \Rightarrow$ in degree of each node = constant

Maximum degree of concurrency = c

Critical path length = l

Then, the task dependency graph can be executed on a PRAM ??

of processors?

Time?



Task dependency graph to Parallel Program (2)

Idea

Note: TDG \leftrightarrow DAG

DAG \rightarrow Organize into levels $0, 1, \dots, l$ (level by level ordering)
 $(l + 1)$ levels

Execute level by level, 0 to l

Total number of processors needed $\leq c$



DAG To Parallel Program (3)

Correct parallel program – all dependencies are satisfied

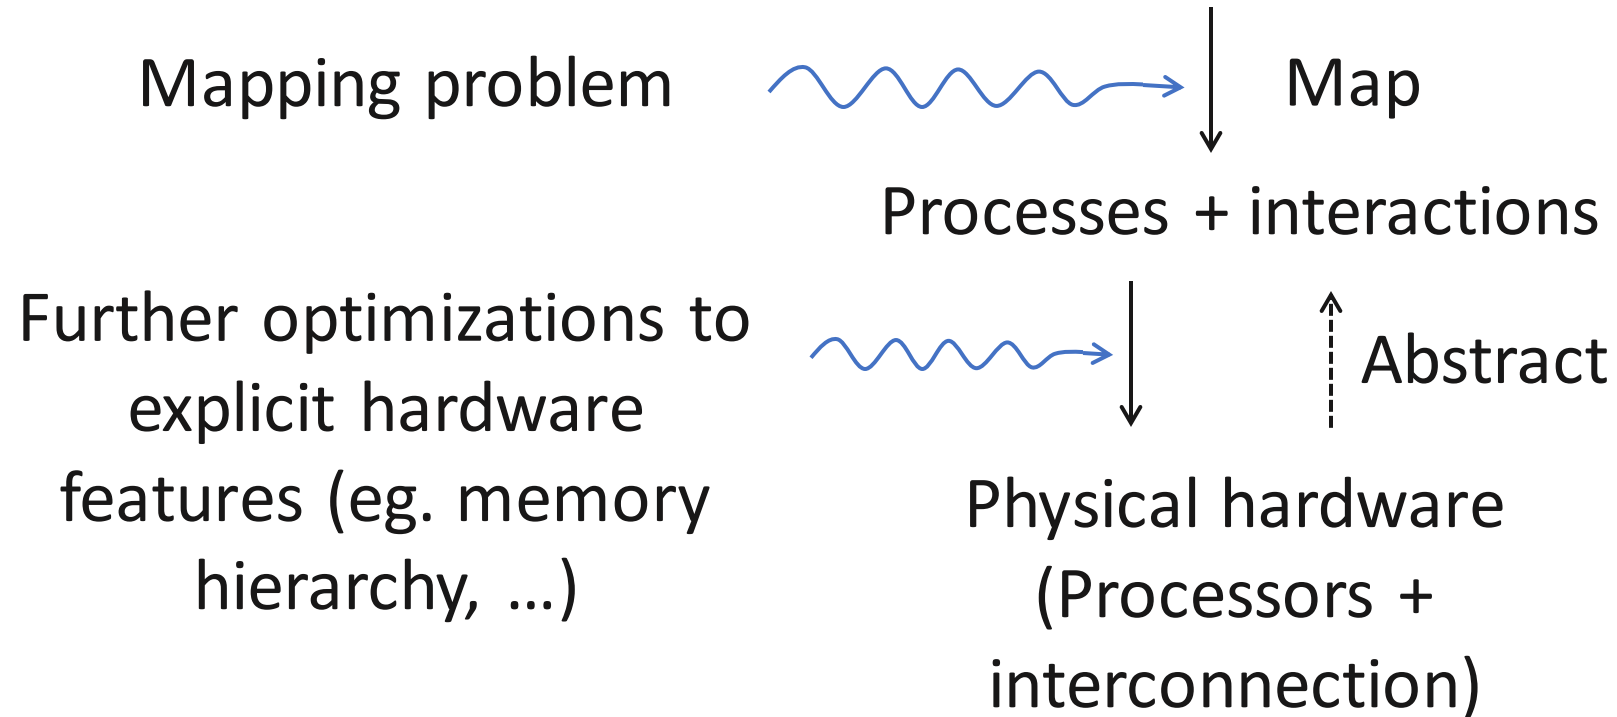
Parallel time $T_p = O(l)$

Using $p = c$ processors



Tasks, Processes, and Mapping (1)

Parallel algorithm → Tasks + interactions






Tasks, Processes, and Mapping (2)

Mapping problem

- Assign each task to a process (processor)
 - Where to execute



Task dependency graph



Tasks, Processes, and Mapping (3)


Scheduling problem

- Determine the execution order of each task
 - When to execute a task (after satisfying the dependencies)




Tasks, Processes, and Mapping (4)

First-come-first-serve
scheduler



Shortest-task-first-serve
scheduler

← List of executable tasks →



Map tasks to processes

← Processes →



Task, Processes, and Mapping (5)

Cluster tasks into processes
(minimize interaction among processes)

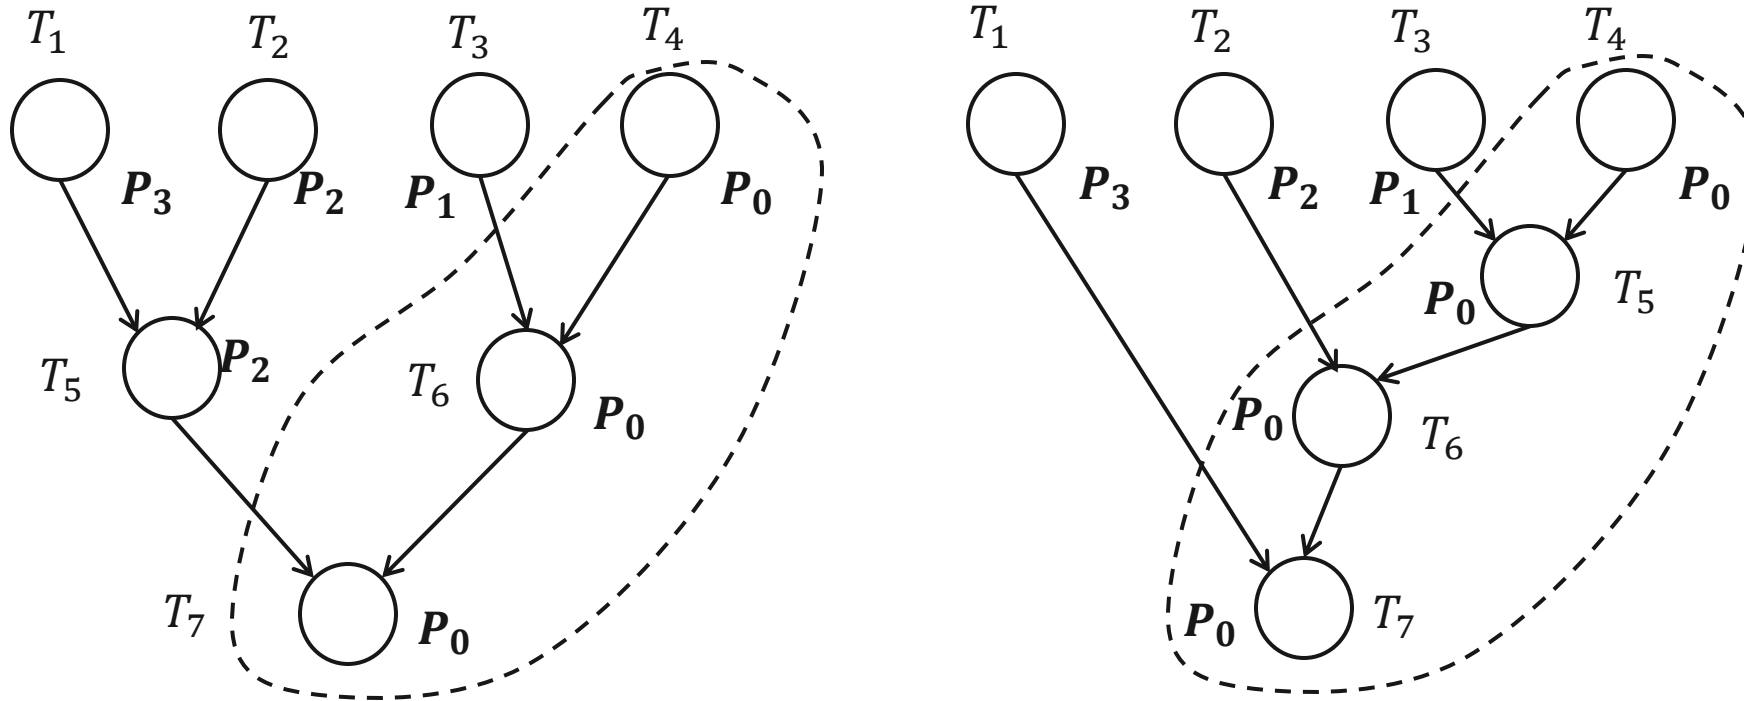
Map processes onto processors
(physical resources)

Process = Code + Data
Execute to produce output



Task, Processes, and Mapping (6)

Mapping tasks to processes






Use of Task Dependency Graph

- Program analysis
- Compiler optimizations (mapping, code generation)
- Automatic parallelization



Example



Pthreads program?



Level by Level Ordering (1)

All nodes with $\text{indeg} = 0 \longrightarrow \text{Level } 0$

$i \leftarrow 1$

Repeat

Delete all nodes with $\text{indeg} = 0$ and all
their outgoing edges

Place all nodes with $\text{indeg} = 0$ at level i

$i \leftarrow i + 1$

Until no nodes left

Can do in $O(n + e)$ time

$n = \# \text{ of nodes}$

$e = \# \text{ of edges}$

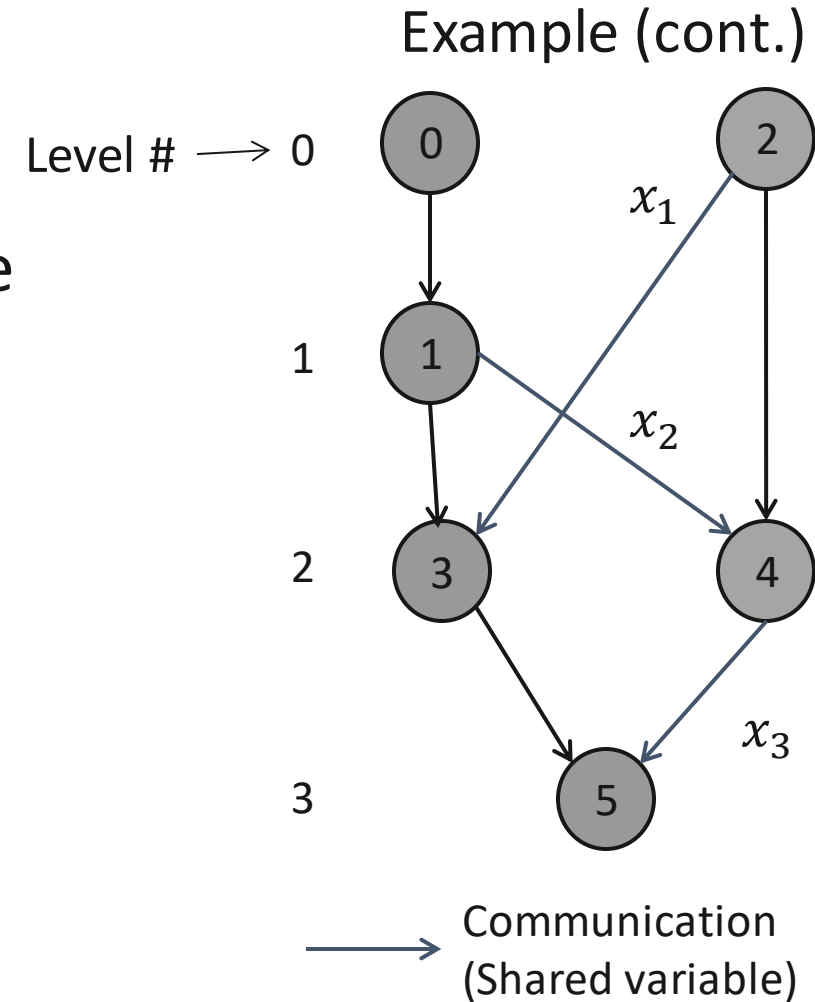


Level by Level Ordering (2)

Thread = Path (chain) from an input node
to an output node


Partition TDG into collection of paths

Insert synchronization primitives
(disjoint)

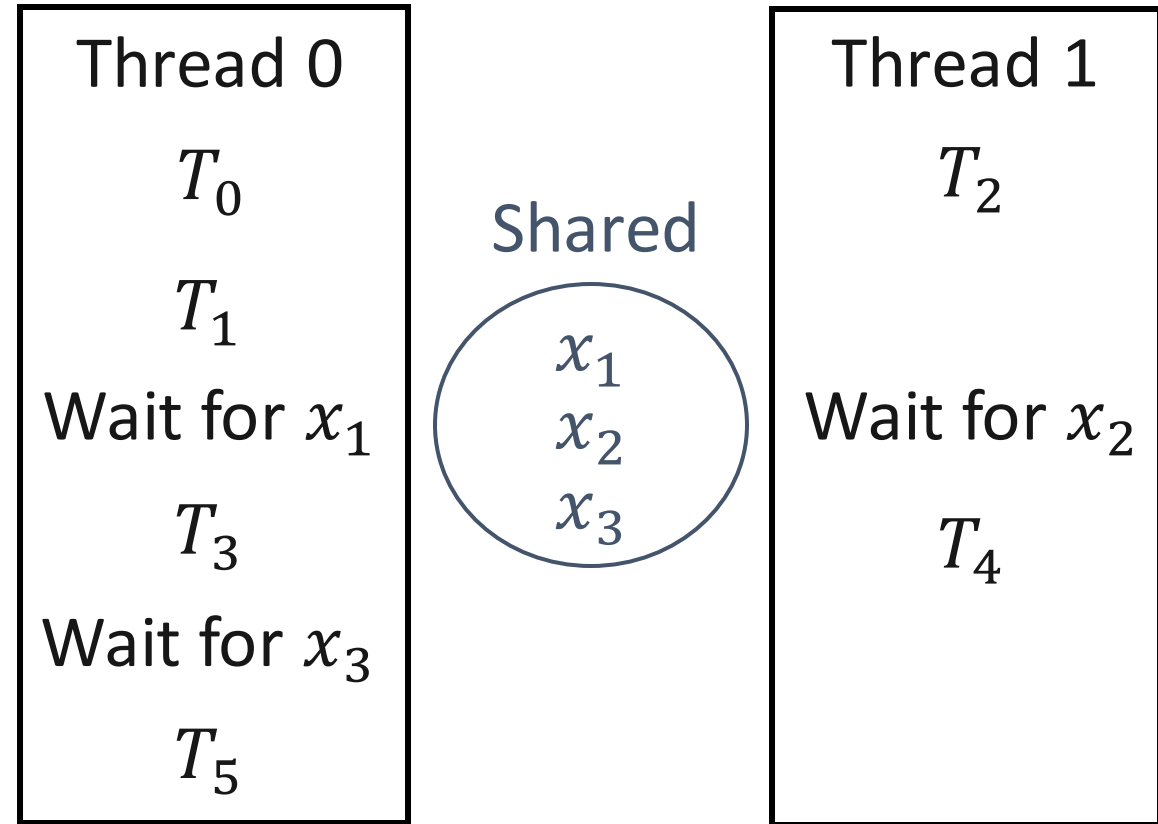




Example (Parallel Program – Shared Memory)



→ Communication
(Shared variable)



Producer Consumer Synchronization



Task dependency graph to parallel program

Level by level ordering

Parallel program

- Execute level 0 tasks (in parallel)

- Barrier

- Communicate

- Execute level 1 tasks

- Barrier

- Communicate

- ...

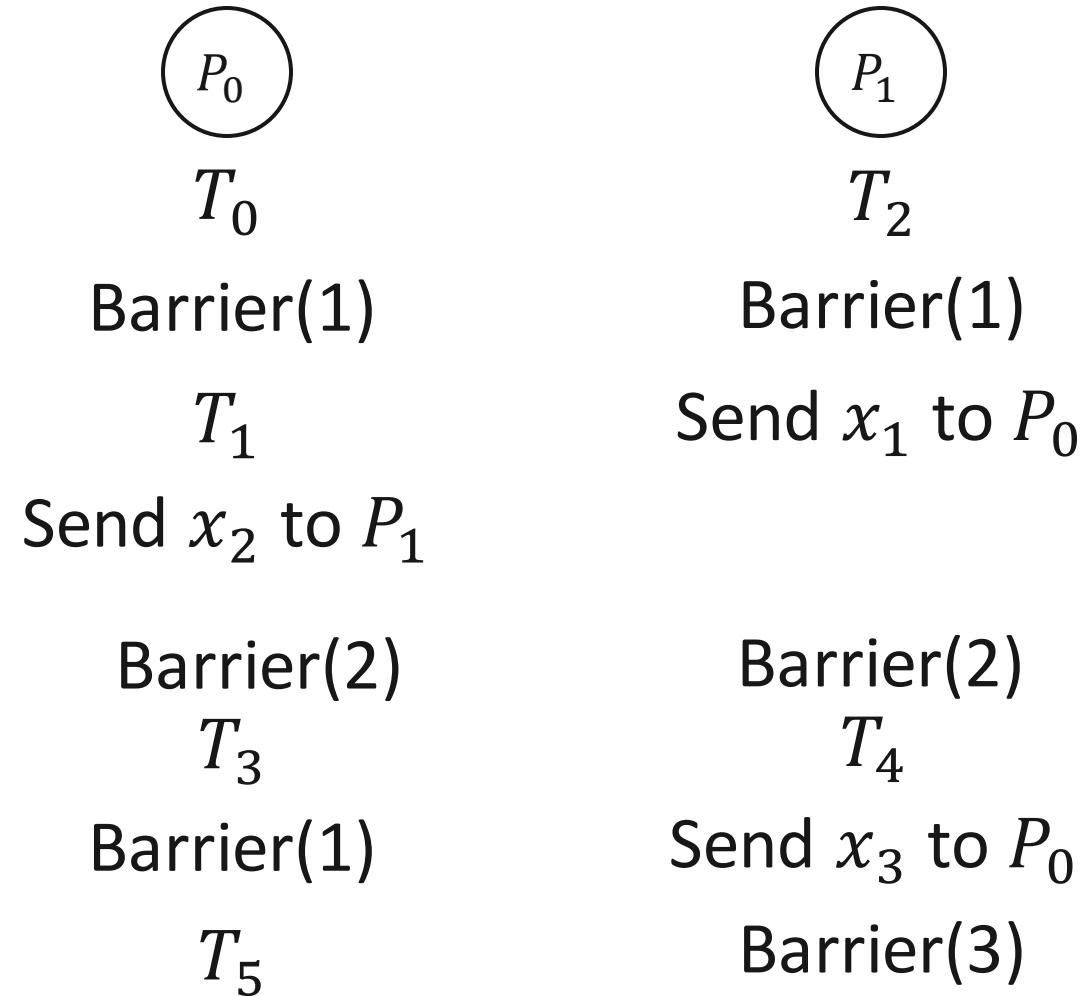
- ...

Map tasks in each level

Schedule tasks in each level



Example (Parallel Program – Message Passing)





Summary

- Parallel algorithm design
 - Task and dependency
 - Critical path
 - DAG to parallel program