



EE/CSCI 451: Parallel and Distributed Computation

Lecture #17

10/15/2020

Viktor Prasanna

prasanna@usc.edu

ceng.usc.edu/~prasanna

University of Southern California



Outline

Last class

- Task Dependency graph
- Critical path
- Max degree of concurrency
- Task Dependency graph (TDG) → Parallel Program

Today

- Decomposition techniques
 - Algorithms
 - Data
- Data distribution (Chapter 3.2)
 - Array data
 - Block distribution
- Graph partitioning
- Mapping (Chapter 3.4)
- Parallel algorithm models (Chapter 3.6)



Announcement

- PHW5 due on 10/22 (Thursday)
- HW6 solution released
- HW7 due on 10/16 (Friday)
- Project Proposal due: 10/18 (extended!)
 - Submit on blackboard
- HW8 will be out on 10/16 and due on 10/22 (1 day before midterm2)
- Final exam date: 2-4 PM Thursday, November 19



Announcement: Midterm 2

- Time: Oct. 23 (Friday) 3:30-5:30 PM
- A sample exam is posted on Piazza!
- Covers material from Sept. 22 to Oct. 16
 - Program mapping questions (covered in Week 6) will be on midterm 2!
- Logistics: same as Midterm 1



Decomposition Techniques

Parallel Solution = Tasks + Concurrency + Interactions

Representation:

Tasks

Dependencies

Interactions

Representation:

Task dependency graph

Task interaction graph



Task interaction graph

Node – task

Edge – interaction between tasks


Type of interactions:

- Static / dynamic
 - ← interactions known at compile time
- Data Access
 - Read data from another task
 - Read/write



Recursive Decomposition

Divide and Conquer






Example (1)

Quick sort

A : n input array

Choose a pivot element x


Partition A using x



Sort A_0 and A_1 in parallel (recursively)



Example (2)



Each node: Choose pivot

Identify data for left and right subtree



Example - Task Dependency Graph (3)

QS (A)

$|A| = 1$ return

Choose pivot x

$A_0 \leftarrow$ those elements in $A \leq x$

$A_1 \leftarrow$ those elements in $A > x$

Do in parallel \longleftarrow **Parallelism**

 QS (A_0)

 QS (A_1)

End



Analysis

Parallel Time

PRAM, n processors

$T_p(n)$ = parallel time for QuickSort using p processes
on n data items

$$T_n(n) = T_{n/2}(n/2) + O(n) \quad \text{best case}$$

$$T_n(n) = O(n) \quad \text{best case}$$

$$T_n(n) = \max\{T_{n-1}(n-1), T_1(1)\} + O(n) \quad \text{worst case}$$

$$T_n(n) = O(n^2) \quad \text{worst case}$$



Example

Merge Sort

$MS(A(0), \dots, A(n-1))$

Sort \longrightarrow

If $|A| = 1$ return

Decompose

Do in parallel

$MS(A(0), \dots, A(n/2-1))$

$MS(A(n/2), \dots, A(n-1))$


End

Merge \longrightarrow

Merge the two sorted sequences of size $n/2$



Task Dependency Graph





Analysis

Parallel time on n processor PRAM

$T_p(n)$ = parallel time for MergeSort using p processes
on n data items

$$T_n(n) = T_{n/2}(n/2) + O(n)$$

$$T_1(1) = O(1)$$

$$T_n(n) = O(n)$$

Serial merge

If we use one processor (Serial merge sort)

$$T_1(n) = \mathbf{2}T_1(n/2) + O(n)$$

Note: Decomposition is fast

Merge takes most of the time



Data Decomposition

Decompose data (partition the data)

Operate on data \longrightarrow tasks

Usually tasks on the partitioned data are similar and can be performed in parallel

Data decomposition \longrightarrow Concurrent tasks \longrightarrow Performance



Data Decomposition Example

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{pmatrix} \cdot \begin{pmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{pmatrix} \rightarrow \begin{pmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{pmatrix}$$

Partition of input and output matrices into 2×2 submatrices

Task 1: $C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1}$

Task 2: $C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2}$

Task 3: $C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1}$


Task 4: $C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2}$

A decomposition of matrix multiplication into four tasks based on the partitioning of matrices above



Owner-Computes Rule

Data



P_0 performs all computations on the input and output data it owns

Ex: Matrix multiplication $C \leftarrow A \times B$

$P_{i,j}$ — owns blocks $A(i, j), B(i, j), C(i, j)$

Computes block $C(i, j)$

Block size $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$, p = total number of processes



Array Distribution Schemes

Widely used strategy

- Data decomposition using input data partitioning
- Owner computes rule

⇒ Defines tasks, mapping to processes

Key problems: Distributing 2-D arrays

Partitioning graph data




Block Distribution (1)


2-D Array \longrightarrow Processes

Block (continuous portion of array) \longrightarrow Process

Example



row-wise distribution



column-wise distribution



Block Distribution (2)

p = number of processes = $p_1 \times p_2$

Block size $n/p_1 \times n/p_2$

P_0	P_1	P_2	P_3
P_4	P_5	P_6	P_7
P_8	P_9	P_{10}	P_{11}
P_{12}	P_{13}	P_{14}	P_{15}

4×4 process grid


P_0	P_1	P_2	P_3	P_4	P_5	P_6	P_7
P_8	P_9	P_{10}	P_{11}	P_{12}	P_{13}	P_{14}	P_{15}

2×8 process grid


Impact of Distribution on Data Sharing (Communication)



Ex.1



Ex.2



Data needed for computing the shaded portion of output matrix (output matrix partitioning)



Load Balancing for MM by Output Partitioning (1)

$n \times n$ matrix multiplication

p processes

Output partitioning : One dimensional: $\frac{n}{p} \times n$

2-dimensional: $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$

Owner Computes Rule:

Each process does the same amount of work (n^3/p)

Total amount of communication varies



Load Balancing for MM by Output Partitioning (2)

Ex.1: 1-dimensional p output processes


$$C \leftarrow A \times B$$

$$P_i \rightarrow \frac{n}{p} \text{ Rows of } A$$
$$\frac{n}{p} \text{ Columns of } B$$

Each process computes $\frac{n}{p}$ output rows

Total Communication = $O(n^2)$ for each output process


Entire B Matrix is needed





Load Balancing for MM by Output Partitioning (3)

Ex.2: 2-dimensional p output processes



$P_i \rightarrow \frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$ blocks of A and B

Each process computes output block of size $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$



$$\begin{aligned} \text{Total Communication} &= O\left(\sqrt{p} \cdot \left(\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}\right)\right) \leftarrow \boxed{\text{Block MM}} \\ &= O\left(\frac{n^2}{\sqrt{p}}\right) \text{ per each output process} \end{aligned}$$




Cyclic and Block Distribution

Problem: distribute 1-D array $n > p$ elements to p processes

Ex. $n = 8, p = 2$



Cyclic



Block



$$\text{Block size} = \frac{n}{p} = 4$$



Block Cyclic Distribution (1)


- A variation of block distribution
- Partition the work into many more blocks (αp) than the number of processes
 - n : problem size (size of input data)
 - p : # of processes
 - L : block size, $= \frac{n}{\alpha p}$ $\{= 1 \text{ to } \frac{n}{p}\}$
 - α : $1 \leq \alpha \leq \frac{n}{p}$
- Distribute blocks in a wraparound fashion
 - Block $b_i \rightarrow P_{i \% p}$ (% is modulo operator)



Block Cyclic Distribution (2)

$$n = 16, p = 4 \text{ processes, } \alpha = 1$$

$$\text{Block size } L = n / \alpha p = 4$$



What if we do more work from left to right?

P_0 does less work compared with P_3




Block Cyclic Distribution (3)

$n = 16, p = 4$ processes, $\alpha = 2$

Block size $L = n / \alpha p = 2$

cyclically distribute




More balanced than $\alpha = 1$



Block Cyclic Distribution (4)

$$n = 16, p = 4 \text{ processes, } \alpha = 4$$

$$\text{Block size } L = n / \alpha p = 1$$



Load balance?


Note: block size = 1 \longrightarrow Cyclic distribution




Example Block Cyclic Distributions

$p = 4$ processes

$n \times n = 16 \times 16$



(a) One dimensional distribution




(b) 2 dimensional distribution



Block Cyclic Distribution

Array size $n \times n$

Number of processes = p





Graph Partitioning (1)

Array data distribution (ex. Block cyclic distribution)

Suited for dense matrices

Other class of problems:

- Sparse data structures
- Interactions are (input) data dependent
- Interactions are **irregular**

⇒ Represented using a (undirected) graph




Graph Partitioning (2)

Example

(Numerical) Simulation of physical phenomenon

Physical domain \longrightarrow Discretized
Mesh of elements



Model of
Lake Superior


Each Point \rightsquigarrow Physical attributes
(eg. Temp, chemical composition,
water flow value, ...)



Graph Partitioning (3)

General Computation

Values at node i at time $t + 1$



$$= f\{\text{Value at time } t \text{ at node } j \mid j \text{ is adjacent to } i\}$$

Note: Computation at each node is the same



Graph Partitioning (4)



Such that

- Load balance
(approx. same number of nodes in each partition)
- Reduce communication cost
Eg. # of edges between partitions

In general, computationally expensive (NP-hard, NP-complete)

Classic problem, many heuristics See METIS software



Mapping based on Task Partitioning (1)

Task-dependency graph (TDG) $\xrightarrow{\text{map}}$ p Processes
(node weights , edge weights)

Goals: Load balance

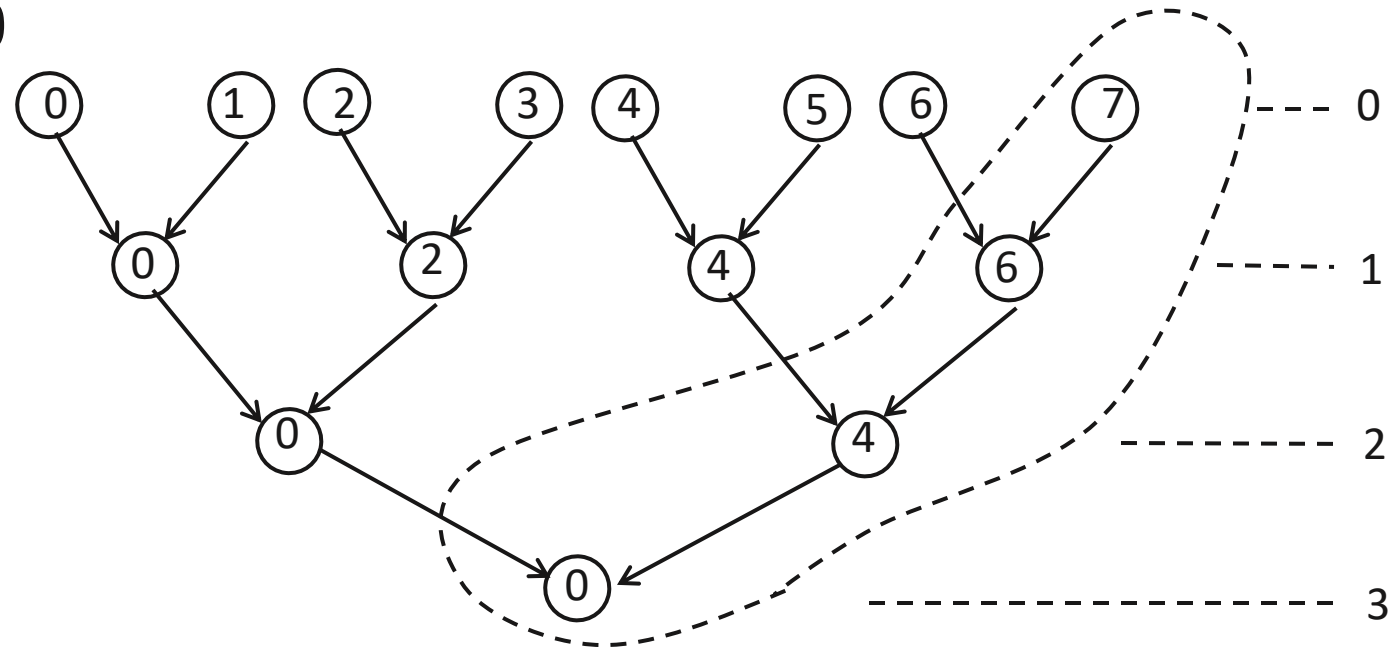
Reduce Idle time

Reduce interaction
(communication) time



Mapping based on Task Partitioning (2)

Process 0




Heuristic:

Level by level ordering

Assign to processes



Mapping for Load Balancing




Objectives

Reduce overhead

Reduce idle time

Reduce inter process communication

Evenly distribute the load among the processes





Static Mapping (1)

Mapping: Tasks \longrightarrow Processes
performed before the execution
of the algorithm begins

Use Task-dependency graph or
Task-interaction graph
+ Meta-data
└─ Task size
 Data size

Most problems are computationally expensive (NP-Complete)



Static Mapping (2)

In most cases,

Static mapping
of tasks



Decomposition
based on data
partitioning

Many scientific computations (eg. Dense matrix algebra)

- Computations are known (statically) at compile time
- Dependency graph is “static”
- Data decomposition is useful in achieving high performance
- Owner-computes rule



Dynamic Mapping (1)

Static mapping may not be effective:

- May lead to unbalanced load
- Task interactions are data dependent



Dynamic Mapping (2)

Distribute tasks among the processes at runtime
(during execution of the algorithm)


May also move data

Run-time system monitors and performs
dynamic task assignment



Dynamic Mapping (3)

Centralized scheme




Master process is responsible to map tasks to processes

Can become the bottleneck



Dynamic Mapping (4)

Self scheduling




When a process becomes idle (completes execution),
access the list and fetch a task to execute

Used for scheduling independent iterations of a loop



Dynamic Mapping (5)

Distributed scheme



Processes maintain list of executable tasks

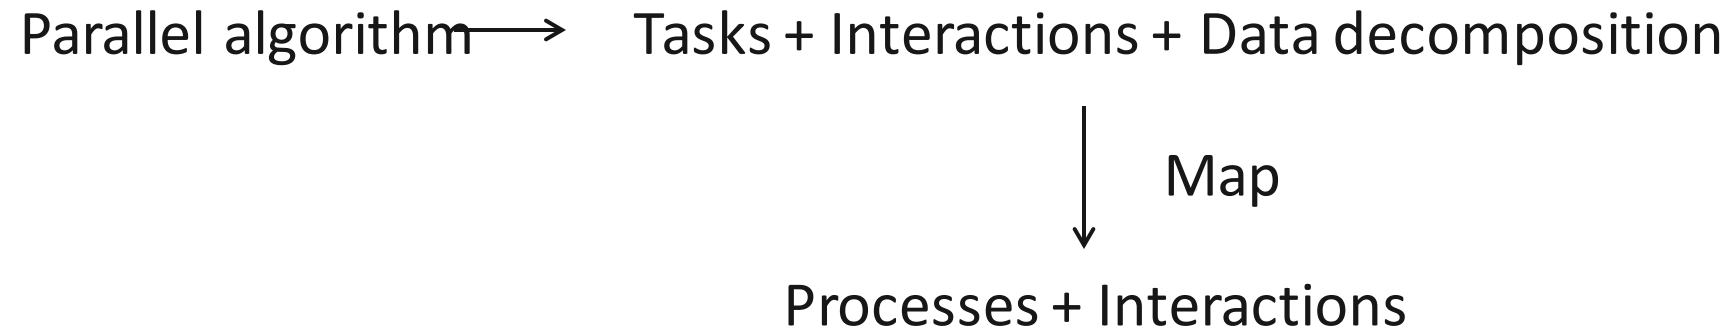
Interact to send/receive work to balance work load

Peer-to-peer system



Parallel Algorithm Models (1)


A parallel algorithm model is a way of structuring a parallel algorithm by selecting a decomposition and mapping technique and applying appropriate strategy to minimize interactions





Parallel Algorithm Models (2)

Data-parallel model (1)



Tasks \longrightarrow processes

Similar tasks operate (in parallel) on different sets of data

- Data partitioning is important to achieve good performance
- Static mapping



Parallel Algorithm Models (3)

Data-parallel model (2)

Examples:

- $$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} \leftarrow \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}$$

Block MM
4 processes

C_{11} ↓ ↓ ↓ ↓ C_{22}

- SIMD execution model

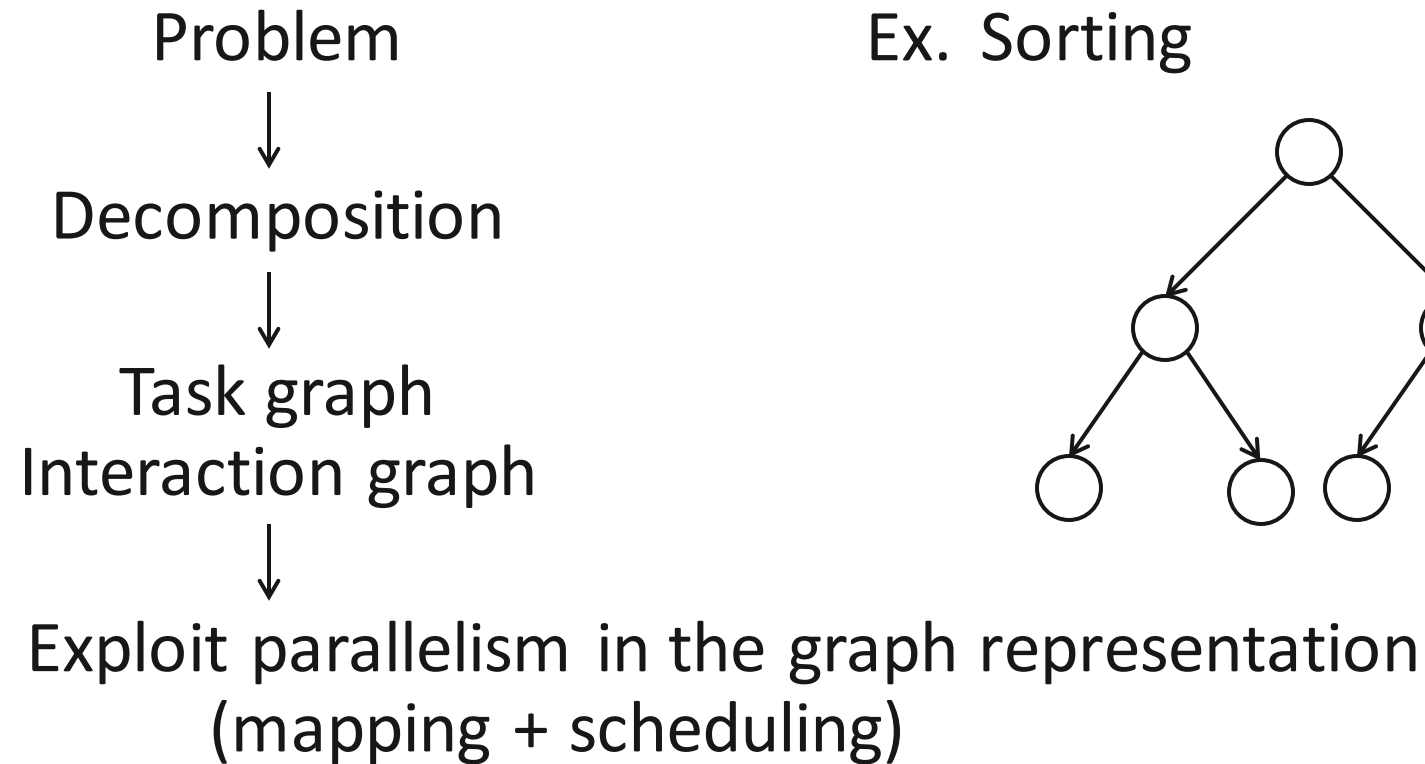
Single Instruction Multiple Data

Synchronous model



Parallel Algorithm Models (4)

Task Parallel Model (Task graph model)





Parallel Algorithm Models (5)

Work Pool Model

Collection of tasks

Any task can be performed by any process

Dynamic mapping of tasks to processes for load balancing


Note: Work pool can be centralized or distributed

Typically: (Small data, large amount of work) per task



Parallel Algorithm Models (6)

Master-Worker Model






Parallel Algorithm Models (7)

Producer-Consumer Model (Pipeline model)

Data (intermediate results) is passed through a series of processes



Stream parallelism

Arrival of data triggers the process

Throughput oriented implementation



Summary

- Parallel Algorithm design
- Data decomposition
- Block Cyclic distribution
- Graph partitioning
- Mapping
 - Static mapping
 - Dynamic mapping
- Parallel algorithm models
 - Task parallel model
 - Work pool model
 - Master-Worker model
 - Producer-Consumer model